

Lecture Notes for **Machine Learning in Python**



Professor Eric Larson
Practical Introductory CNNs

Class logistics and Agenda

- Wide/Deep Lab (late turn in)
- Agenda:
 - CNN Demo
 - CNN Town Hall
 - History of CNNs
 - with Modern CNN Architectures
 - However far we make it...
- Next Time:
 - Transformers

Class Overview, by topic

Table Data
Visualization

Numpy, Pandas, Seaborn
Overviews with some in-depth discussion

Dimension
Reduction and
Image Processing

Scikit-learn, Scikit Image,
Intuition only, Some mathematics

Linear and
Logistic
Regression

Numpy, Recreate API for Scikit-learn
Detailed mathematics for simple optimization
intuition for advanced optimization

Neural Networks
and Back Prop.

Numpy
Detailed mathematics for NN operations

Wide and Deep
Networks

Convolutional
Networks

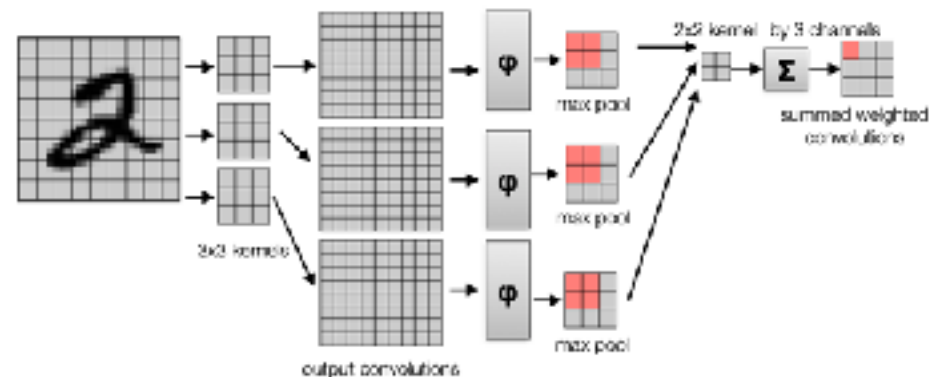
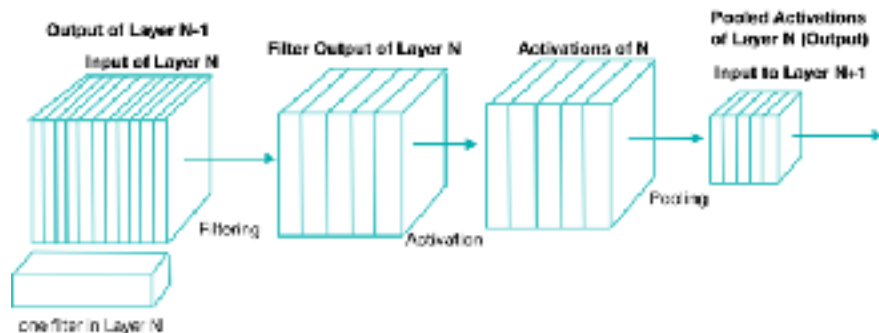
Recurrent
Networks

Keras, Tensorflow
Intuition, Detailed implement.

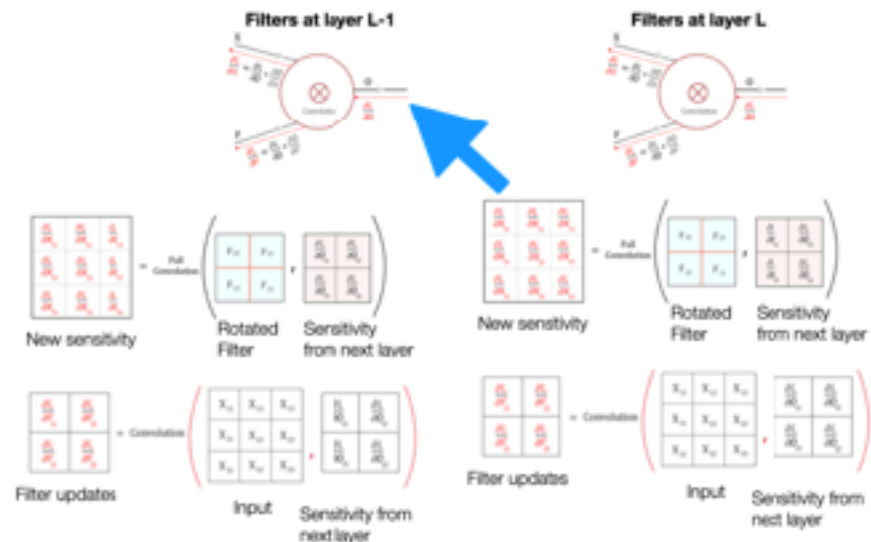

Ethics in
Language Models

ConceptNet
Case studies

Last Time:



Structure of Each Tensor: Channels x Rows x Columns



TensorFlow and Basic CNNs

**Last
Time!**

Convolutional Neural Networks
in TensorFlow
with Keras

with Sequential API!

If needed:

**Finish
Demo**



11. Convolutional Neural Networks.ipynb

Image Data Augmentation

```
cnn = Sequential()  
# add in augmentations directly  
cnn.add( RandomFlip("horizontal") ) # flip horizontally  
cnn.add( RandomRotation(0.05) ) # rotate by 5%  
cnn.add( RandomTranslation(height_factor=0.1, width_factor=0.1) )  
cnn.add( RandomBrightness(factor=0.1, value_range=(0.0, 1.0)) ) #  
cnn.add( RandomContrast(0.1) ) # add or decrease contrast
```




**Image
Augmentation**

RandomRotation()

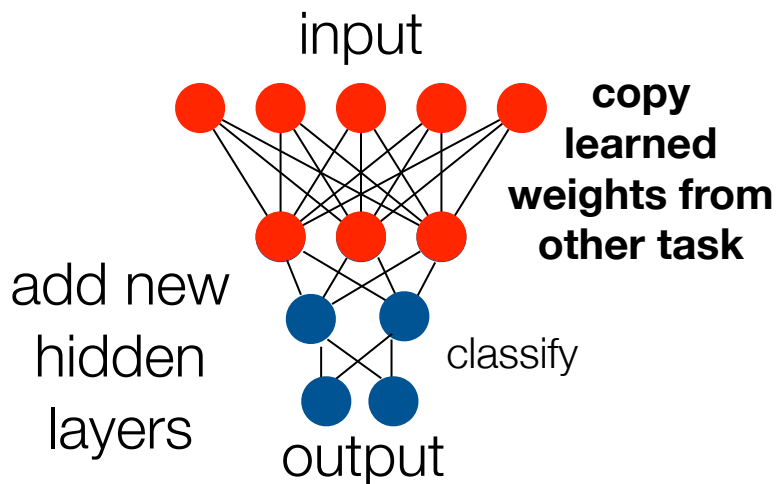
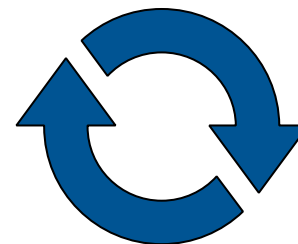


Transfer Learning

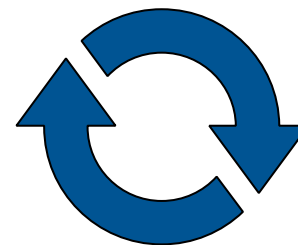
- transfer learning: a basic primer



train with lots of
data (like ImageNet)



train with fewer
labeled data (new task)



Many Pre-trained Models to choose from!

AlexNet

A landmark in computer vision, this 2012 winner of ImageNet has over 50,000 citations.



AlexNet (Places)

The same architecture as the classic AlexNet model, but trained on the Places365 dataset.



Inception v1

Also known as GoogLeNet, this network set the state of the art in ImageNet classification in 2014.



Inception v1 (Places)

The same architecture as the classic Inception v1 model, but trained on the Places365 dataset.



VGG 19

Introduced in 2014, this network is simpler than Inception variants, using only 3x3 convolutions and no branches.



Inception v3

Released in 2015, this iteration of the Inception architecture improved performance and efficiency.



Inception v4


Released in 2016, this is the fourth iteration of the inception architecture, focusing on uniformity.



ResNet v2 50

ResNets use skip connections to enable stronger gradients in much deeper networks. This variant has 50 layers.





CNN Town Hall

Thanks to
Machine Learning the
robot apocalypse was
short lived!

History of Convolutional Neural Networks



Machine Learning 101

Types of CNN, 1988-1998




Yann LeCun
Heads Facebook
AI Team

- **LeNet-1** (1988)
 - ~2600 params, not many layers
- **LeNet-5** (1998)
 - 7 layers, gets excellent MNIST performance
- Major contribution, general structure:
 - conv=>pool=>non-linearity=> ...=>MLP

avg

tanh or sigmoid



CNN History

- List of major breakthroughs from 1998 through 2010 in convolutional networks:



- 2010



Types of CNN, 2010



Dan Ciresan

AI Researcher
IDSA, Switzerland

- **Ciresan Net**
- Publishes code for running CNN via GPU
 - Subsequently wins 5 international competitions
 - from stop signs => cancer detection
- Major contribution: NVIDIA parallelized training algorithms


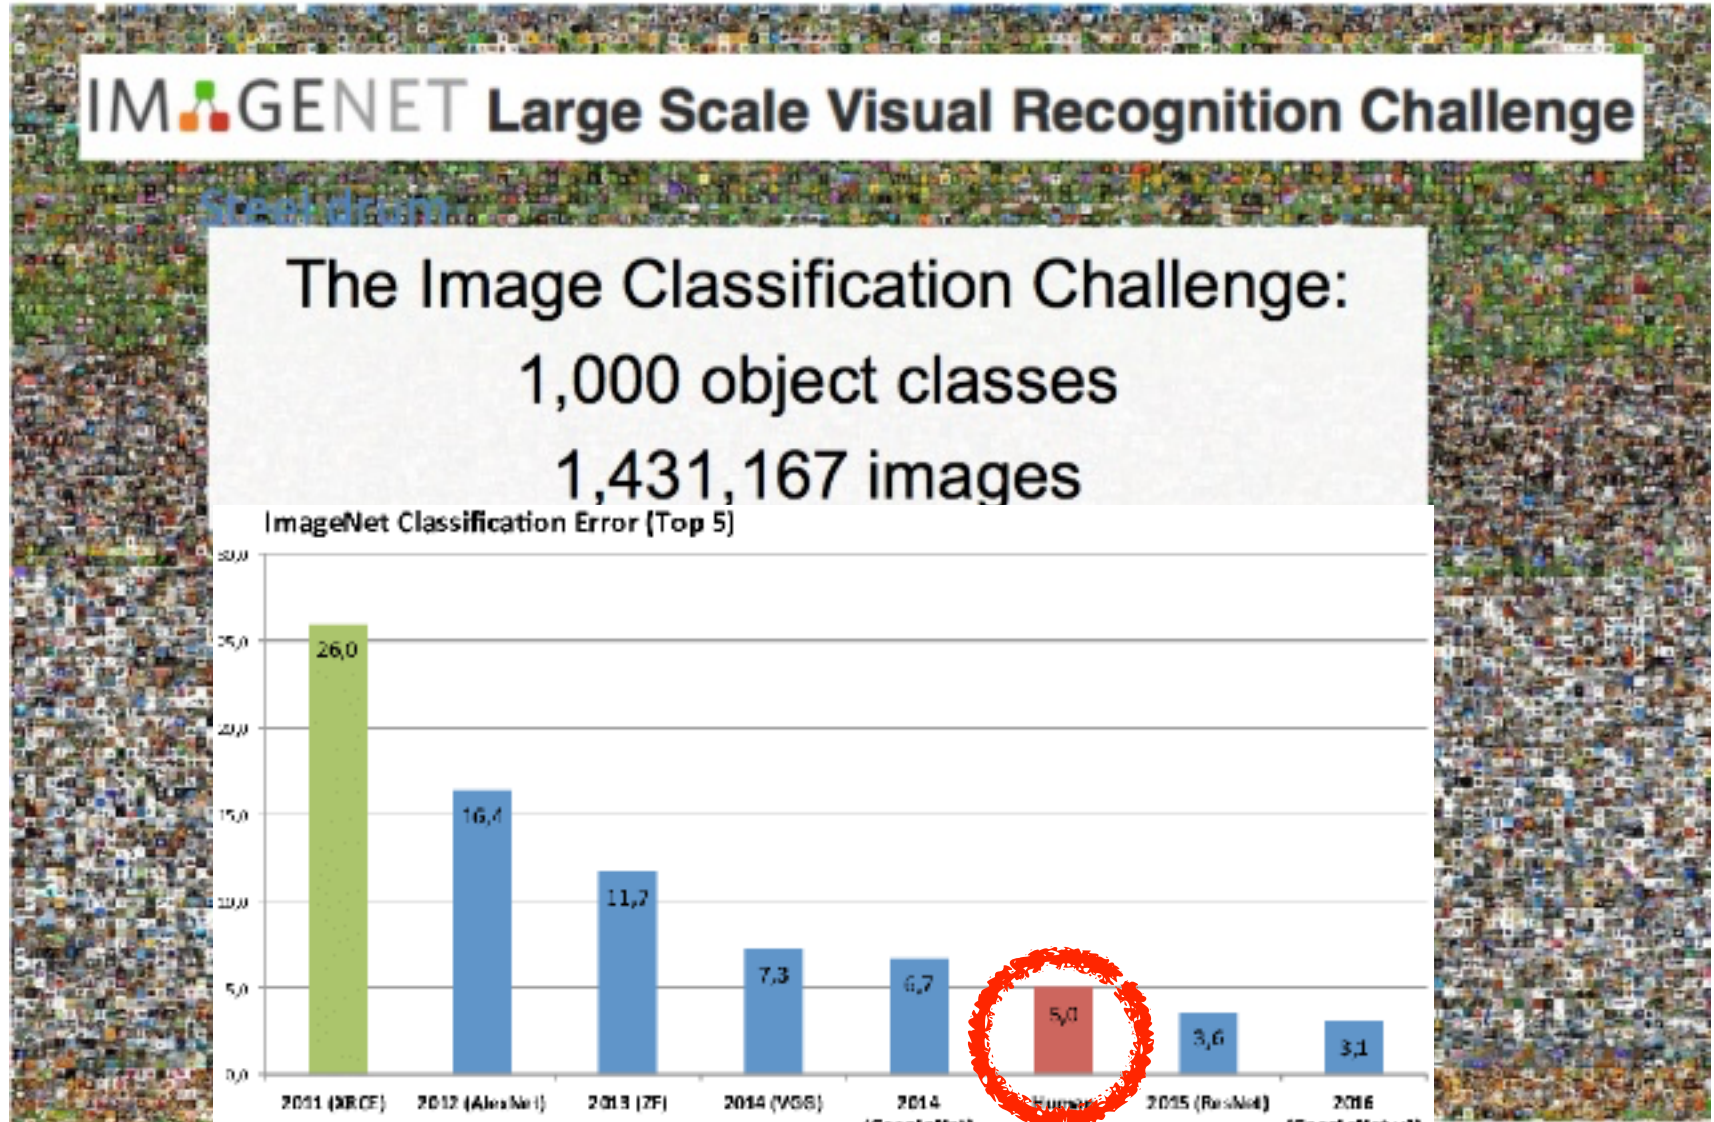


Figure 2: Forward propagation: a) mapping of kernel 1 grid onto the padded weight matrix; b) mapping the kernel 2 grid onto the partial dot products matrix; c) output of forward propagation.


ImageNet Competition (2010-2016)



https://www.researchgate.net/figure/Winner-results-of-the-ImageNet-large-scale-visual-recognition-challenge-LSVRC-of-the_fig7_324476862

<https://www.slideshare.net/nmhkahn/case-study-of-convolutional-neural-network-61556303>


Types of CNN, 2012



Alex Krizhevsky

Google

- **AlexNet**, Hinton is mentor
 - wins ImageNet competition
- Major contributions:
 - dropout for regularization
 - systematic use of ReLU
 - data expansion
 - ***overlapping max pool***



AlexNet

FC 1000

FC 4096 / ReLU

FC 4096 / ReLU

Max Pool 3x3s2

Conv 3x3s1, 256 / ReLU

Conv 3x3s1, 384 / ReLU

Conv 3x3s1, 384 / ReLU

Max Pool 3x3s2

Local Response Norm

Conv 5x5s1, 256 / ReLU

Max Pool 3x3s2

Local Response Norm


Conv 11x11s4, 96 / ReLU

Warning




WeKnowMemes

Types of CNN, 2013



Karen Simonyan
and
Andrew Zisserman



- Oxford **VGG Net** (Visual Geometry Group)
- Major contributions:
 - small cascaded kernels
 - way more layers (19 versus ~7)
 - “emulates” biology “better”
 - trained on NVIDIA GPUs for 2-3 weeks


ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Table 2: Number of parameters (in millions).


Network	A,A-LRN	B	C	D	E
Number of parameters	133	133	134	138	144

- Network in Network **NiN**
 - or MLPConv

Min Lin^{1,2}, Qiang Chen², Shuicheng Yan²
¹Graduate School for Integrative Sciences and Engineering
²Department of Electronic & Computer Engineering
National University of Singapore, Singapore
{linmin, chenqiang, eleyans}@nus.edu.sg




(a) Linear convolution layer



weighted sum
of channels

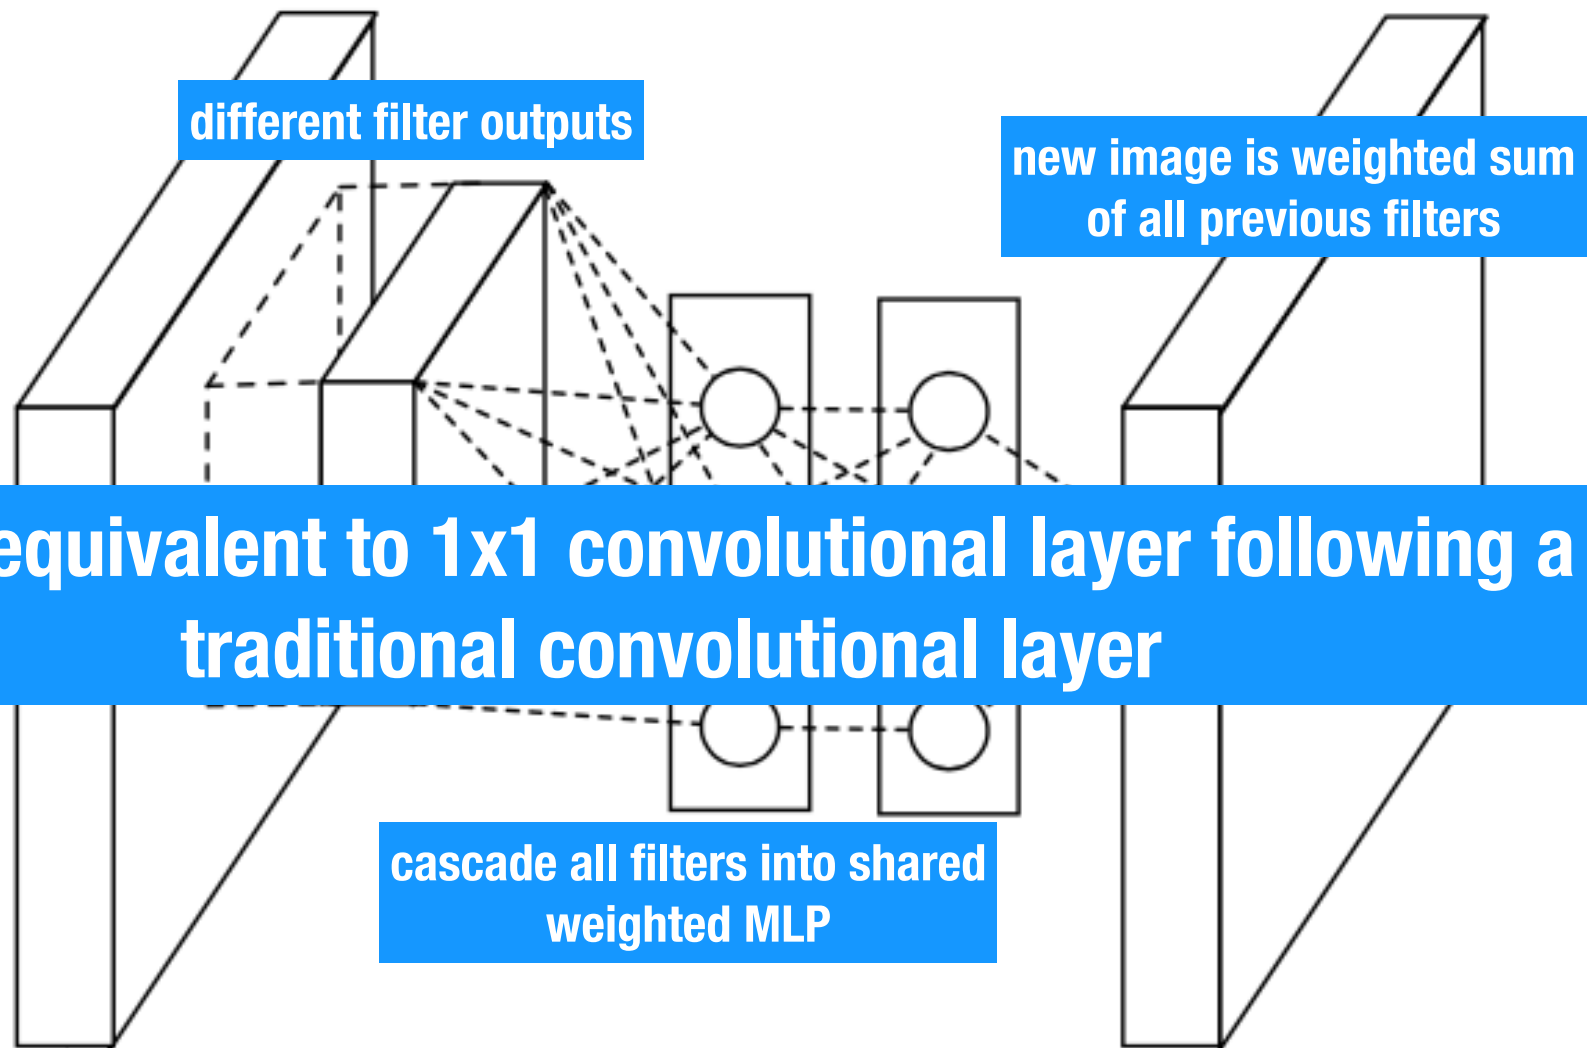
(b) Mlpconv layer



<https://arxiv.org/pdf/1312.4400.pdf>

Types of CNN, 2014


- Network in Network



NiN, expanded view

J and $M \gg K$ and L

Common Choice: $J=M$ and $K=L$




Convolve with K , 1×1 Filters

Equivalently: each new channel is weighted sum of convolutions complete control of channels size



one 1×1 filter input filter



one full filter, $k \times k$



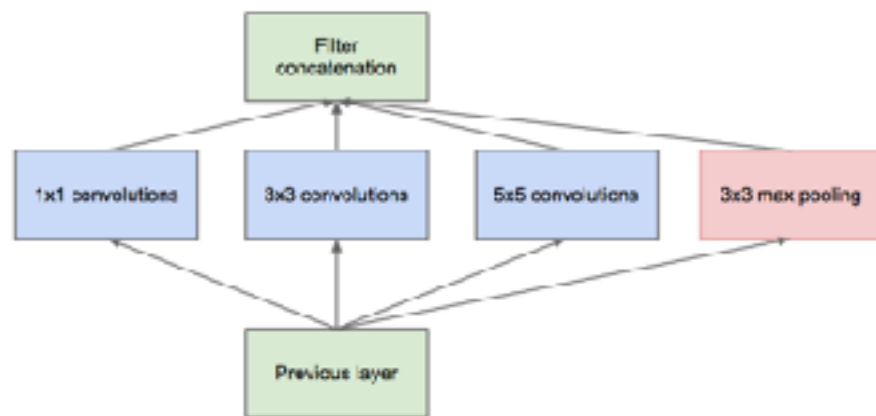
optional one 1×1 filters to control output size

Structure of Each Tensor: Channels x Rows x Columns


Types of CNN, 2014



- **GoogLeNet**
 - or **Inception V1**
- Major contribution:
 - bottleneck layering
 - parallel NiN



(a) Inception module, naïve version



(b) Inception module with dimension reductions





Figure 2: Inception module

Types of CNN, 2015 February and December






- **Inception V2**, Inception V1 with batch normalization
- **Inception V3:**
 - replace 5x5 with multiple 3x3




- Major Contributions:
 - “ensembles” not strictly sequential
 - “bio-plausible” with feedback

- ResNet**


- Parametric ReLU
- PPReLU: adaptive trained slope



- NiN: triple bypass layer
 - similar to bottleneck






Residual Connection, expanded view



Back Propagation: Two paths, including one without ANY operations that cause the gradient to vanish...

How big are these networks?

How big are these networks?



Transition Period in Convolutional Networks

- 2012 - 2017:
 - Add more layers! 🤪
 - How can we train it even deeper? 🧐
 - Can we run out of memory? Let's try! 🤔
- 2017-present:
 - How can we get similar performance with reduced parameters? 🤔
 - How should the number of parameters scale for competing resource? Is there an optimum scaling for a given set of resources? 📈

Types of CNN, 2017

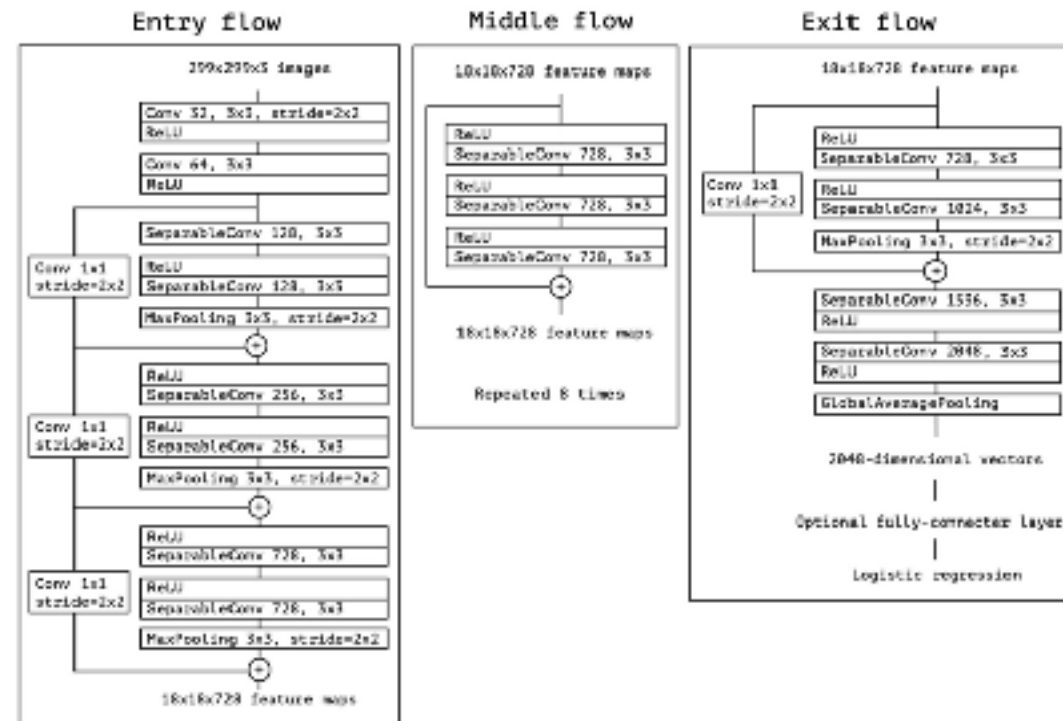
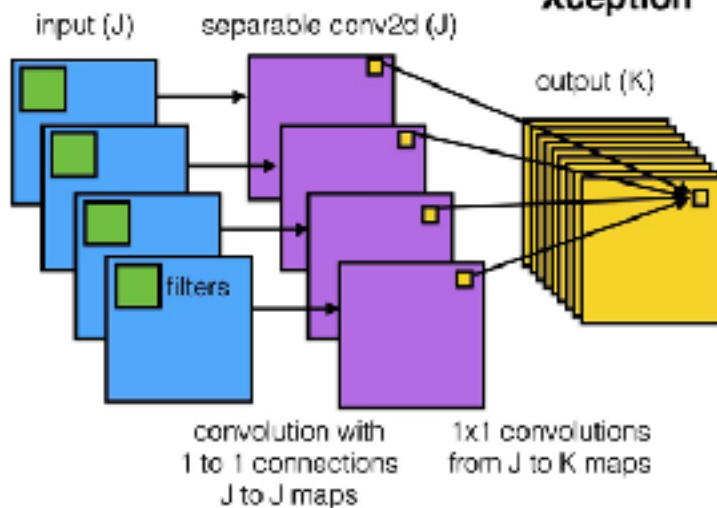
Xception · Major Contributions:

- combining branching / residual blocks
- separable convolutions (fewer trainable params)




Francois Chollet
Google

Xception




<https://arxiv.org/pdf/1610.02357.pdf> 62

Separable Convolution Explanation




Inputs, From Layer N-1
Num Channels = J



Filters, Layer N
Convolve Each Channel Separately

Trainable params:
Same as one filter in traditional convolution!




Concat Outputs
Num Channels = J

Perform K, 1x1 Traditional Convolutions
K Outputs


Trainable params: **K x J**

K Outputs




SqueezeNet (2018)

- Idea: squeeze and expand in each layer
 - Use mostly 1x1 filters
 - downsample later in network



$$SR = \frac{|F_{s1x1}|}{|F_{e1x1}| + |F_{e3x3}|}$$

$$PCT_{3x3} = \frac{|F_{e3x3}|}{|F_{e1x1}| + |F_{e3x3}|}$$



SQUEEZENET: ALEXNET-LEVEL ACCURACY WITH 50X FEWER PARAMETERS AND <0.5MB MODEL SIZE

Forrest N. Iandola¹, Song Han², Matthew W. Moskewicz¹, Khalid Ashraf¹, William J. Dally², Kurt Keutzer¹

¹DeepScale^{*} & UC Berkeley ²Stanford University

{forresti, moskewicz, kashraf, keutzer}@eecs.berkeley.edu

{songhan, dally}@stanford.edu

In paper:

- Good SR = 12.5% up to 100%
- Good PCT_{3x3} from 25% up to 100%

Efficient Net (2019)

Start with so

Observation 1 – Scaling up any width, depth, or resolution improve accuracy gain diminishes for bigger mo

Observation 2 – In order to pursue efficiency, it is critical to balance all width, depth, and resolution during

Depth Scaling

Resolution Scaling: If we use larger resolution

depth: $d = \alpha^\phi$

width: $w = \beta^\phi$

res.: $r = \gamma^\phi$

s.t. $\alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$

$\alpha, \beta, \gamma \geq 1$

ϕ user specified scaling coefficient

$\alpha = 1.2$

$\beta = 1.1$

$\gamma = 1.15$




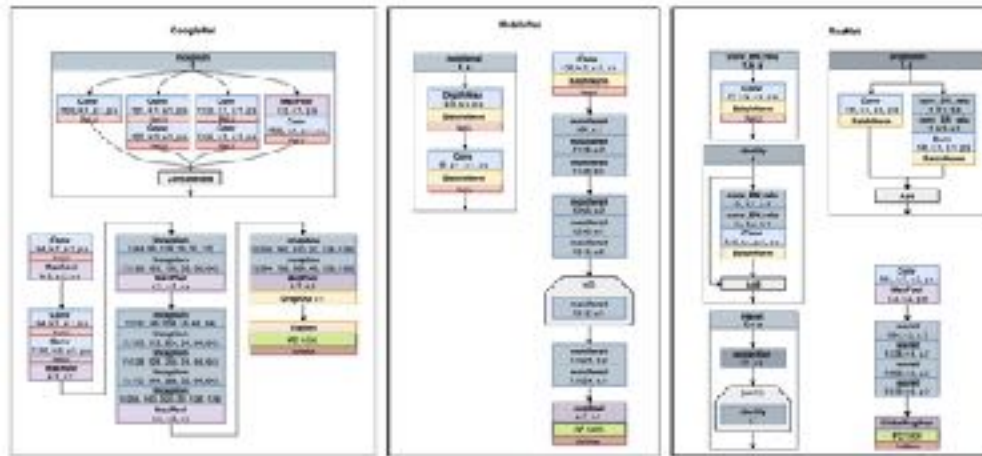
Figure 8. Scaling Up EfficientNet-B0 with Different Methods.

where α, β, γ are constants that can be determined by a small grid search. Intuitively, ϕ is a user-specified coefficient that controls how many more resources are available for model scaling, while α, β, γ specify how to assign these extra resources to network width, depth, and resolution re-

optimal values found in paper!

Even more Convolutional
Neural Networks
...in TensorFlow
...with Keras

Self Guided Demo



12. More Advanced CNN Techniques as TFData.ipynb

Next Time:

- Intro to Sequential Neural Network Architectures
 - Word Embeddings, 1D CNNs, Transformers
 - Ethics by Case Study