

Lecture Notes for **Machine Learning in Python**

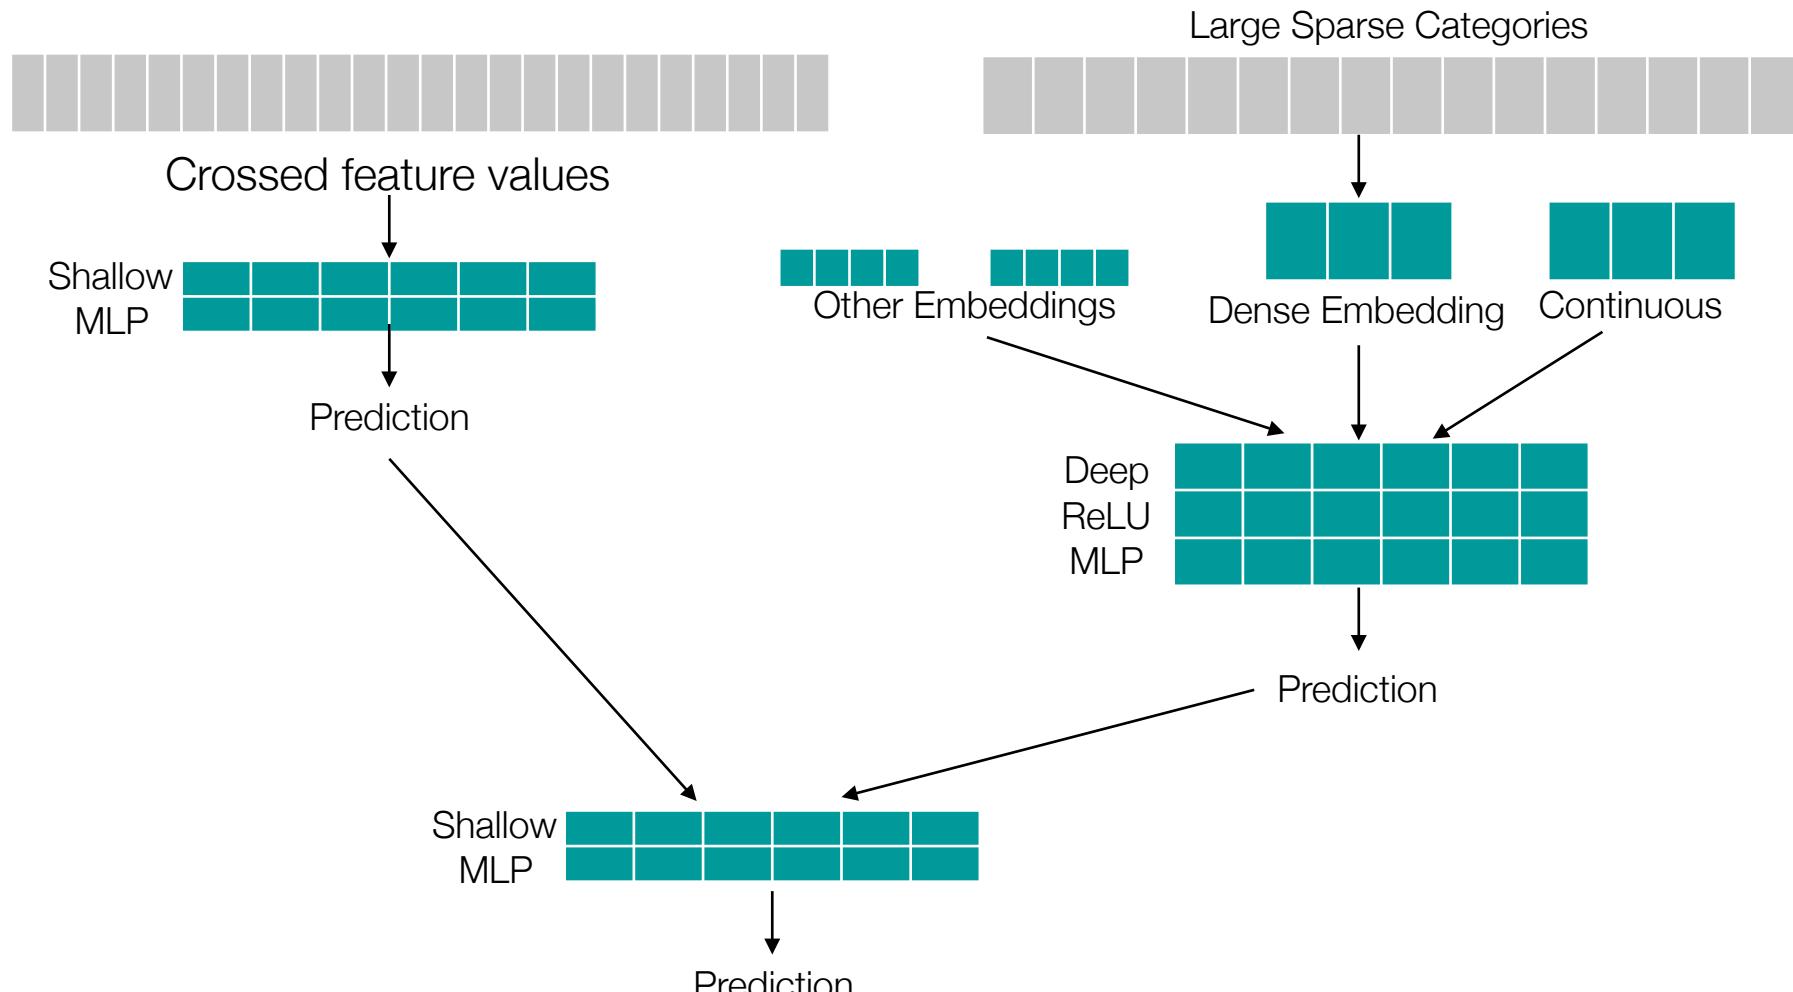
Professor Eric Larson
Basic Convolutional Neural Networks

Logistics and Agenda

- Logistics
 - Wide/Deep due date coming up!
- Agenda
 - Finish Town Hall
 - Basic CNN architectures and Demo

Last Time:

- Deep refers to increasingly smaller hidden layers
- Embed into sparse representations via ReLU



Town Hall, Wide and Deep



WHEN VISITING A NEW HOUSE, IT'S
GOOD TO CHECK WHETHER THEY HAVE
AN ALWAYS-ON DEVICE TRANSMITTING
YOUR CONVERSATIONS SOMEWHERE.

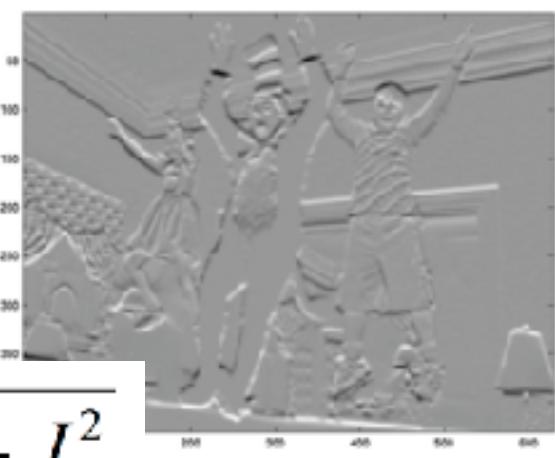
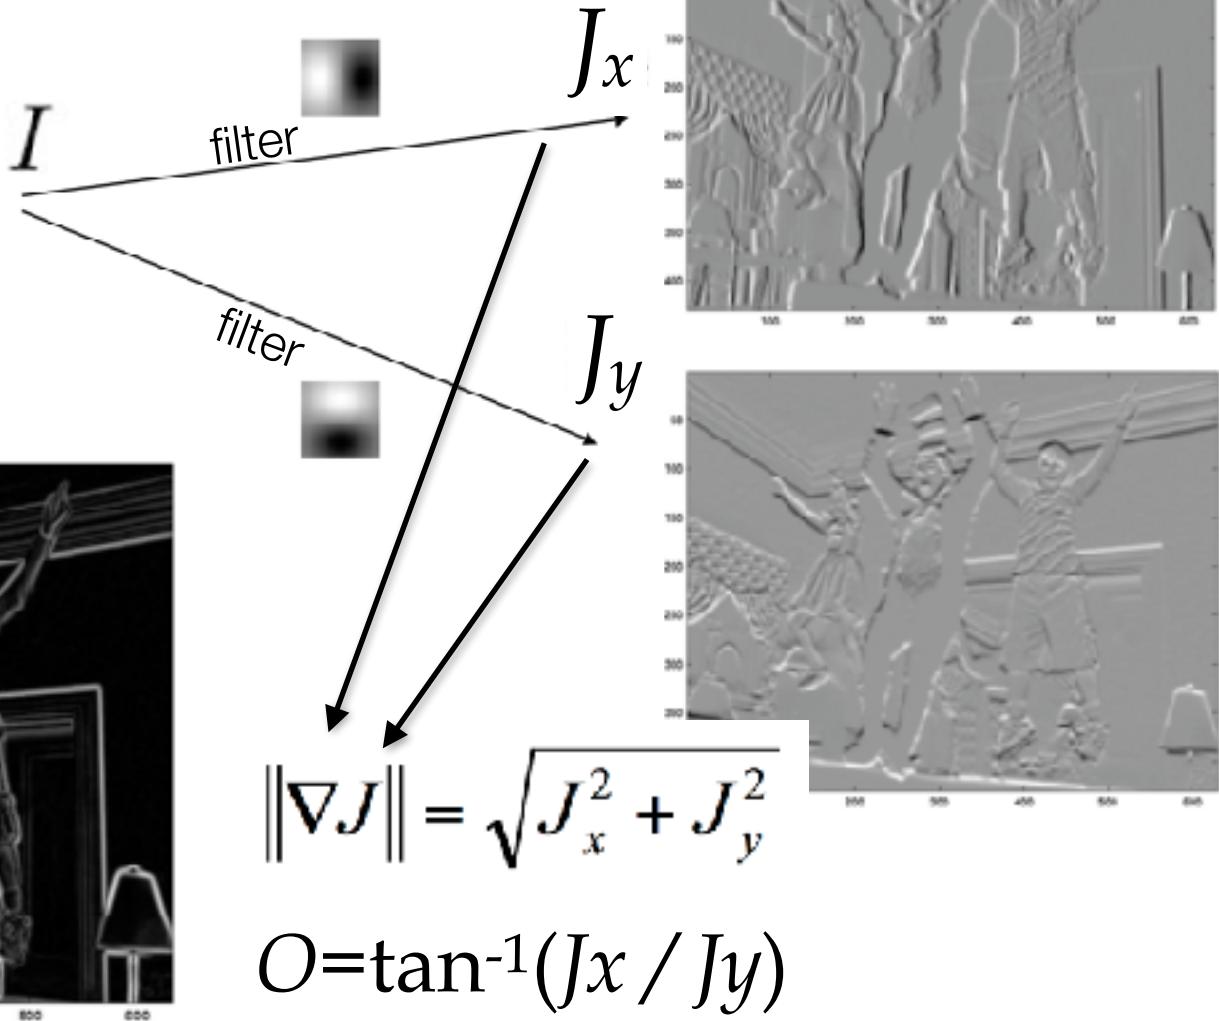
Convolutional Neural Networks



IN CS, IT CAN BE HARD TO EXPLAIN
THE DIFFERENCE BETWEEN THE EASY
AND THE VIRTUALLY IMPOSSIBLE.

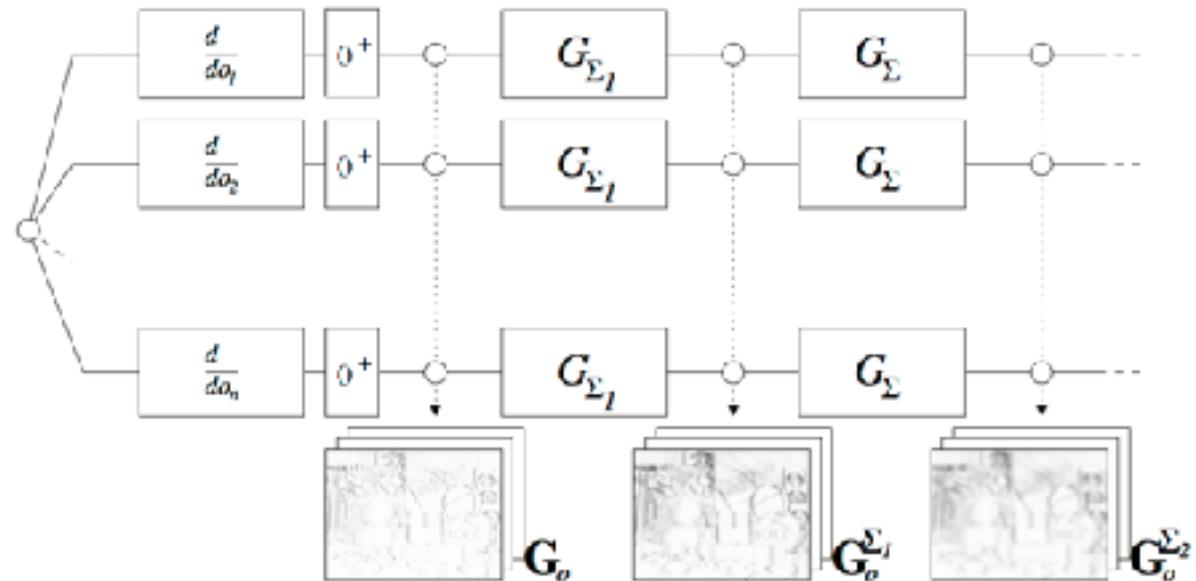
What we did before

- the gradient (2D derivative)



images: Jianbo Shi, Upenn

What we did before



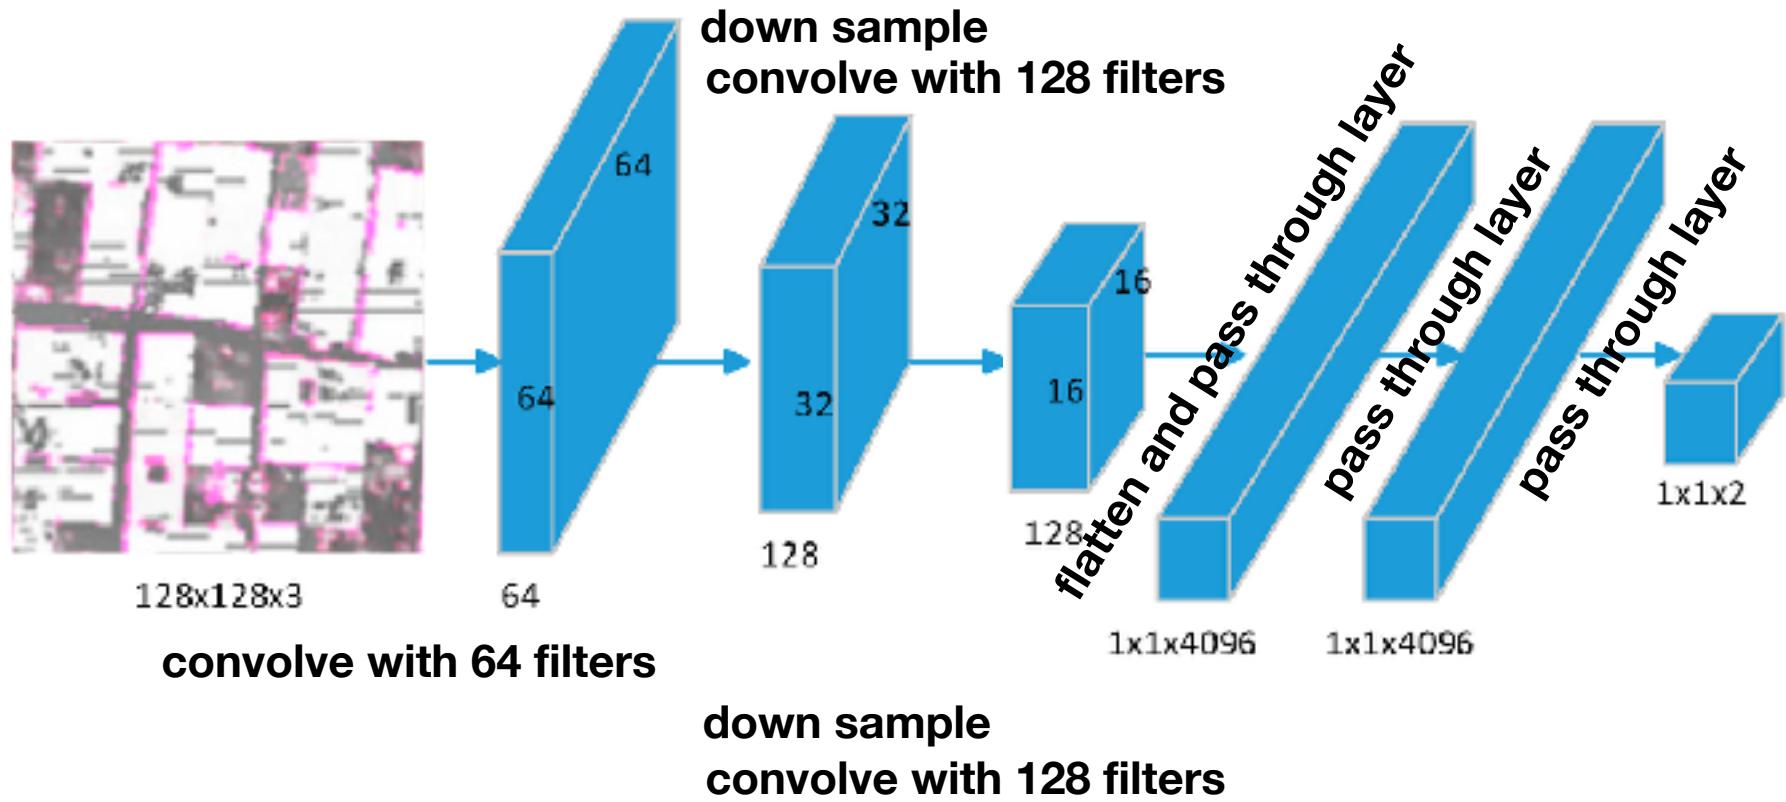
take normalized histogram at point u, v

$$\tilde{\mathbf{h}}_{\Sigma}(u, v) = \left\| [\mathbf{G}_1^{\Sigma}(u, v), \dots, \mathbf{G}_H^{\Sigma}(u, v)]^T \right\|$$

$$\begin{aligned} \mathcal{D}(u_0, v_0) = \\ \left[\begin{aligned} & \tilde{\mathbf{h}}_{\Sigma_1}^{\top}(u_0, v_0), \\ & \tilde{\mathbf{h}}_{\Sigma_1}^{\top}(\mathbf{l}_1(u_0, v_0, R_1)), \dots, \tilde{\mathbf{h}}_{\Sigma_1}^{\top}(\mathbf{l}_T(u_0, v_0, R_1)), \\ & \tilde{\mathbf{h}}_{\Sigma_2}^{\top}(\mathbf{l}_1(u_0, v_0, R_2)), \dots, \tilde{\mathbf{h}}_{\Sigma_2}^{\top}(\mathbf{l}_T(u_0, v_0, R_2)), \end{aligned} \right] \end{aligned}$$

Tola et al. "Daisy: An efficient dense descriptor applied to wide-baseline stereo." Pattern Analysis and Machine Intelligence, IEEE Transactions

Anatomy of a convolution

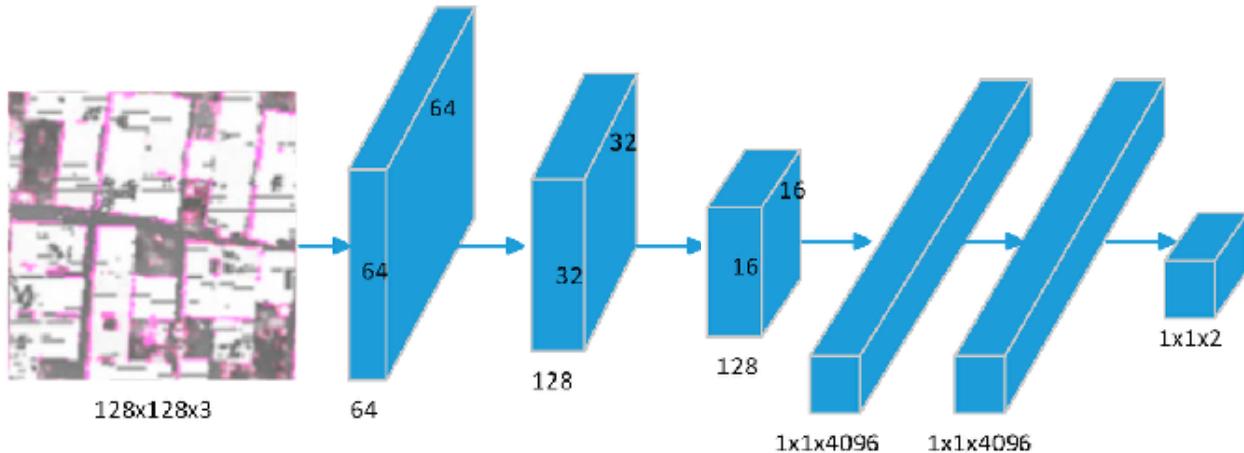


Blue Tensors: Outputs of Each Layer

Learned Params: Weights in Each Filter and Fully Connected Layer

CNN Overview

- First layer(s):
 - convolution
 - nonlinearity
 - pooling
 - Each pooling layer can make the input image “smaller”
 - allows for “Information Distillation”
 - less dependence on exact pixels
- Final layers are densely connected
 - typically multi-layer perceptrons

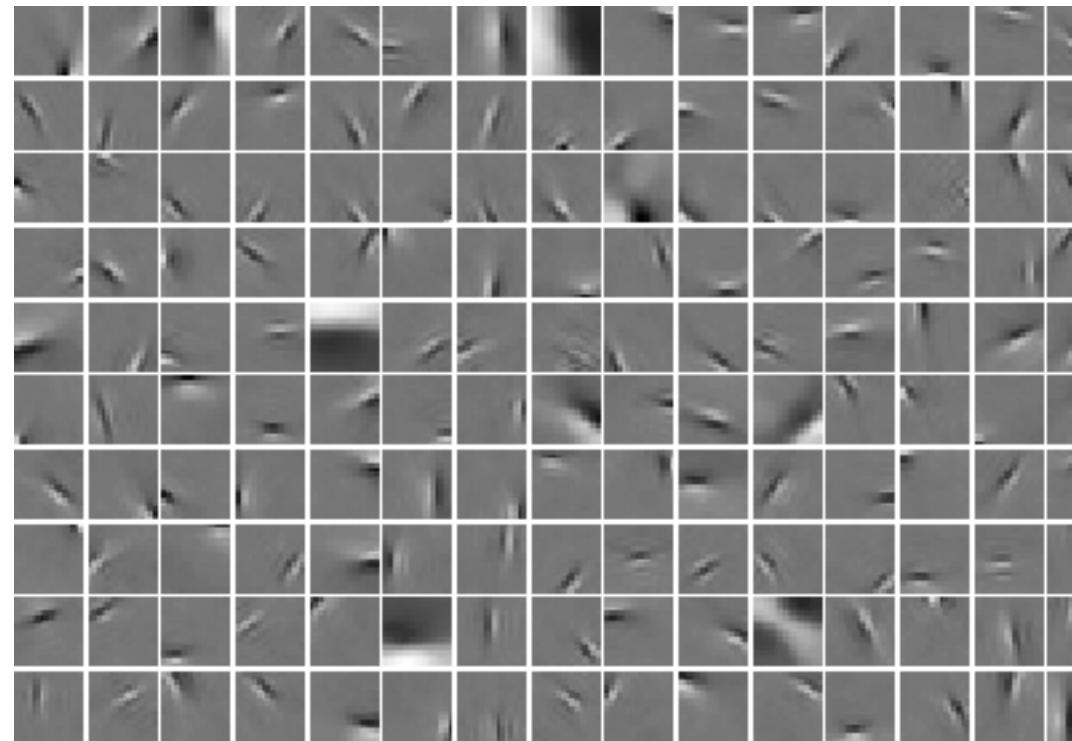
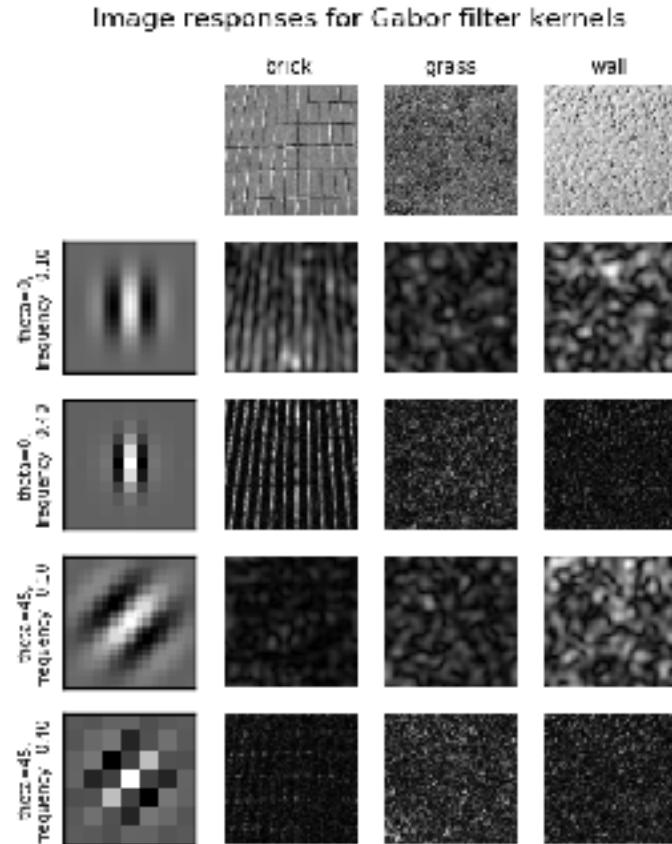


CNN Overview: Self Test

- First layer(s):
 - convolution
 - nonlinearity
 - pooling
 - Each pooling layer *can* make the input image “smaller”
 - allows for “Information Distillation”
 - less dependence on exact pixels
- Final layers are densely connected
 - typically multi-layer perceptrons
- Where are unstable gradients **most** problematic?
 - (A) During Convolution Layer(s) updates
 - (B) During Fully Connected Layer(s) updates
 - (C) Both A and B
 - (D) They are not a problem

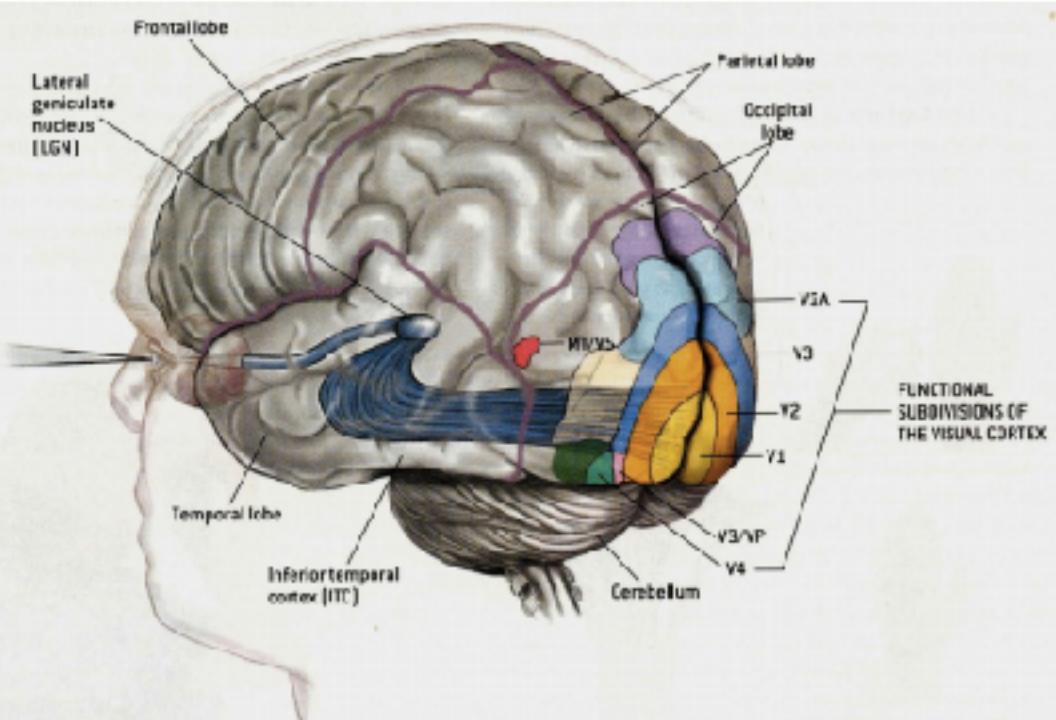
CNN Filtering

- Why perform lots of filtering?
 - recall gabor filtering?



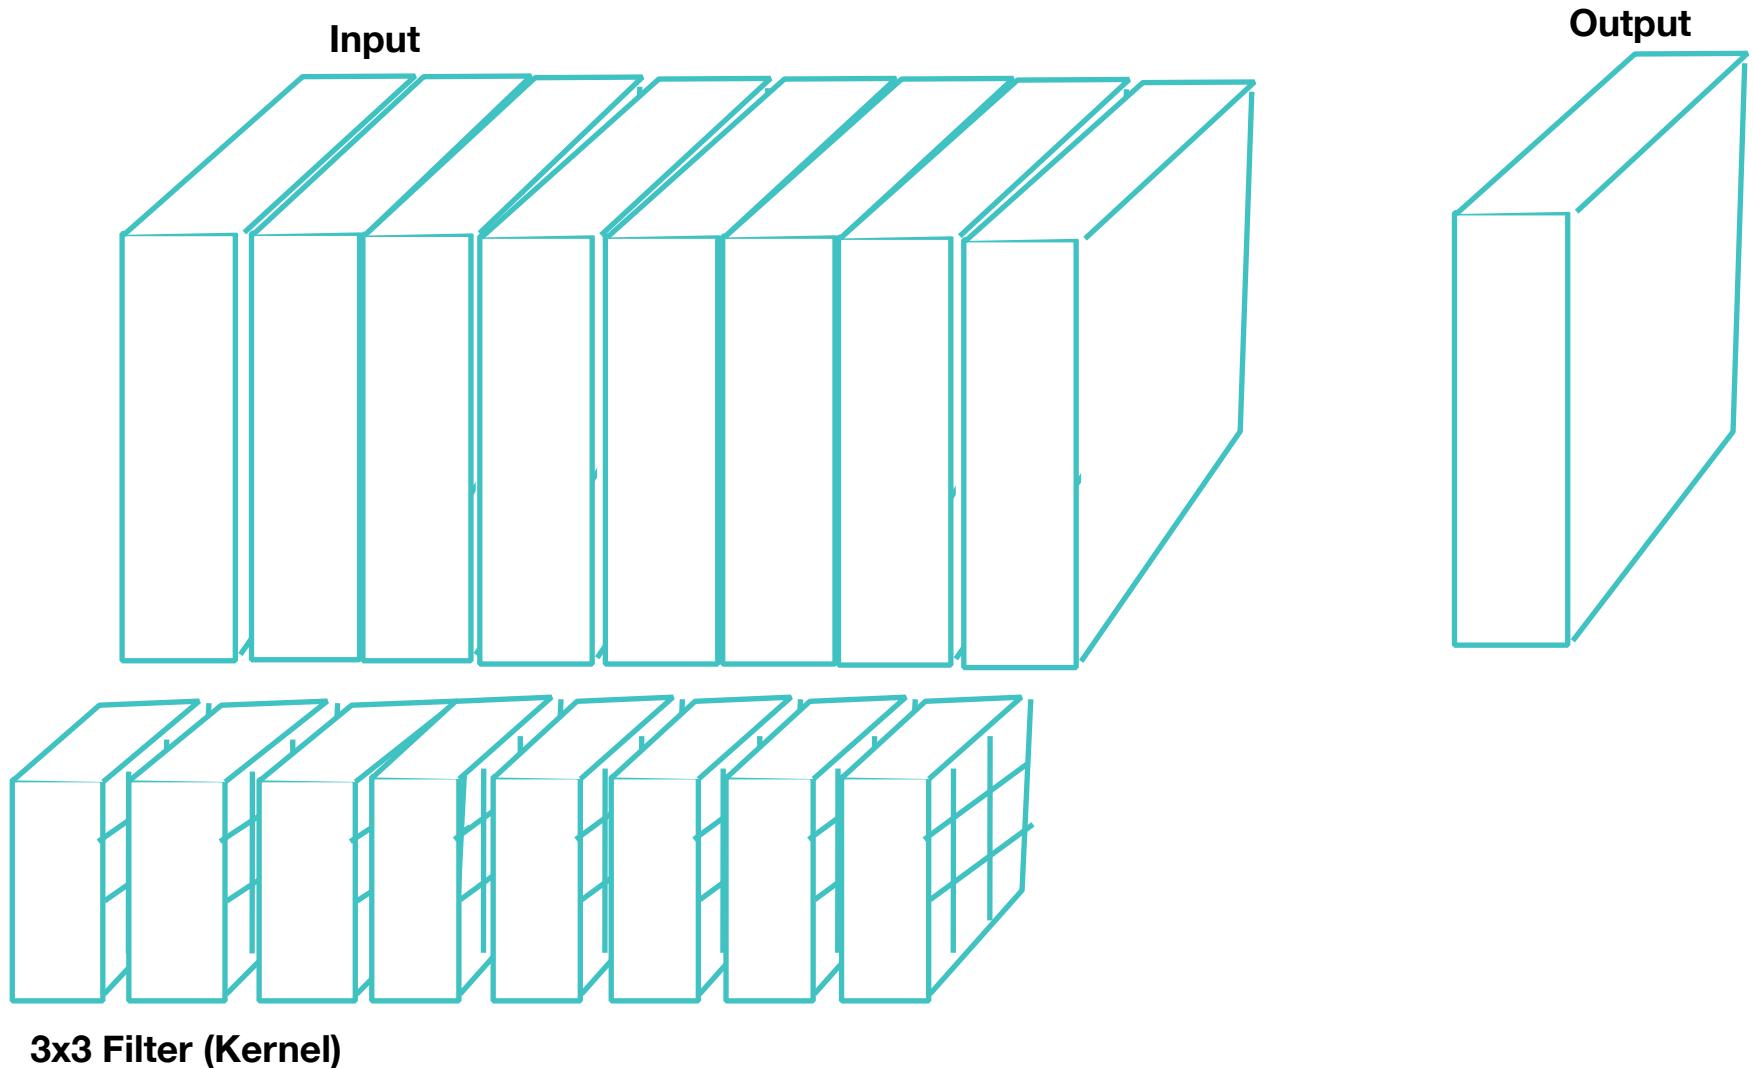
CNN Filtering

- Why perform lots of filtering?
 - recall gabor filtering?



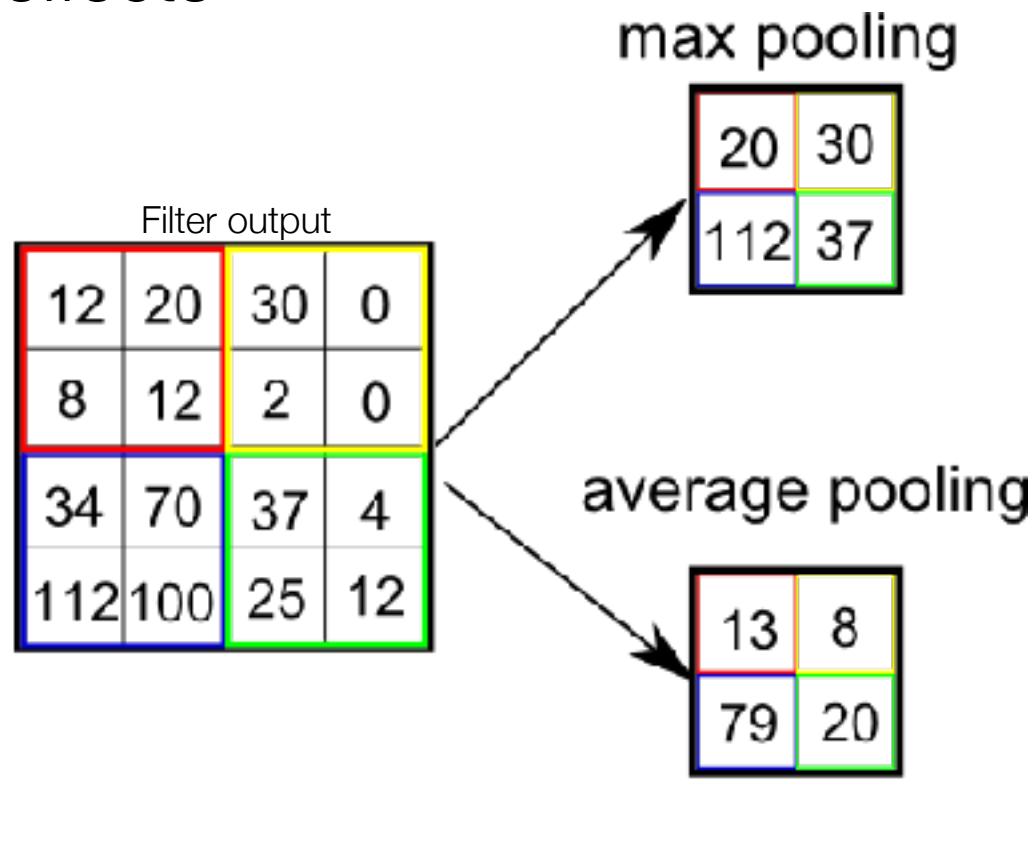
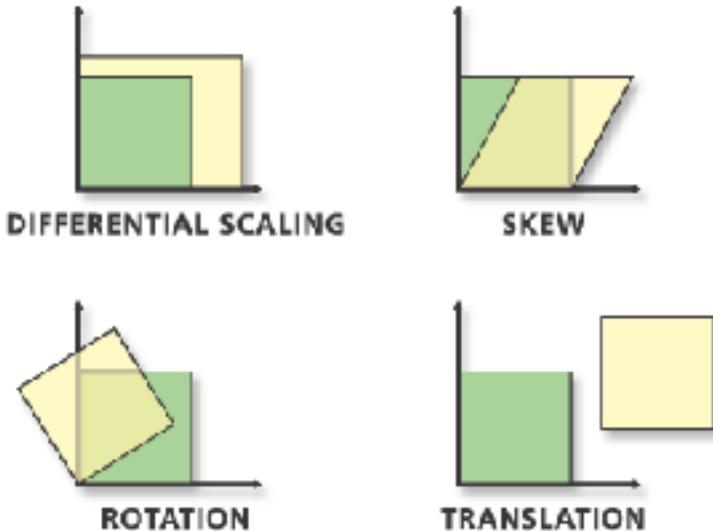
V1	Motion
V2	Stereo
V3	Color
V3a	Texture segregation
V3b	Segmentation, grouping
V4	Recognition
V7	Face recognition
MT	Attention
MST	Working memory/mental imagery

Convolution in a CNN

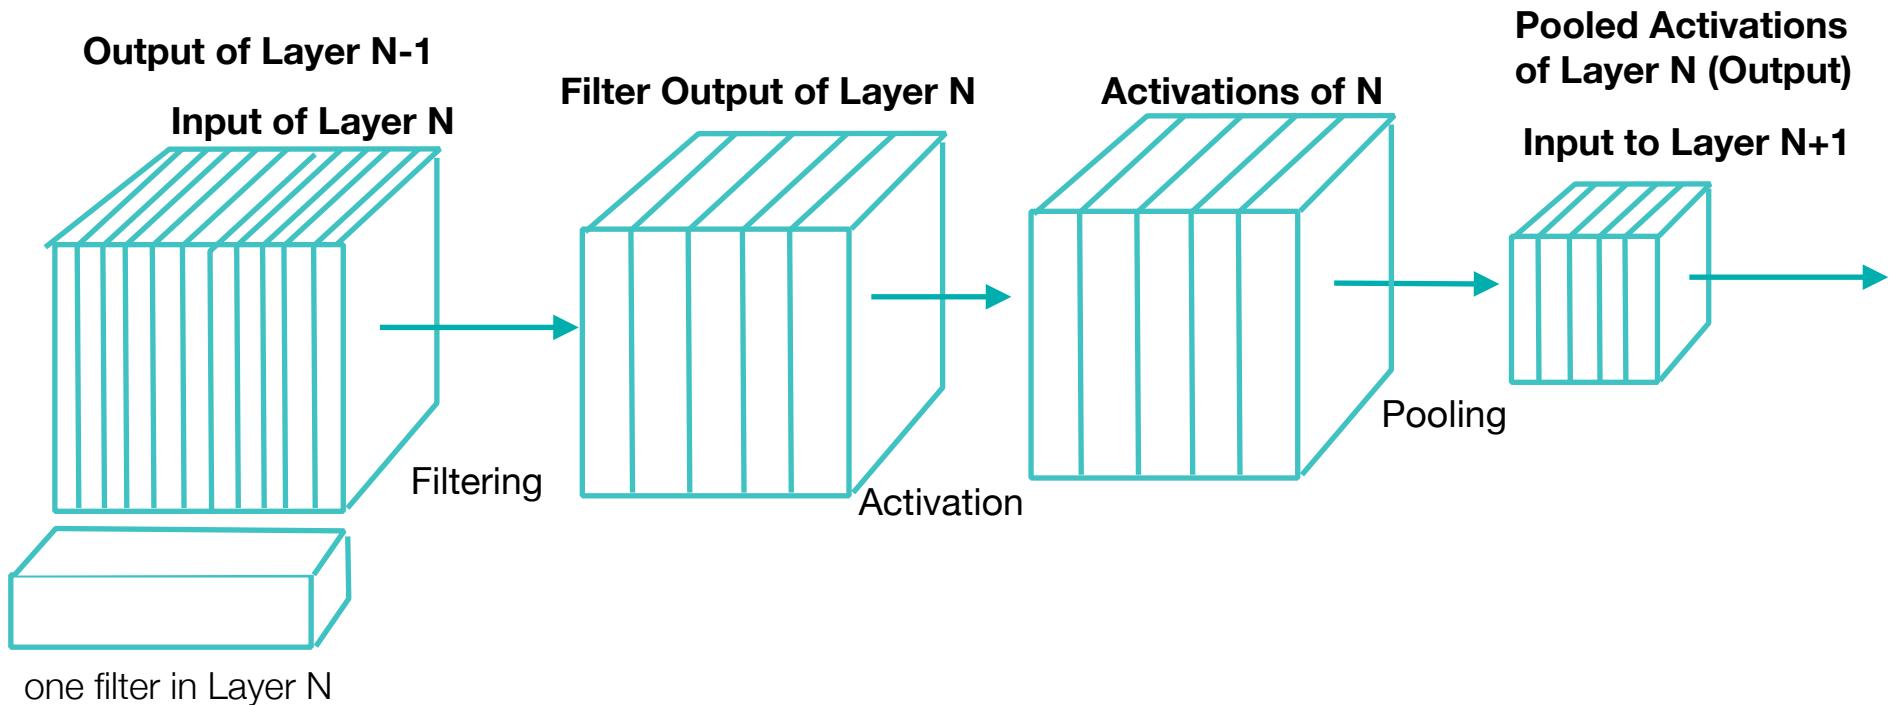


CNN Pooling

- Why perform pooling?
- Why max pooling?
 - reduce translation effects
 - param reduction

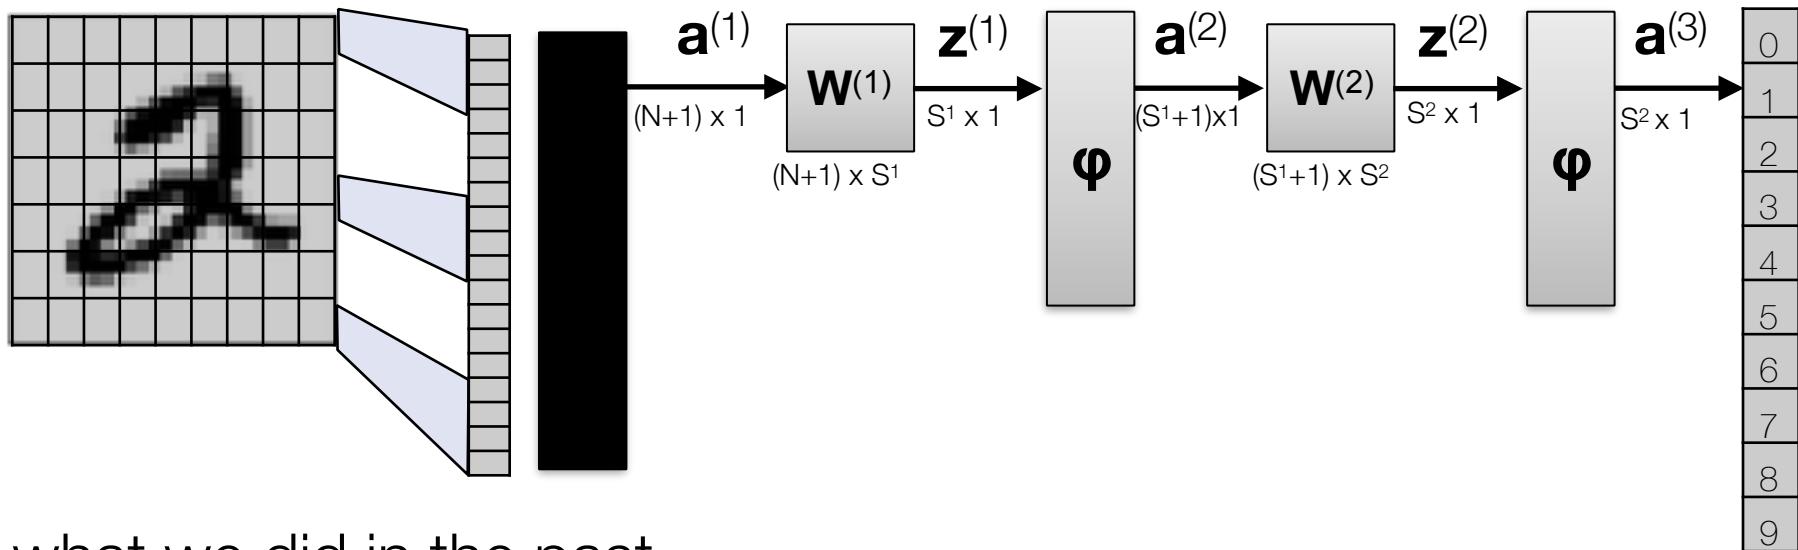


CNNs: Putting it together



Structure of Each Tensor: Channels x Rows x Columns

Simple Example: From Fully Connected to CNN



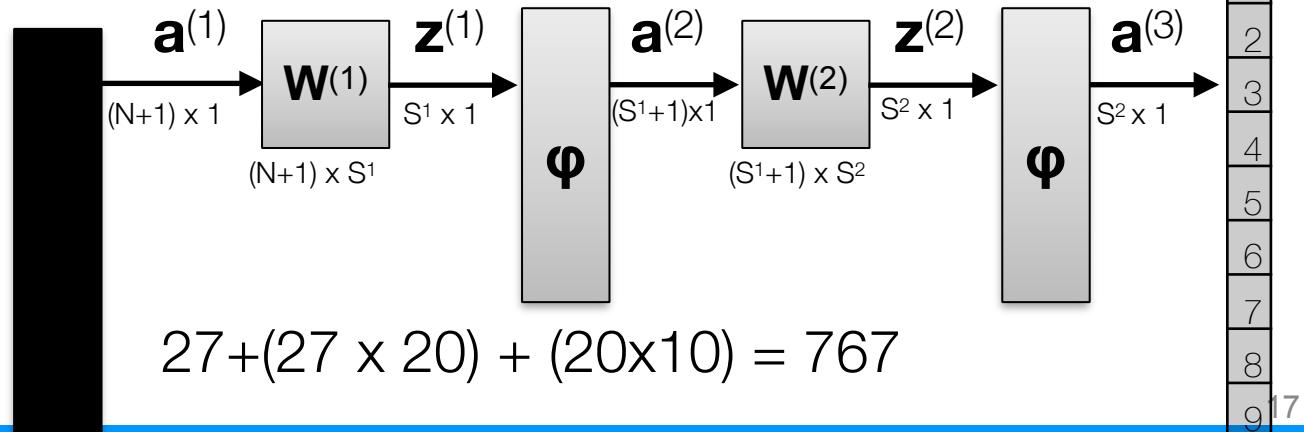
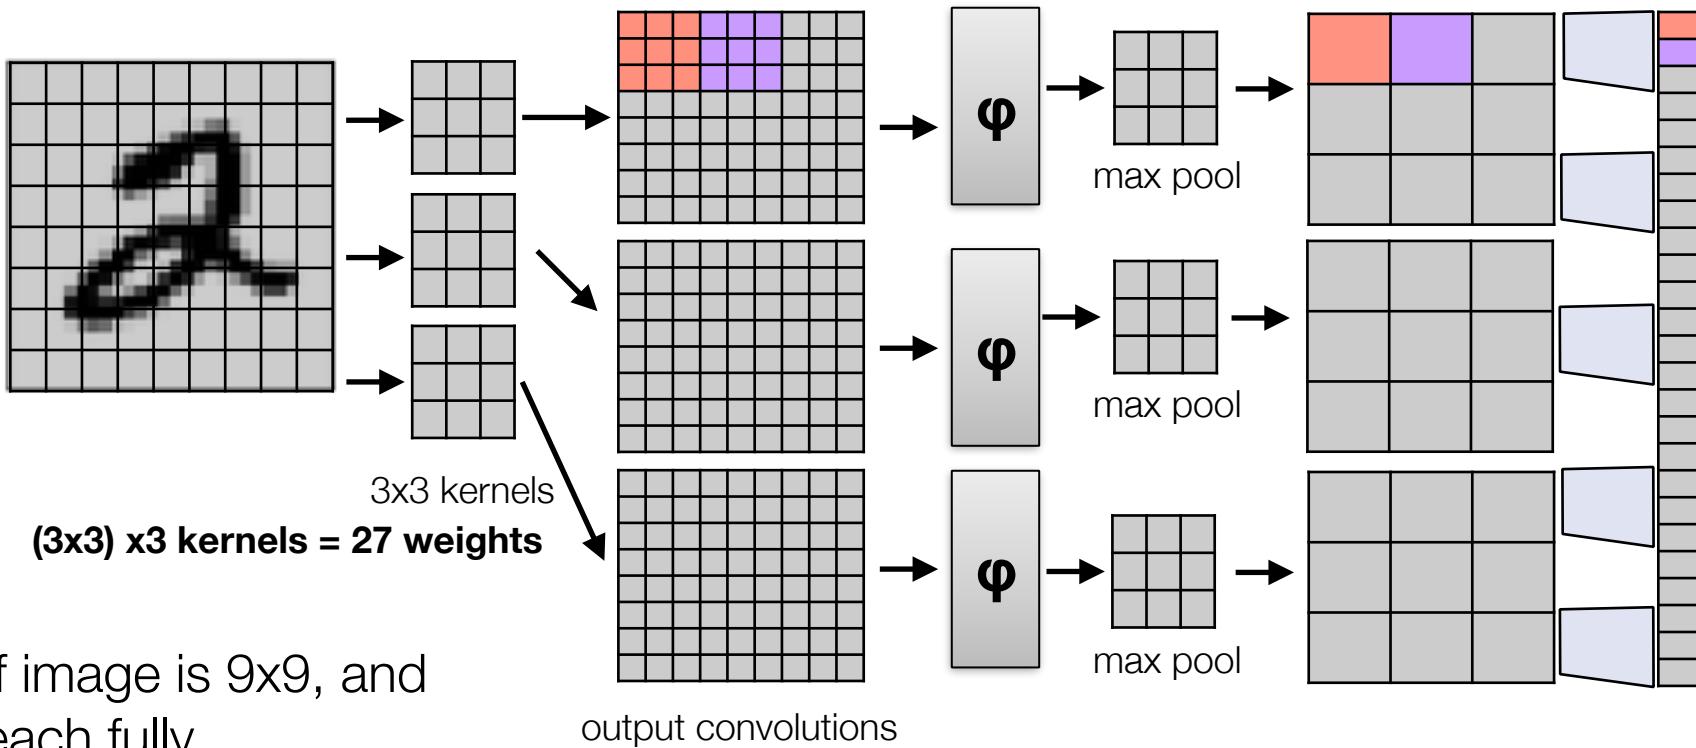
what we did in the past

If image is 9×9 , and each fully connected layer is 20 hidden neurons wide, how many parameters are in this NN (ignore bias)?

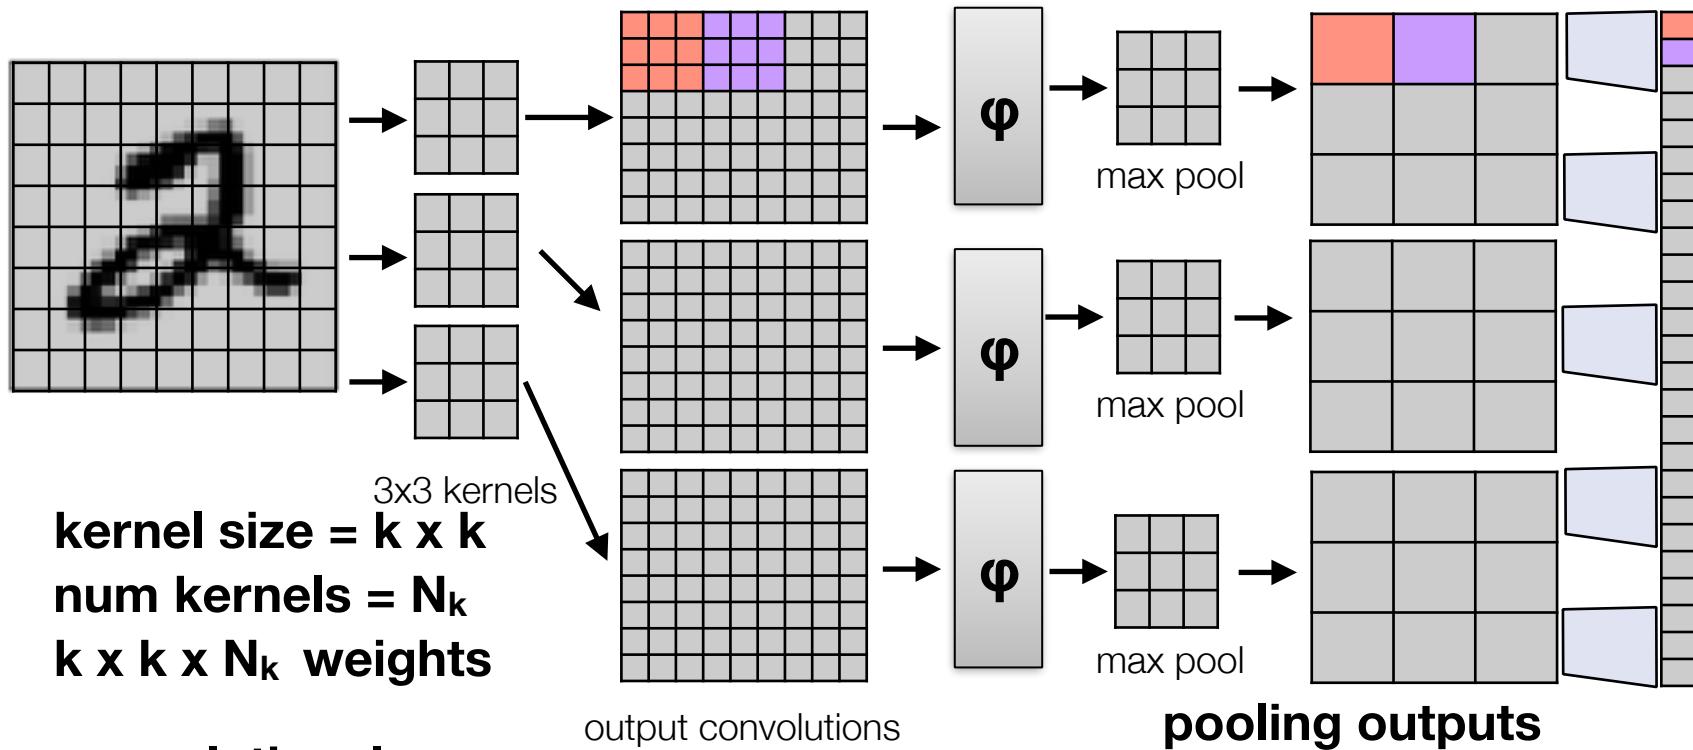
$$(K^2 \times 20) + (20 \times 10) = 200 + 20 K^2$$

$$\text{for } 9 \times 9 = 200 + 20 \times 9^2 = 1,820 \text{ parameters}$$

Simple Example: From Fully Connected to CNN



Simple Example: From Fully Connected to CNN

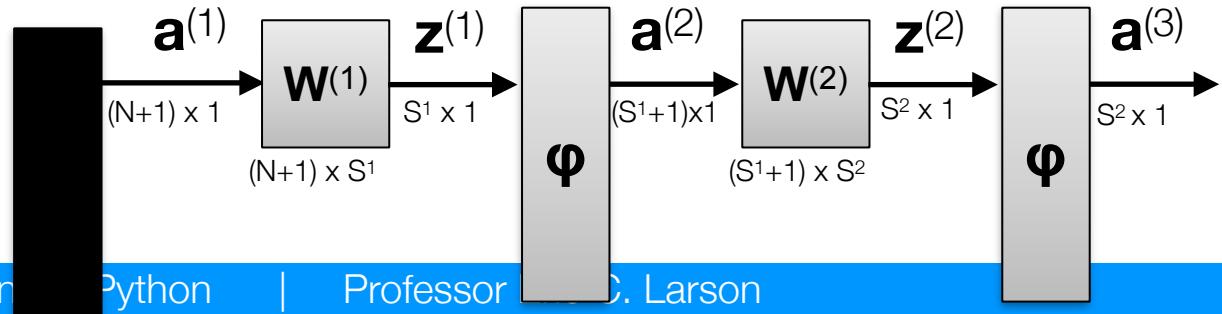


Input to MLP
 $N_k \times (K^2/k^2)$

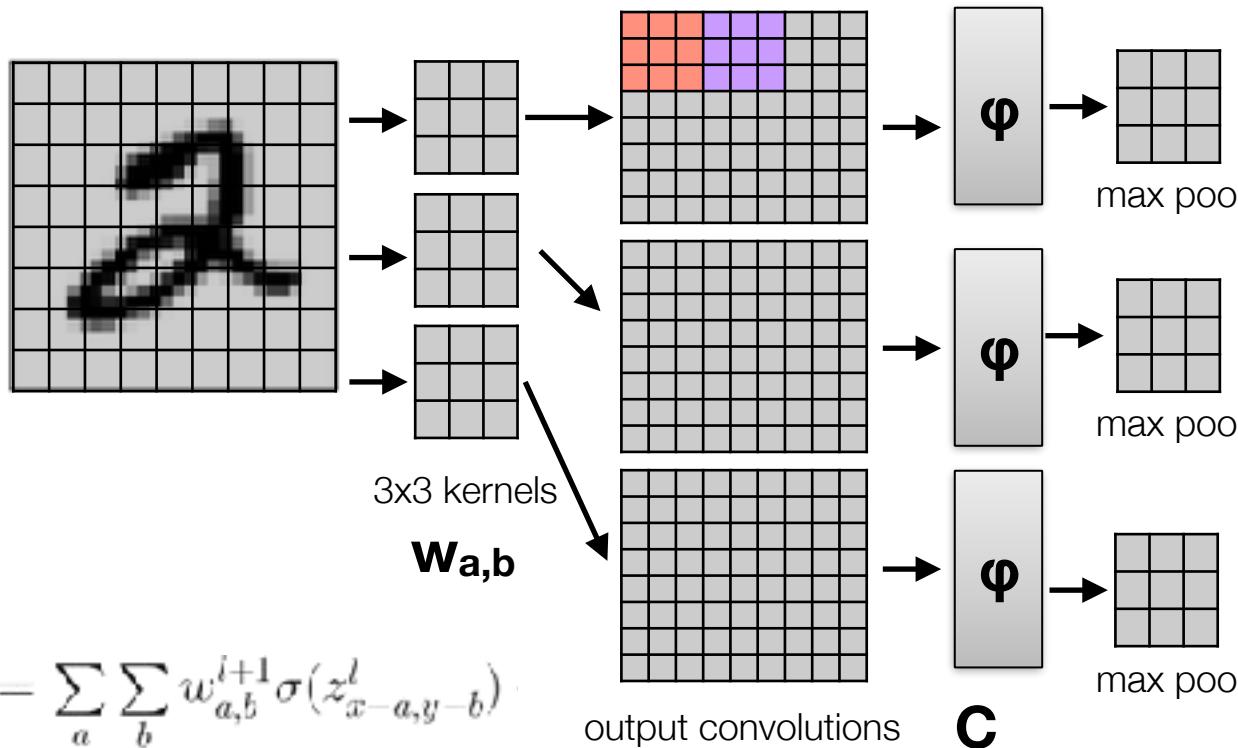
image dimension

filter dimension

num filters



CNN gradient



Derivative of max pool is easy:

for each input x_i

$$f'(x_i) = \begin{cases} 1 & \text{if } x_i \text{ is max} \\ 0 & \text{else} \end{cases}$$

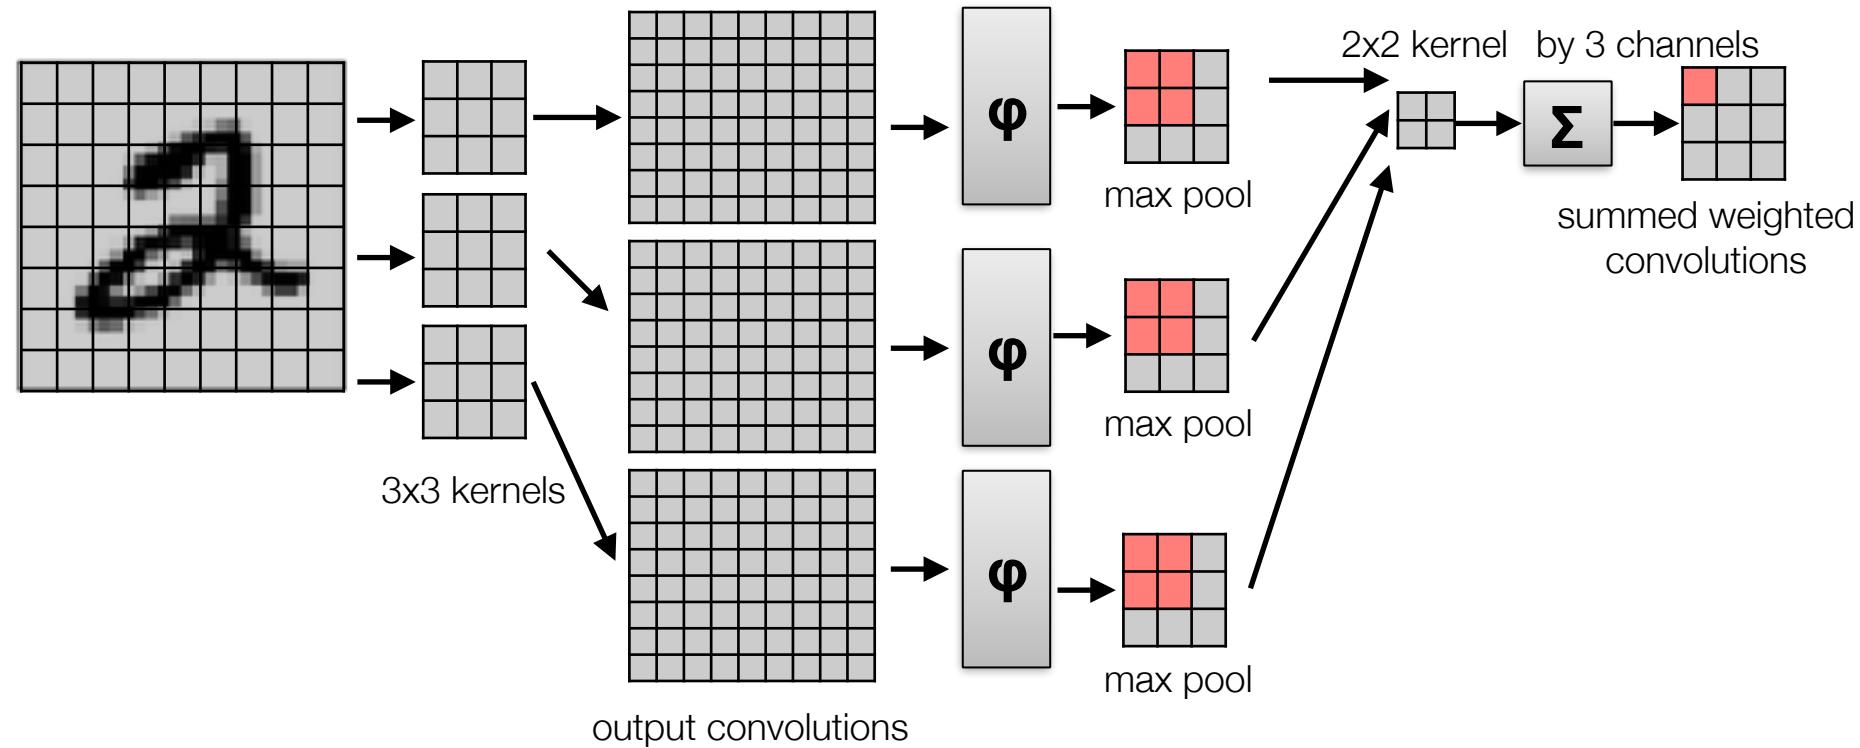
Derivative of convolution is more involved:

$$\frac{\partial C}{\partial w_{a,b}^l} = \sum_x \sum_y \frac{\partial C}{\partial z_{x,y}^l} \frac{\partial z_{x,y}^l}{\partial w_{a,b}^l} = \sum_x \sum_y \delta_{x,y}^l \frac{\partial (\sum_{a'} \sum_{b'} w_{a',b'}^l \sigma(z_{x-a',y-b'}^l) + b_{x,y}^l)}{\partial w_{a,b}^l} =$$
$$\sum_x \sum_y \delta_{x,y}^l \sigma(z_{x-a,y-b}^{l-1}) = \delta_{a,b}^l * \sigma(z_{-a,-b}^{l-1}) = \delta_{a,b}^l * \sigma(ROT180(z_{a,b}^{l-1}))$$

CNN gradient

- But we really want to understand the process!
- These are great guides:
 - [https://grzegorzgwardys.wordpress.com/
2016/04/22/8/](https://grzegorzgwardys.wordpress.com/2016/04/22/8/)
 - [http://andrew.gibiansky.com/blog/machine-
learning/convolutional-neural-networks/](http://andrew.gibiansky.com/blog/machine-learning/convolutional-neural-networks/)

CNN adding more convolutional layers



Input layer

(S1) 4 feature maps

(C1) 4 feature maps

(S2) 6 feature maps

(C2) 6 feature maps



convolution layer

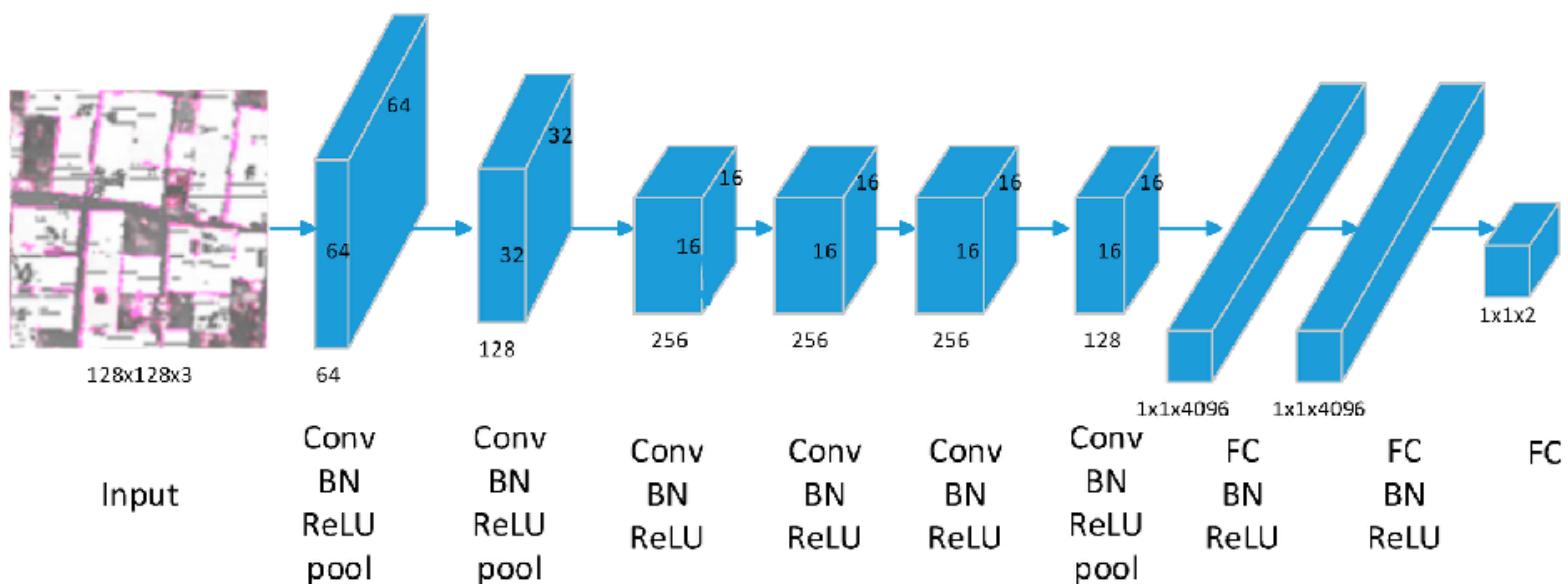
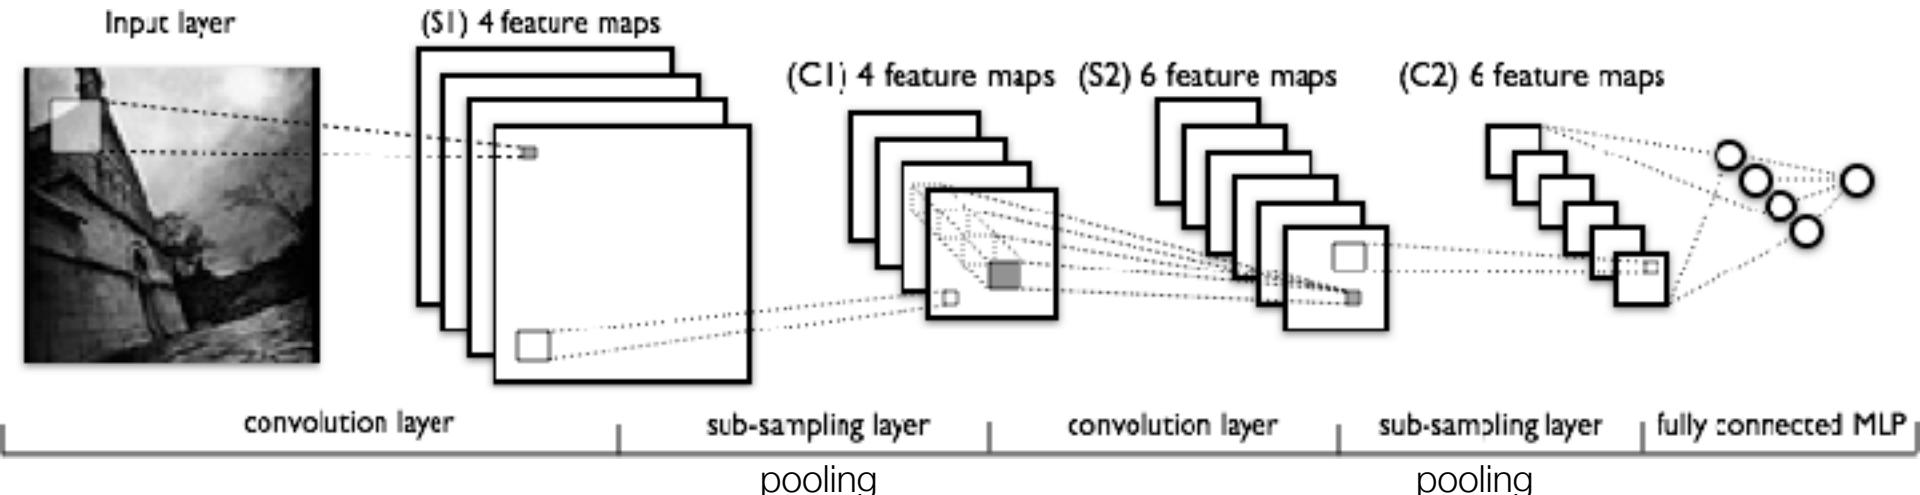
sub-sampling layer

convolution layer

sub-sampling layer

fully connected MLP

Some Example CNN Architectures



CNN: What does it all mean?

Deep Visualization Toolbox

yosinski.com/deepvis

#deepvis



Jason Yosinski



Jeff Clune



Anh Nguyen



Thomas Fuchs



Hod Lipson



Cornell University



Jet Propulsion Laboratory
California Institute of Technology

<https://github.com/yosinski/deep-visualization-toolbox>

TensorFlow and Basic CNNs

Demo

Convolutional Neural Networks
in TensorFlow
with Keras



11. Convolutional Neural Networks.ipynb

Next Lecture

- More CNN architectures and CNN history

Lecture Notes for **Machine Learning in Python**

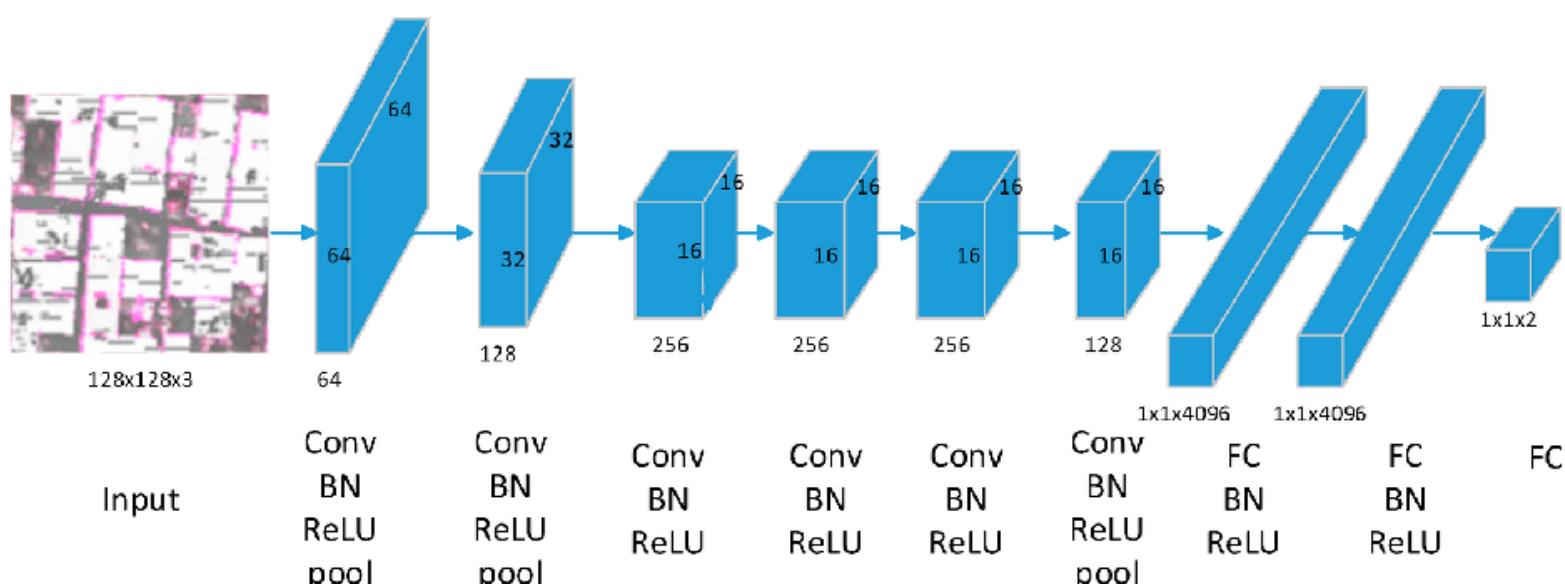
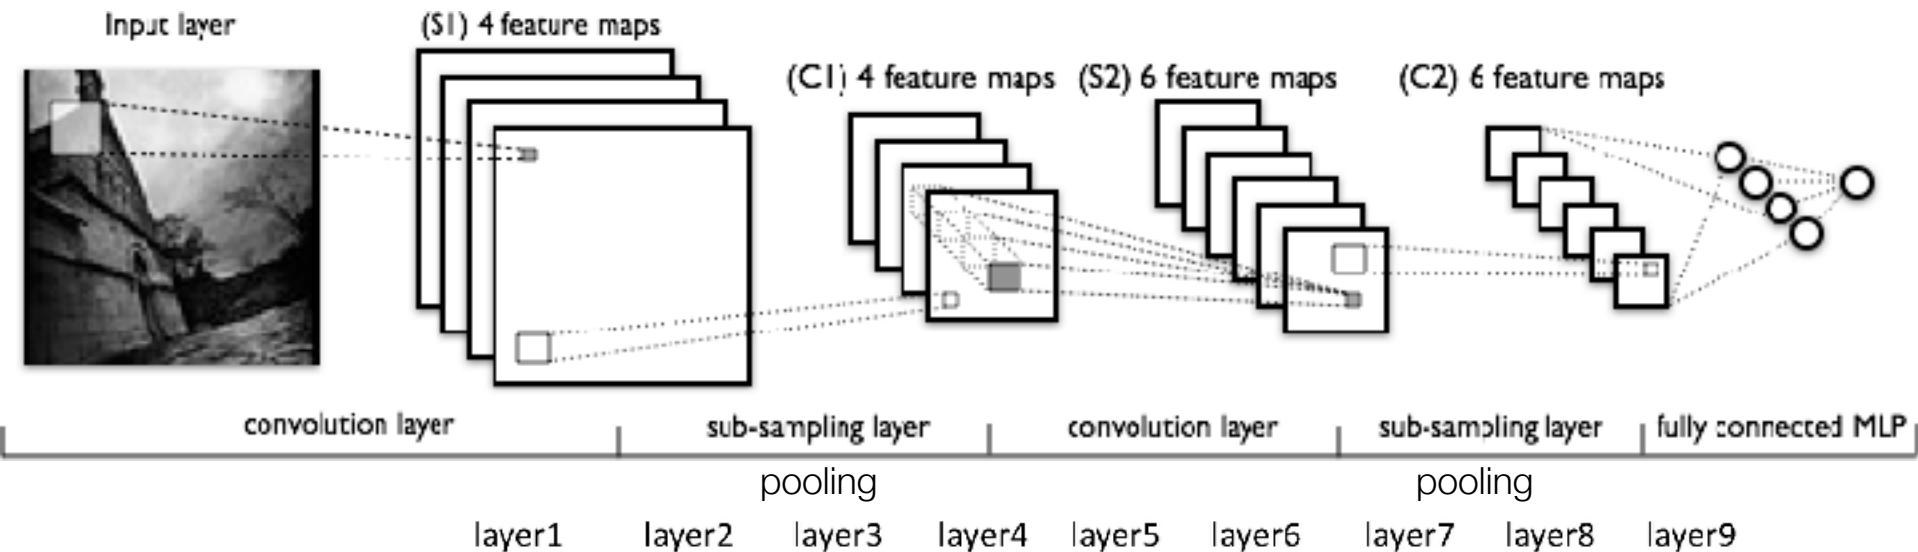
Professor Eric Larson

An Ongoing History of Convolutional Networks

Class logistics and Agenda

- Wide/Deep Lab due!
- Agenda:
 - Finish CNN Demo
 - History of CNNs
 - with Modern CNN Architectures
- Next Time:
 - More Advanced CNN Demo

Last Time:

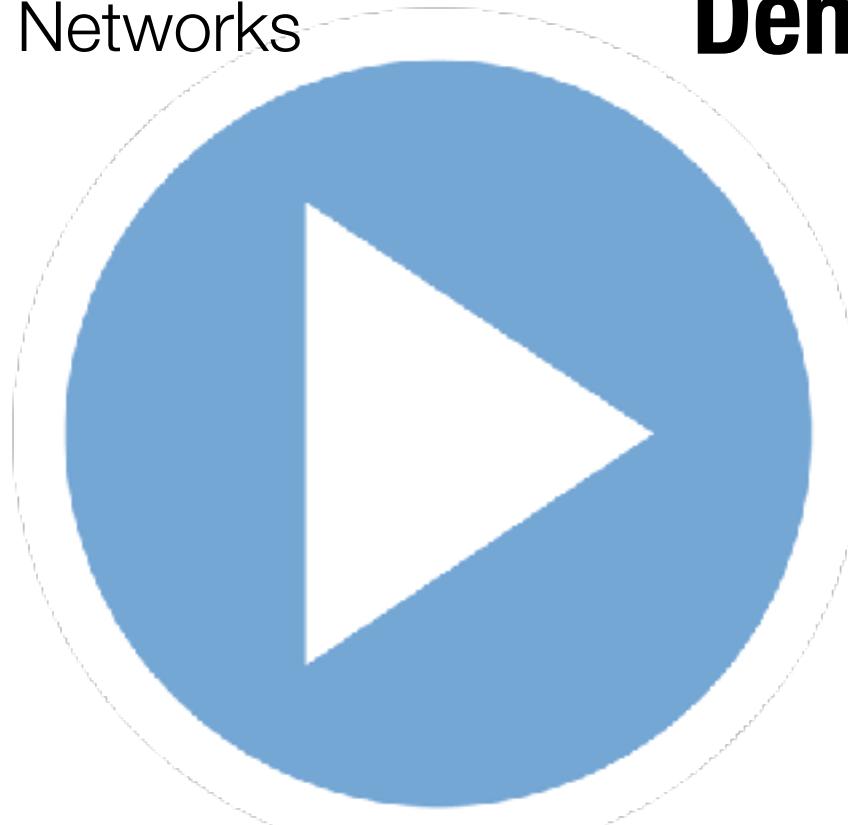


TensorFlow and Basic CNNs

Convolutional Neural Networks
in TensorFlow
with Keras

with Sequential API!

Finish
Demo



11. Convolutional Neural Networks.ipynb

History of Convolutional Neural Networks



Thanks to machine-learning algorithms,
the robot apccalypse was short-lived.

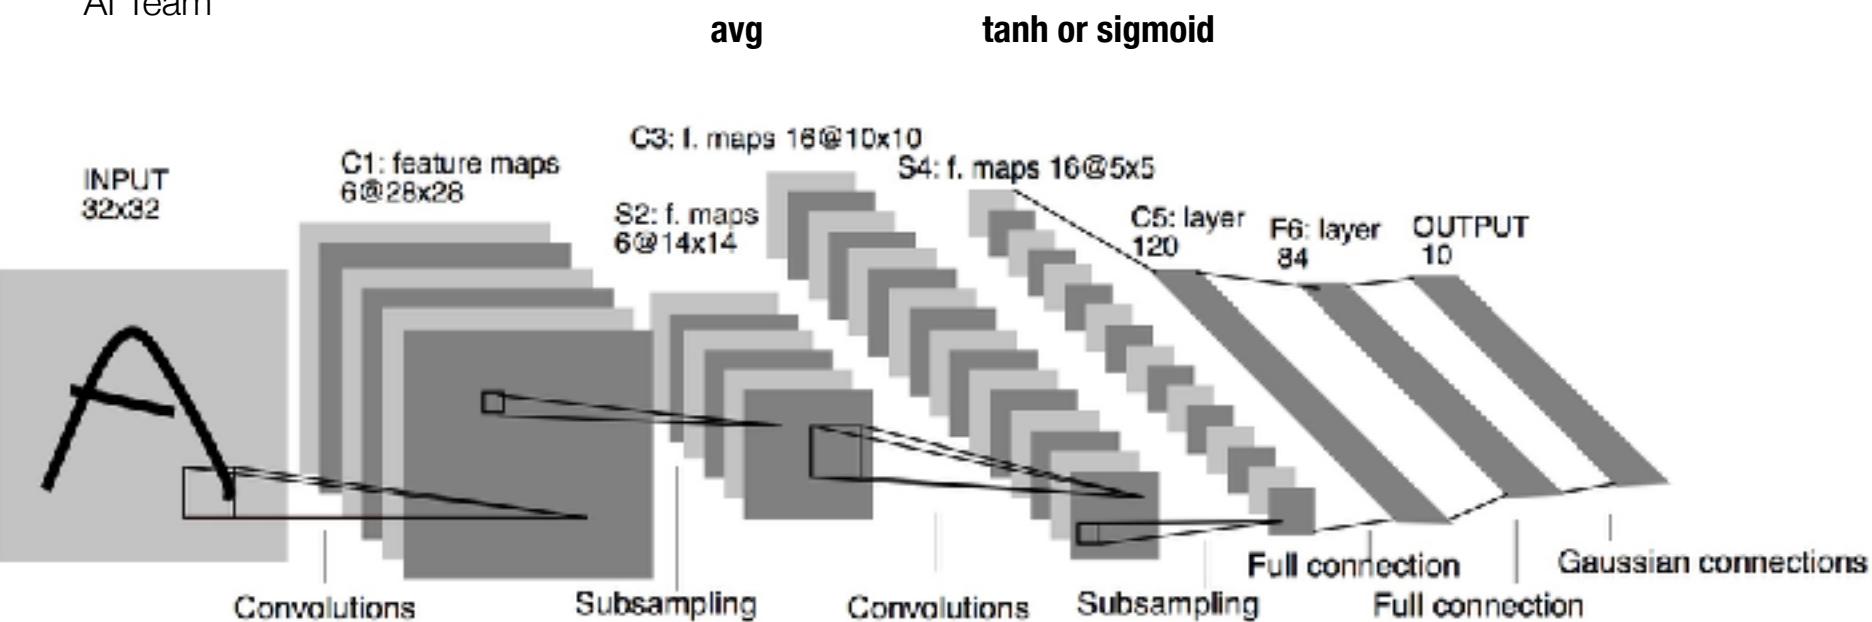
Types of CNN, 1988-1998



Yann
LeCun

Heads Facebook
AI Team

- **LeNet-1** (1988)
 - ~2600 params, not many layers
- **LeNet-5** (1998)
 - 7 layers, gets excellent MNIST performance
- Major contribution, general structure:
 - conv=>pool=>non-linearity=> ...=>MLP

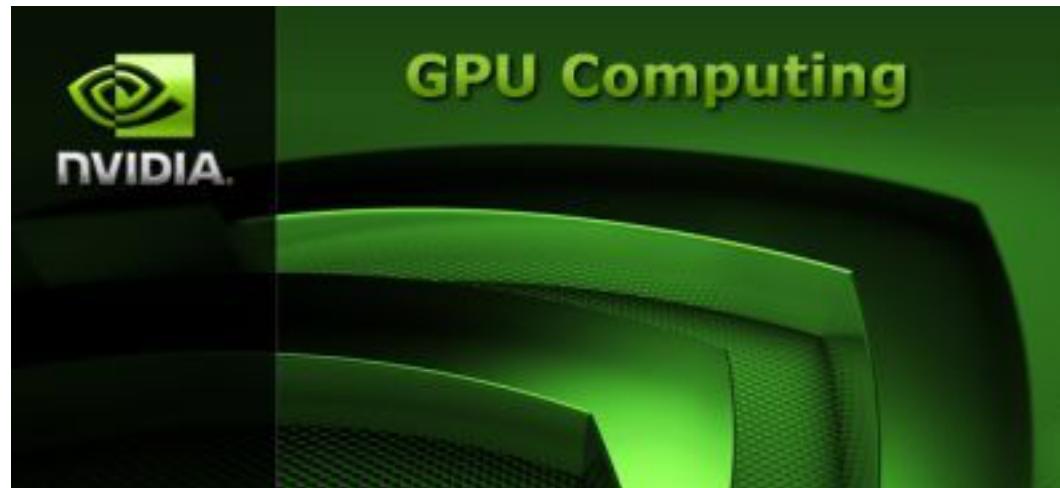


CNN History

- List of major breakthroughs from 1998 through 2010 in convolutional networks:



- 2010



Types of CNN, 2010



Dan
Ciresan

AI Researcher
IDSIA, Switzerland

- **Ciresan Net**
- Publishes code for running CNN via GPU
 - Subsequently wins 5 international competitions
 - from stop signs => cancer detection
- Major contribution: NVIDIA parallelized training algorithms

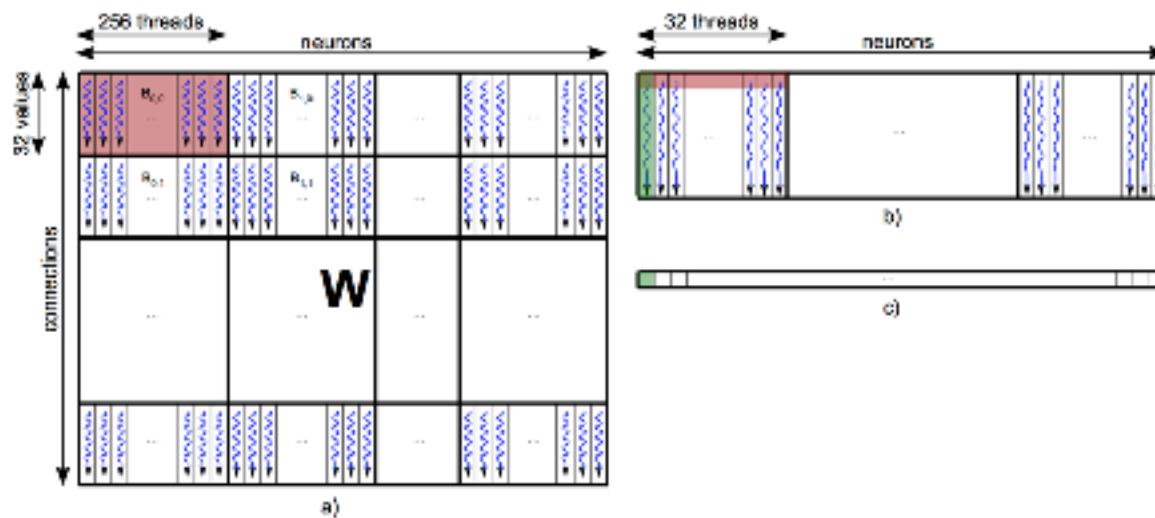
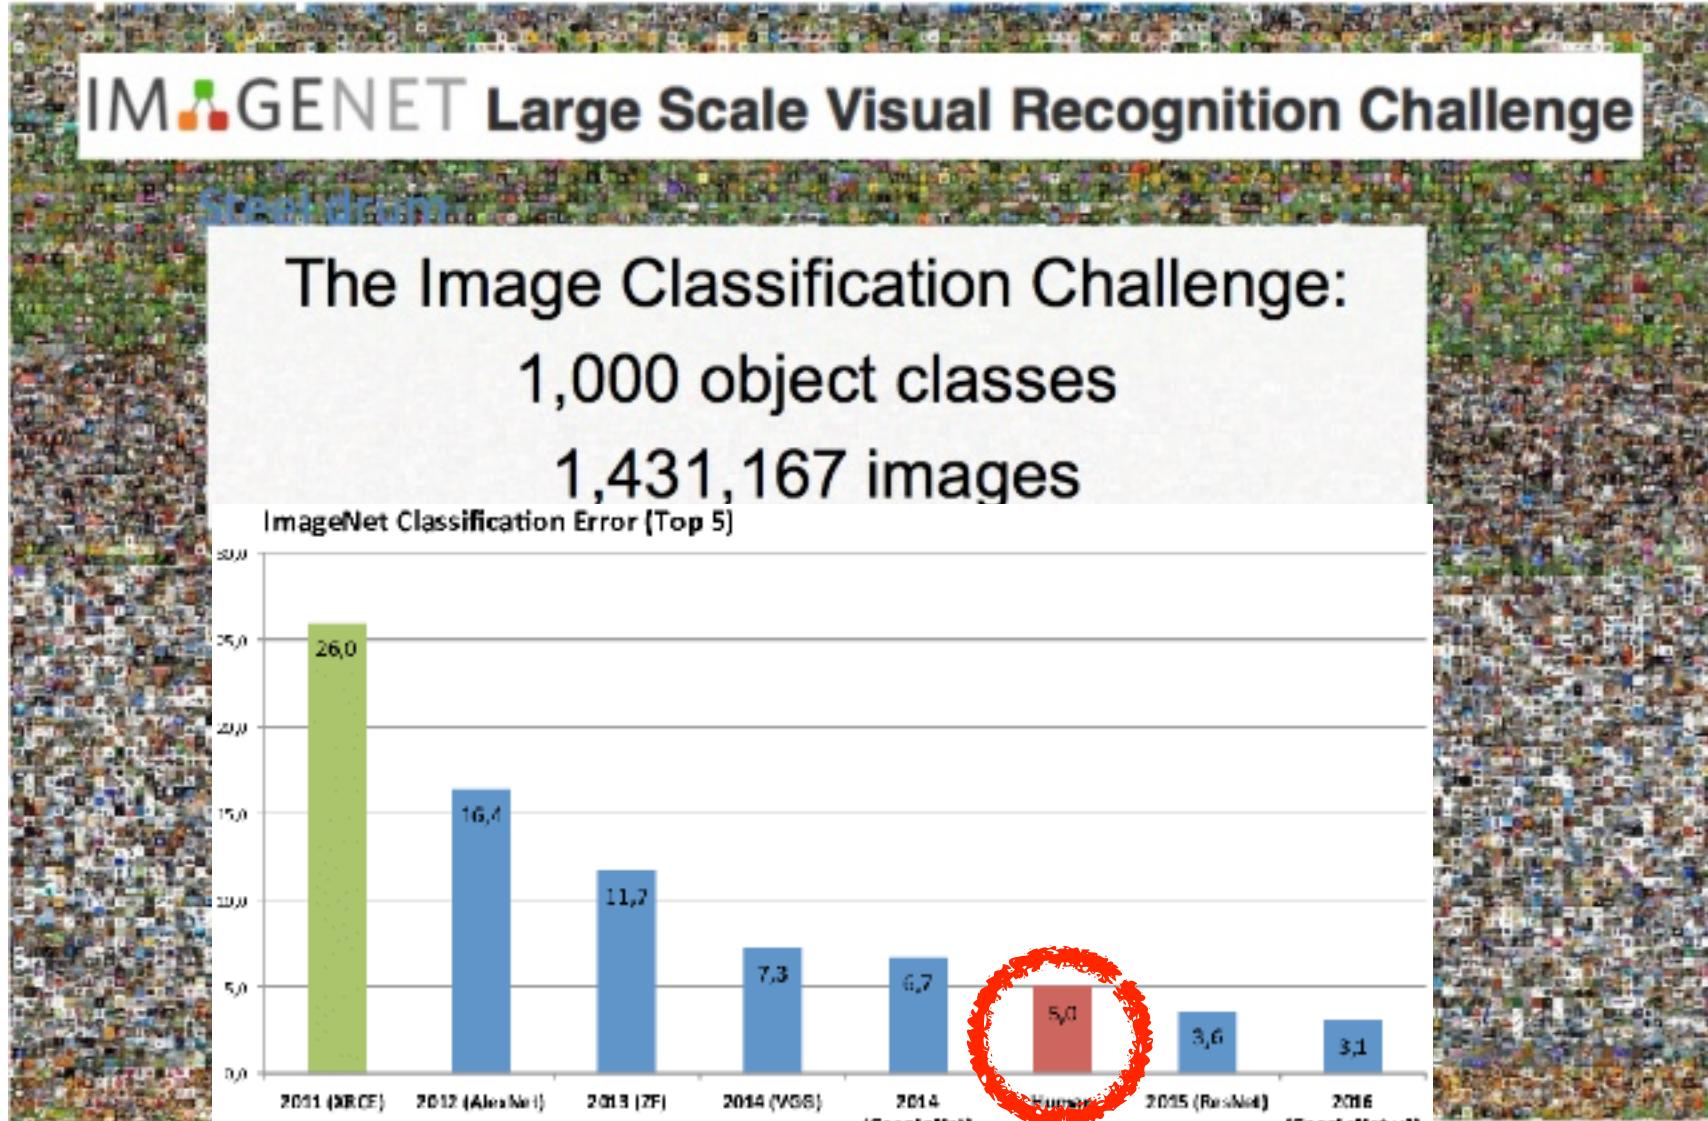


Figure 2: Forward propagation: a) mapping of kernel 1 grid onto the padded weight matrix; b) mapping the kernel 2 grid onto the partial dot products matrix; c) output of forward propagation.

<https://arxiv.org/pdf/1003.0358.pdf>

ImageNet Competition (2010-2016)



https://www.researchgate.net/figure/Winner-results-of-the-ImageNet-large-scale-visual-recognition-challenge-LSVRC-of-the_fig7_324476862

<https://www.slideshare.net/nmhkahn/case-study-of-convolutional-neural-network-61556303>

Types of CNN, 2012

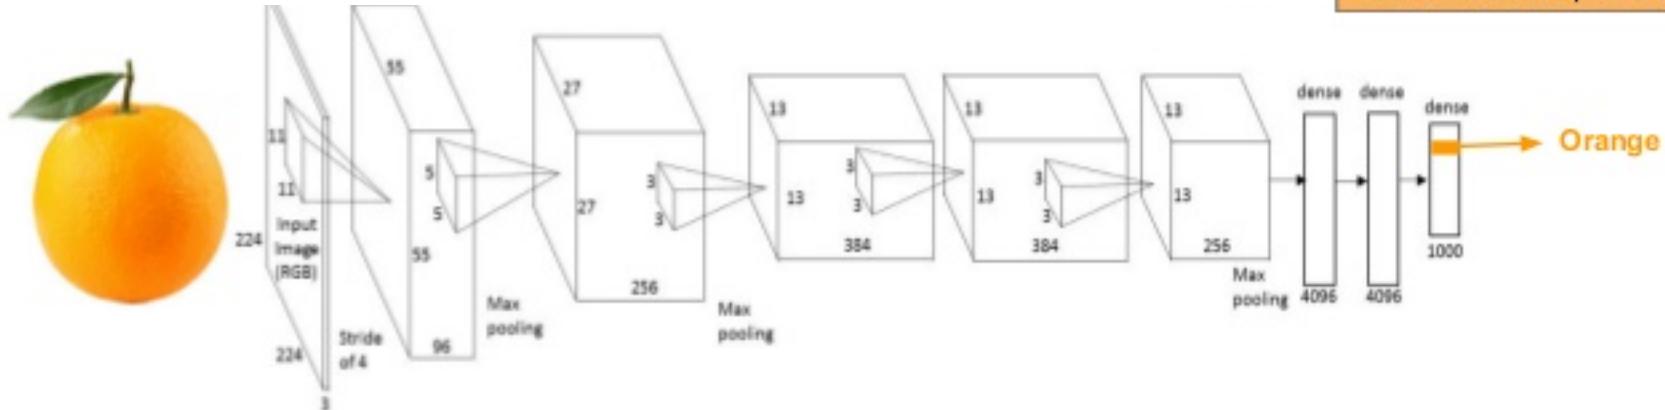
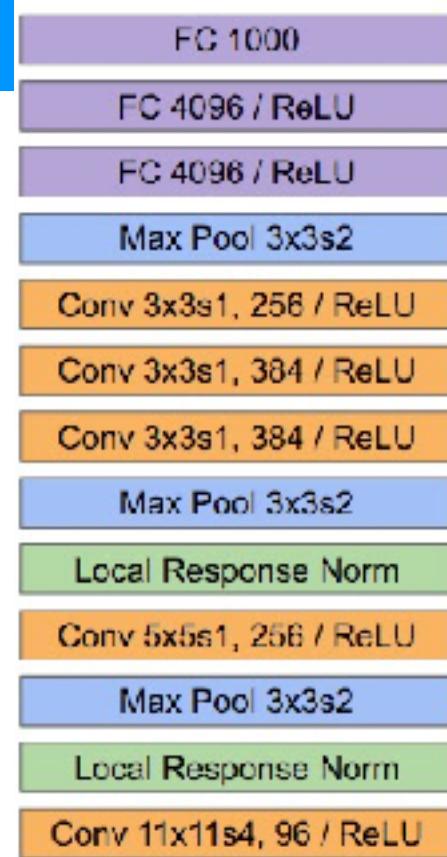
AlexNet



Alex
Krizhevsky

Google

- **AlexNet**, Hinton is mentor
 - wins ImageNet competition
- Major contributions:
 - dropout for regularization
 - systematic use of ReLU
 - data expansion
 - overlapping max pool



Warning



Types of CNN, 2013



Karen Simonyan
and
Andrew Zisserman



- Oxford **VGG Net** (Visual Geometry Group)
- Major contributions:
 - small cascaded kernels
 - way more layers (19 versus ~7)
 - “emulates” biology “better”
 - trained on NVIDIA GPUs for 2-3 weeks

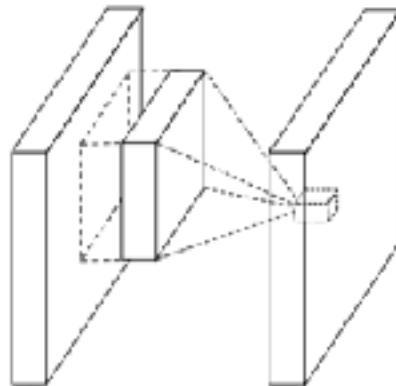
<http://www.robots.ox.ac.uk/~vgg/practicals/cnn/>

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64	conv3-64	conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128	conv3-128	conv3-128
maxpool					
conv3-256	conv3-256	conv3-256	conv3-256	conv3-256	conv3-256
conv3-256	conv3-256	conv3-256	conv3-256	conv3-256	conv3-256
maxpool					
conv3-512	conv3-512	conv3-512	conv3-512	conv3-512	conv3-512
conv3-512	conv3-512	conv3-512	conv3-512	conv3-512	conv3-512
maxpool					
conv3-512	conv3-512	conv3-512	conv3-512	conv3-512	conv3-512
conv3-512	conv3-512	conv3-512	conv3-512	conv3-512	conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Table 2: Number of parameters (in millions).

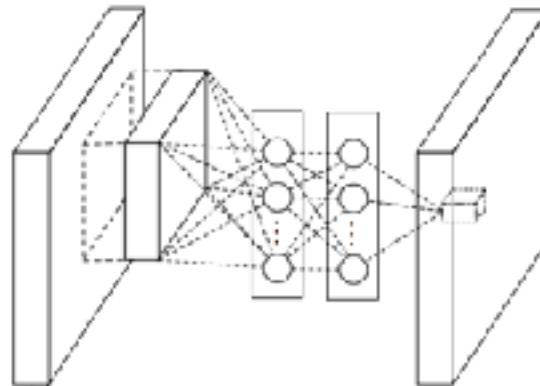
Network	A,A-LRN	B	C	D	E
Number of parameters	133	133	134	138	144

- Network in Network **NiN**
 - or MLPConv

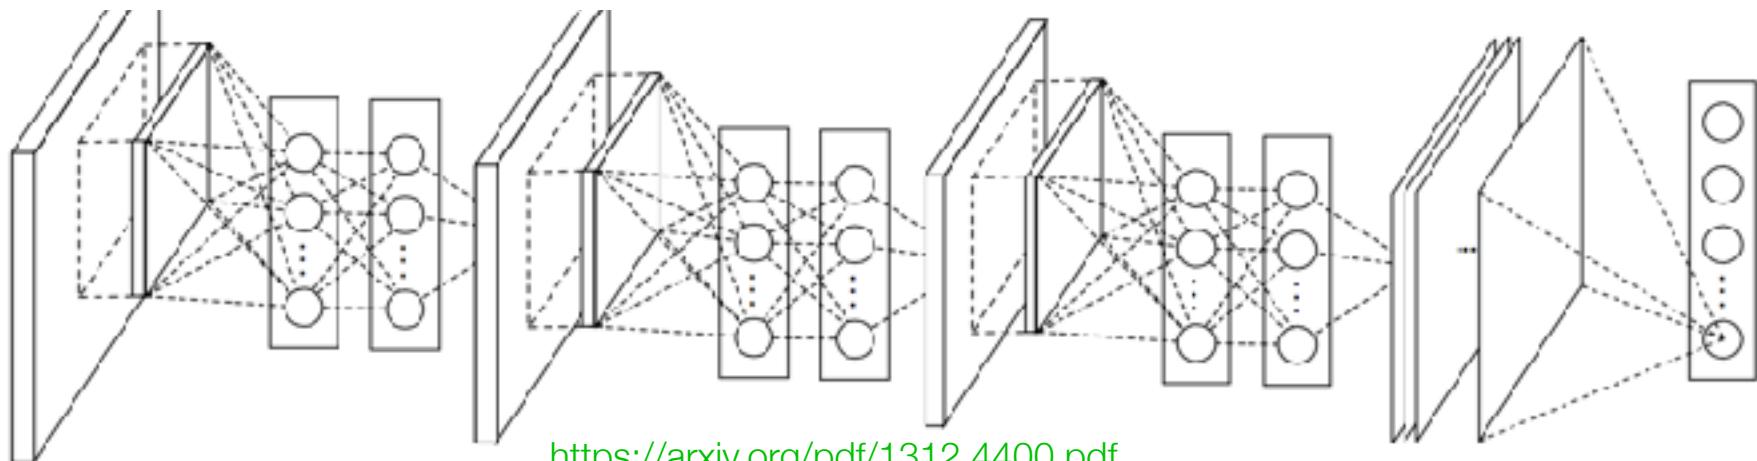


(a) Linear convolution layer

Min Lin^{1,2}, Qiang Chen², Shuicheng Yan²
¹Graduate School for Integrative Sciences and Engineering
²Department of Electronic & Computer Engineering
National University of Singapore, Singapore
`{linmin,chenqiang,elec yans}@nus.edu.sg`



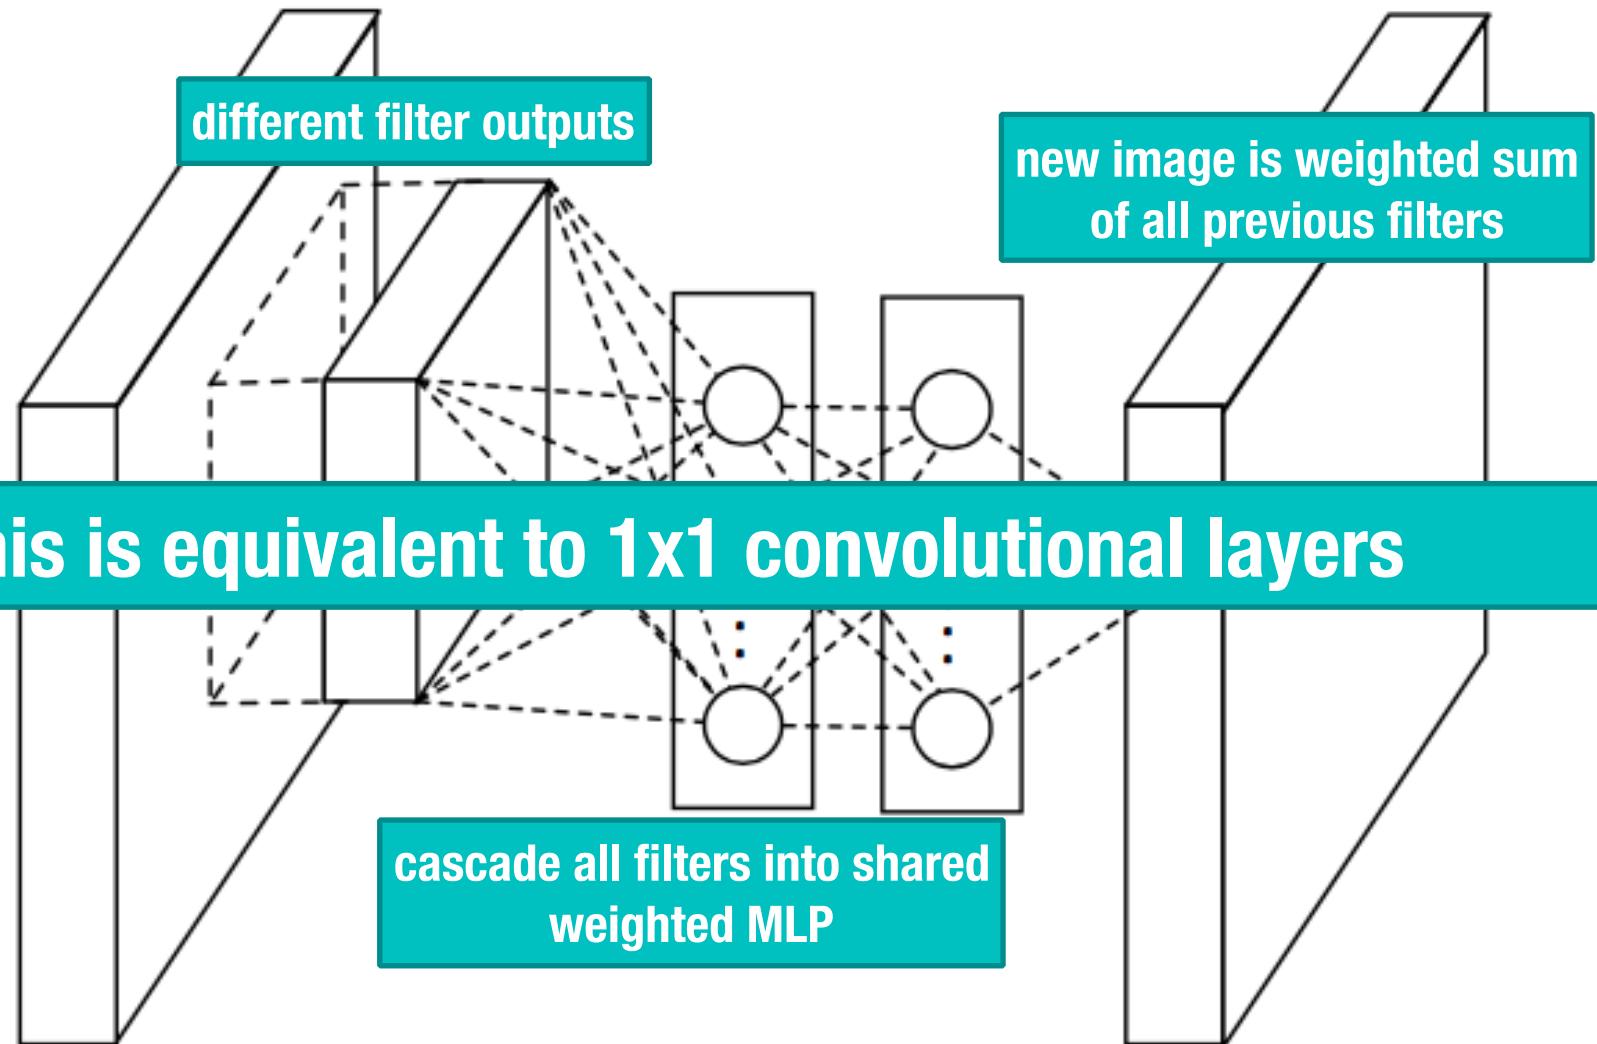
(b) Mlpconv layer



<https://arxiv.org/pdf/1312.4400.pdf>

Types of CNN, 2014

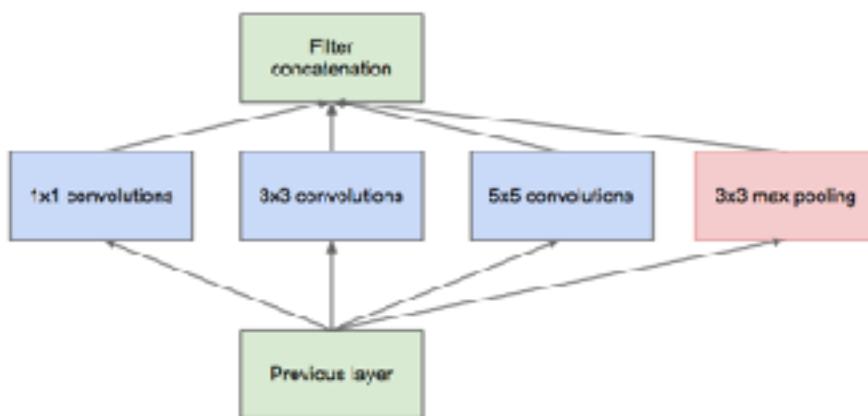
- Network in Network



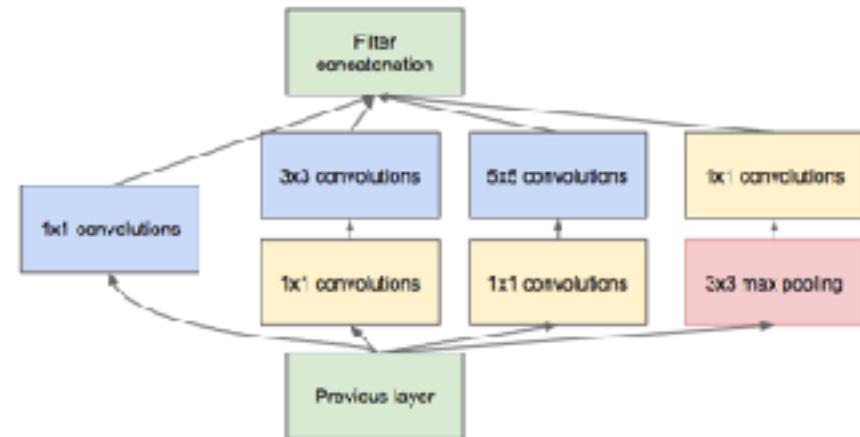
Types of CNN, 2014

Research at Google
Home Publications People Teams Outreach
Christian Szegedy
Research Area(s)
Machine Intelligence
Machine Perception
Photo of Christian Szegedy

- GoogLeNet
 - or Inception V1
- Major contribution:
 - bottleneck layering
 - parallel NiN



(a) Inception module, naïve version



(b) Inception module with dimension reductions

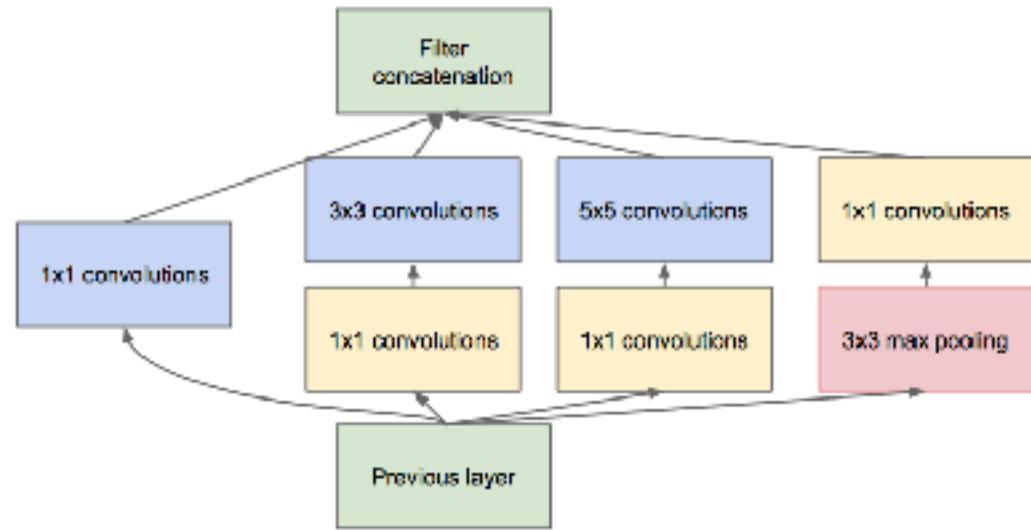
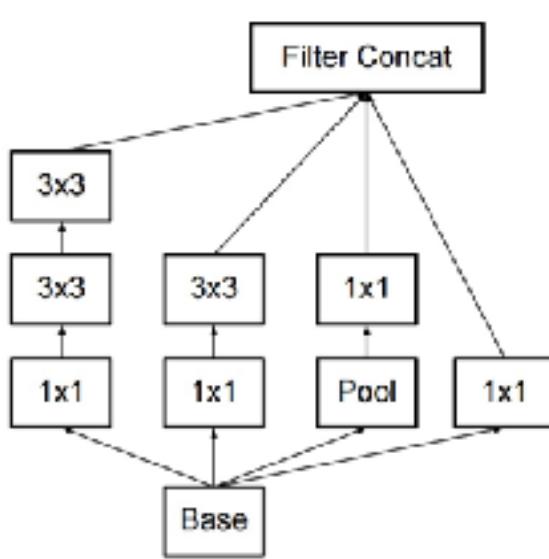
Figure 2: Inception module

<https://arxiv.org/pdf/1409.4842.pdf>

Types of CNN, 2015 February and December

A screenshot of Christian Szegedy's profile on Google Research. It shows his name, a photo, and research areas: Machine Intelligence and Machine Perception.

- **Inception V2**, Inception V1 with batch normalization
- **Inception V3:**
 - replace 5×5 with multiple 3×3



<https://arxiv.org/pdf/1512.00567.pdf>

41

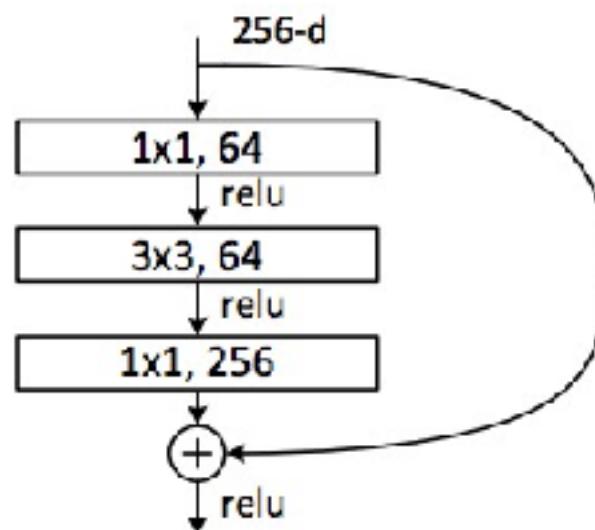
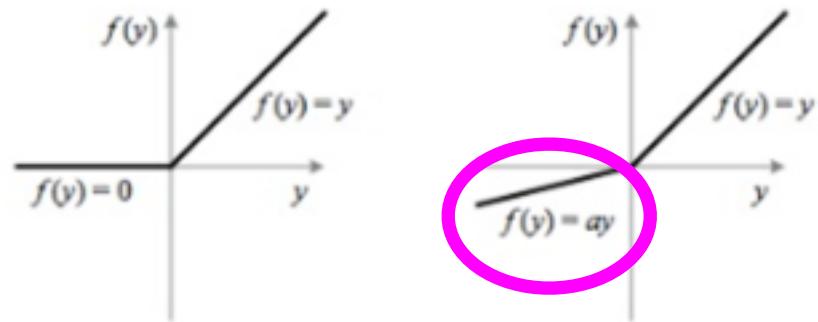
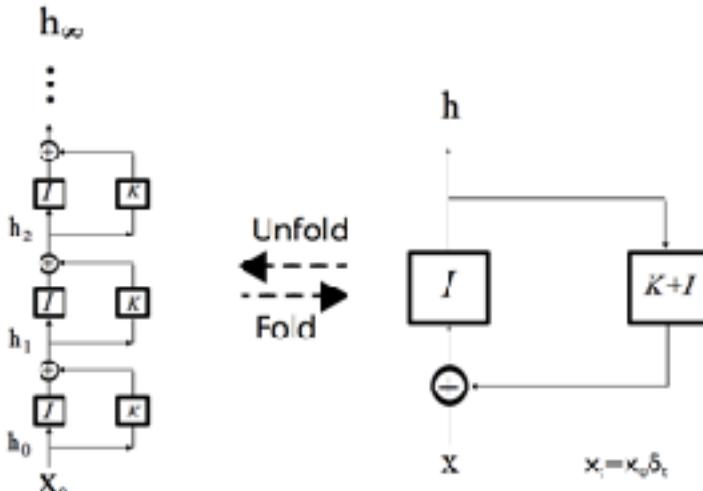
Types of CNN, 2015 December

Microsoft Research

- Major Contributions:
 - ensembles, not strictly sequential
 - bio-plausible with feedback

• ResNet

- PReLU: adaptive trained slope



- NiN: triple bypass layer
 - similar to bottleneck

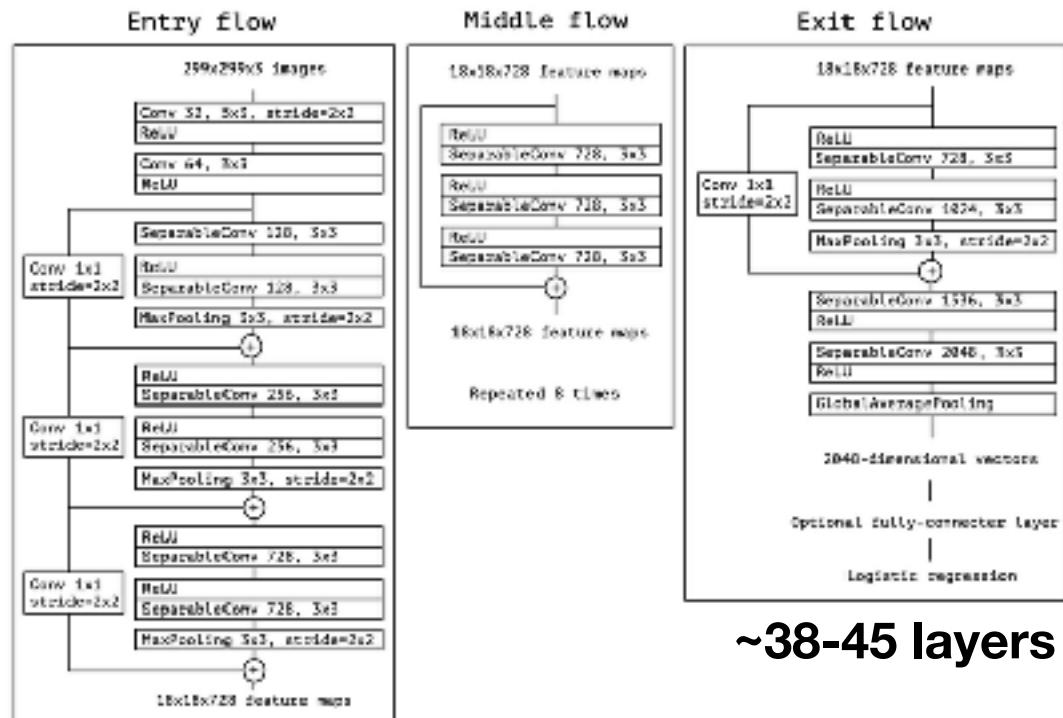
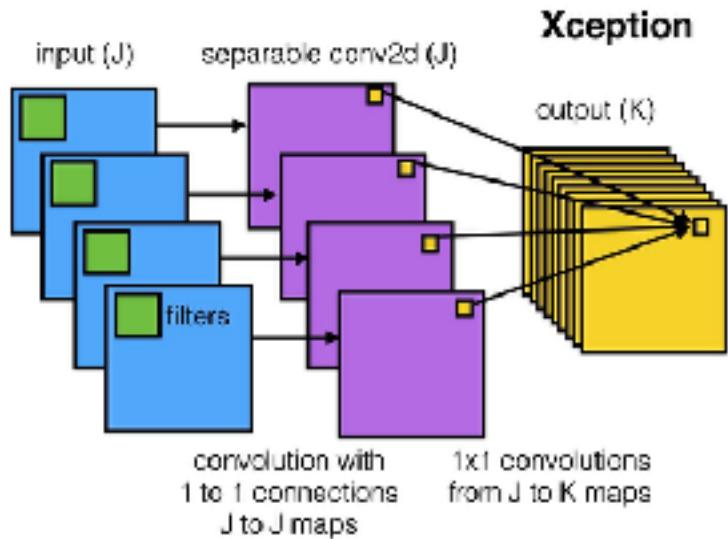
Types of CNN, 2017 April

Xception

- Major Contributions:
 - combining branching / residual blocks
 - separable convolutions



Francois Chollet
Google



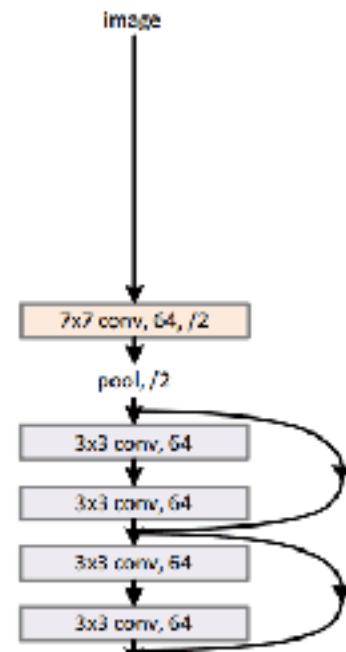
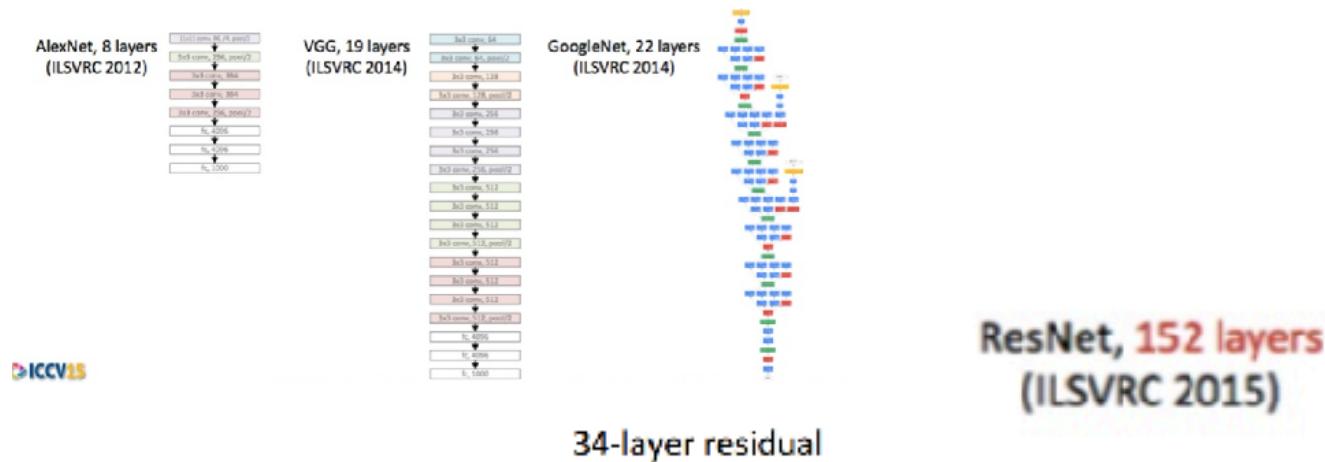
~38-45 layers



<https://arxiv.org/pdf/1610.02357.pdf>

How big are these networks?

How big are these networks?



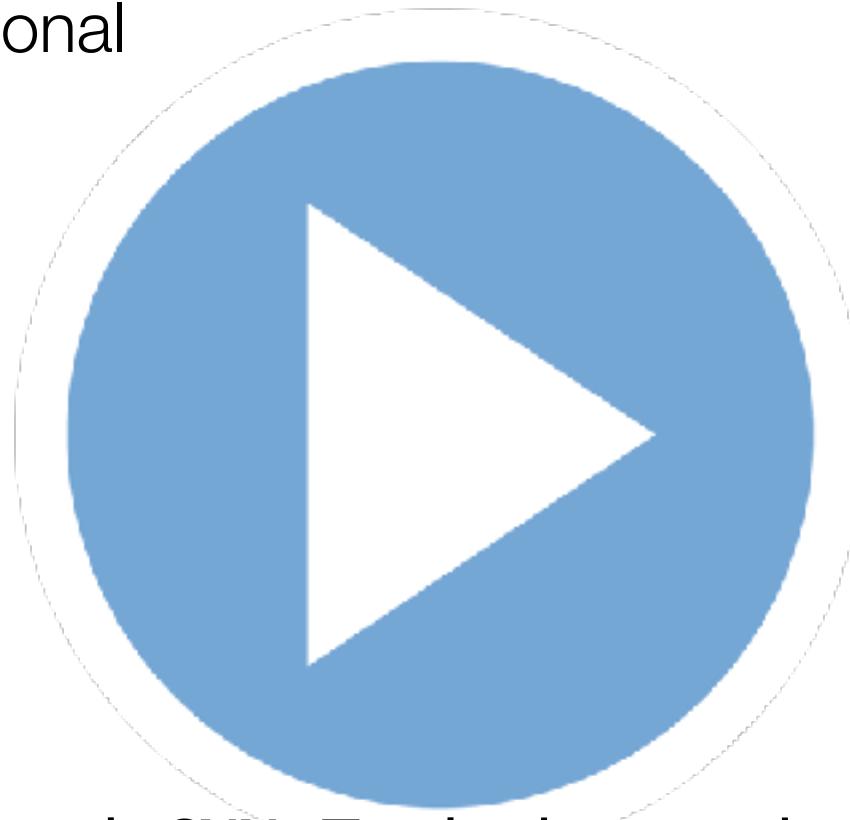
Self Test

- We have seen a lot of different networks.
- The most important concept to understand in using convolutional neural networks is:
 - A. Use proper initialization of layers
 - B. Have plenty of data or use expansion
 - C. Set aside time for training
 - D. Use batch normalization

Demo

Even more Convolutional
Neural Networks

...in TensorFlow
...with Keras



12. More Advanced CNN Techniques.ipynb

Next Time:

- CNN Lab Discussion (Town Hall)
- Intro to Recurrent Neural Network Architectures
 - RNNs, GRUs, LSTMs
 - Training for characters