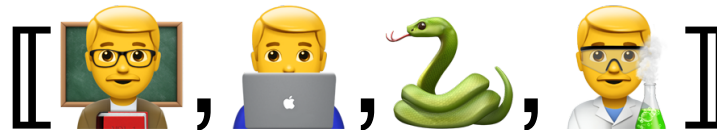


# Lecture Notes for **Machine Learning in Python**



Professor Eric Larson  
**Dimensionality Reduction**

# Class Logistics and Agenda

- Logistics:
  - First flipped module in one week!
  - Grading has commenced!
- Agenda:
  - Dimensionality Reduction
    - PCA
    - Randomized PCA
    - Images Representation with PCA

# Class Overview, by topic

Table Data  
Visualization

Numpy, Pandas, Seaborn  
Overviews with some in-depth discussion

Dimension  
Reduction and  
Image Processing

Scikit-learn, Scikit Image,  
Intuition only, Some mathematics

Linear and  
Logistic  
Regression

Numpy, Recreate API for Scikit-learn  
Detailed mathematics for simple optimization  
intuition for advanced optimization

Neural Networks  
and Back Prop.

Numpy  
Detailed mathematics for NN operations

Wide and Deep  
Networks

Convolutional  
Networks

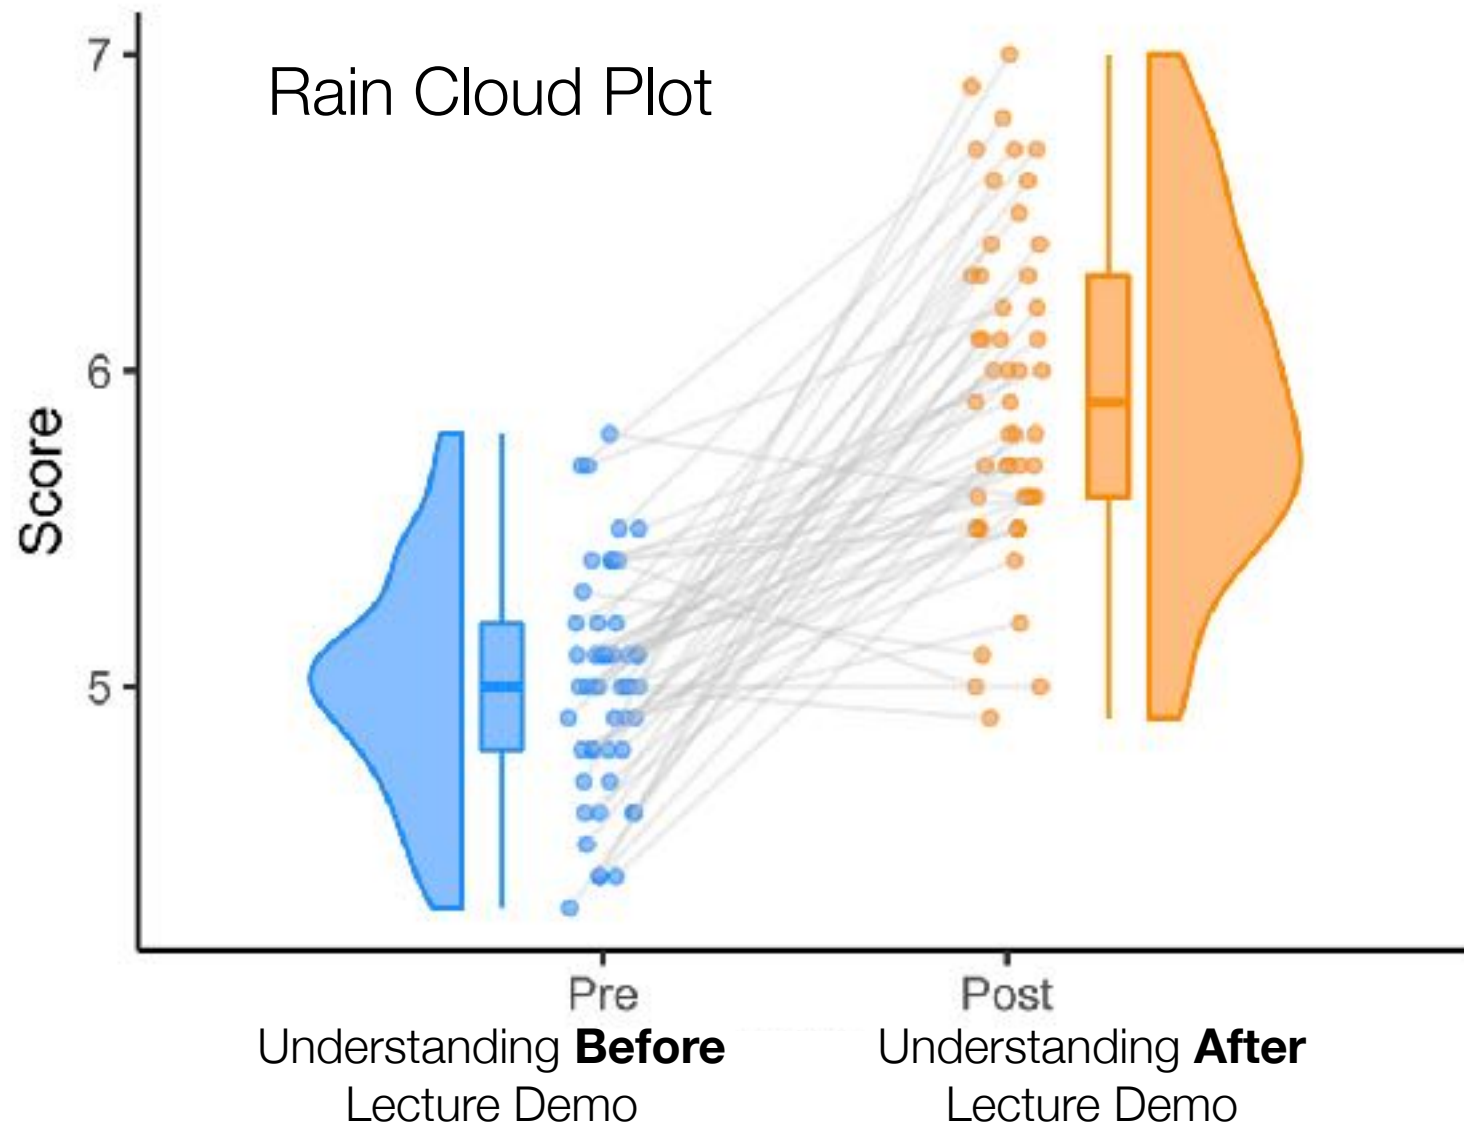
Recurrent  
Networks

Keras, Tensorflow  
Intuition, Detailed implement.

Ethics in  
Language Models

ConceptNet  
Case studies

# Last time: visualization



# Dimensionality Reduction: PCA



Kyle 🚀 🐬 🪐 🦖 @KyleMorgens... · 1d ...

eigenvalues are just the TLDR for a matrix

💬 38

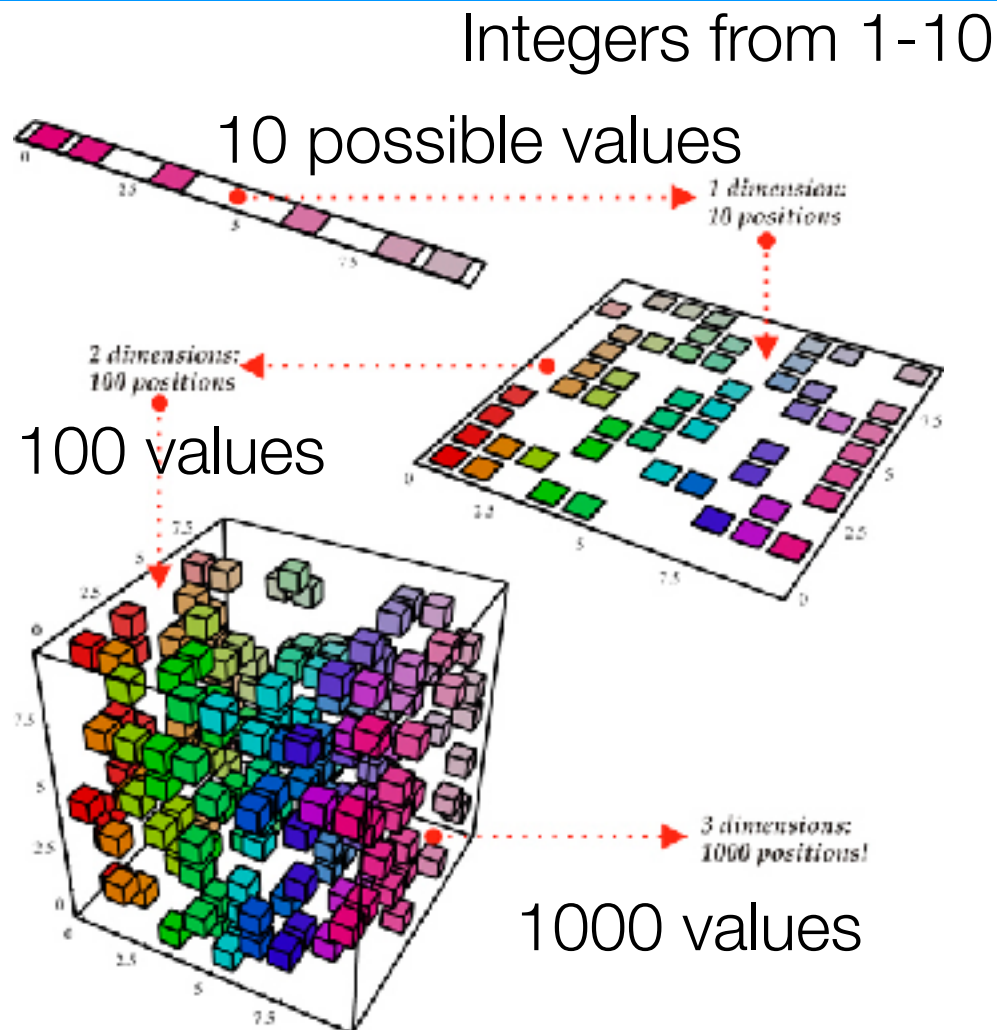
↻ 602

❤️ 6,046



# Curse of Dimensionality

- When dimensionality increases, data becomes increasingly sparse in the space that it occupies
- Definitions of density and distance between points, which is critical for clustering and outlier detection, become less meaningful



# Dimensionality Reduction

- Purpose:
  - Avoid curse of dimensionality
  - Select subsets of independent features
  - Reduce amount of time and memory required by data mining algorithms
  - Allow data to be more easily visualized
  - May help to eliminate irrelevant features or reduce noise

- Techniques

- Principle Component Analysis
- Non-linear PCA
- Stochastic Neighbor Embedding (tSNE)
- Uniform Manifold Approximation (UMAP)



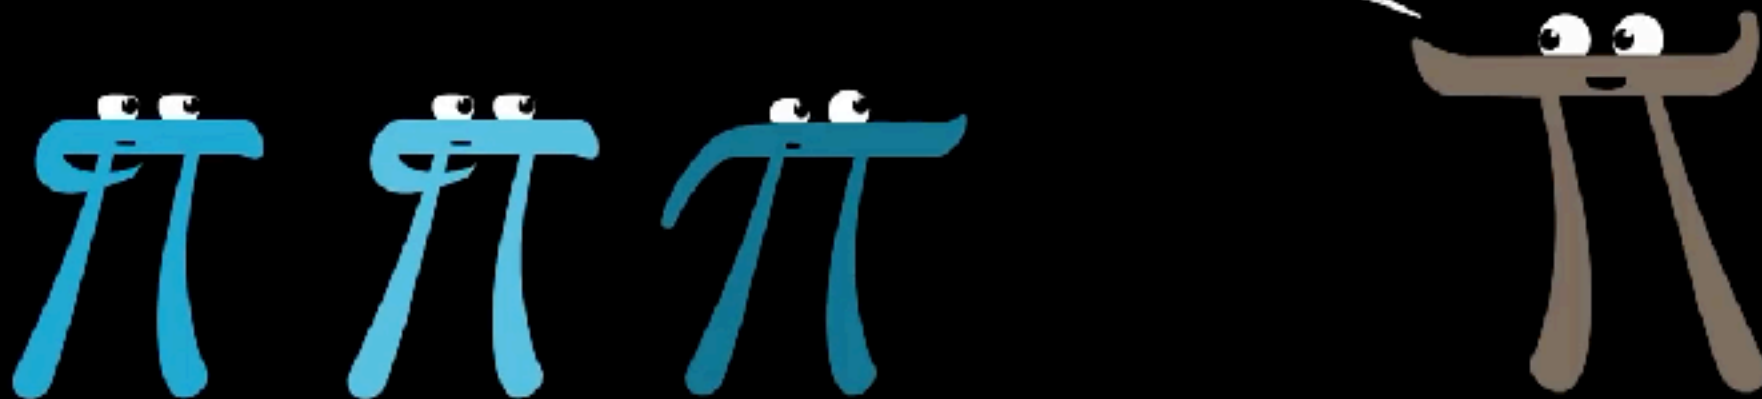
I invented PCA...  
and *Social Darwinism*

# Aside: Eigen Vectors are your friend!

(Grant Sanderson) **Three Blue One Brown:**

<https://www.youtube.com/watch?v=PFDu9oVAE-g>

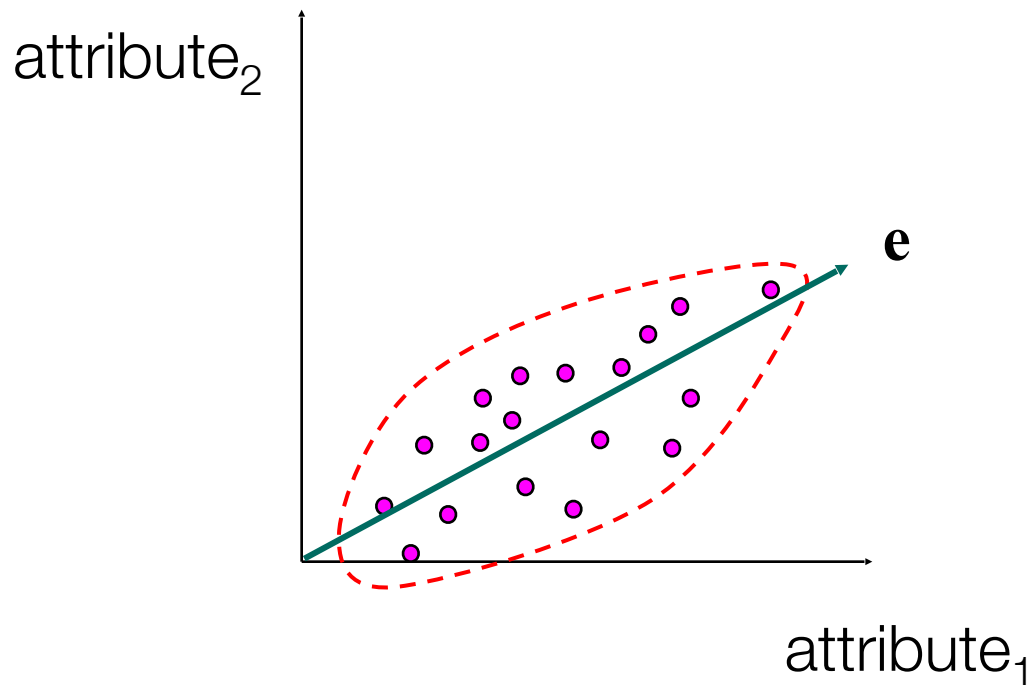
Eigen-things aren't  
actually so bad





# Dimensionality Reduction: PCA

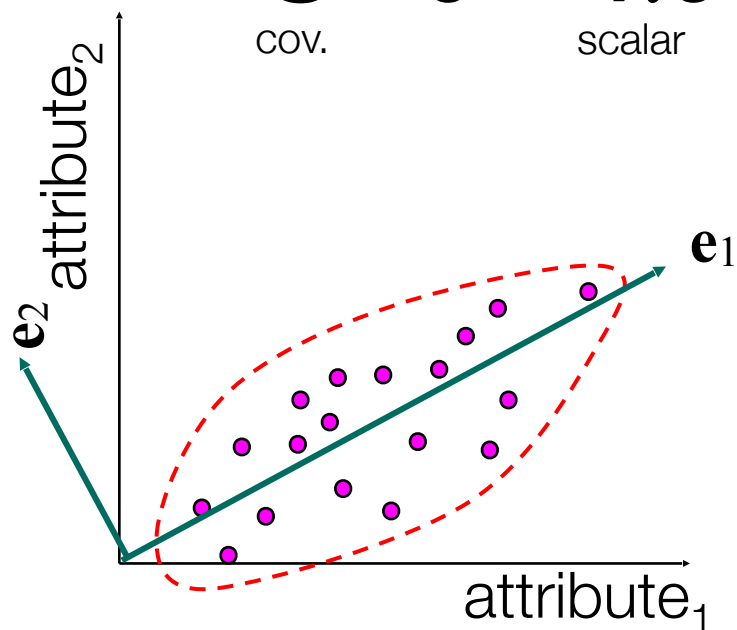
- Goal is to find a projection that captures the largest amount of variation in data



# Dimensionality Reduction: PCA

- Find the **eigenvectors** of the **covariance** matrix
- keep the “k” **largest** eigenvectors

$$\underset{\text{cov.}}{\mathbf{C}} \cdot \underset{\text{scalar}}{\mathbf{e}} = \lambda \mathbf{e}$$



| $E_1$           | $E_2$          |
|-----------------|----------------|
| 0.749           | 0.662          |
| 0.662           | -0.749         |
| $\lambda=268.3$ | $\lambda=1.57$ |

covariance

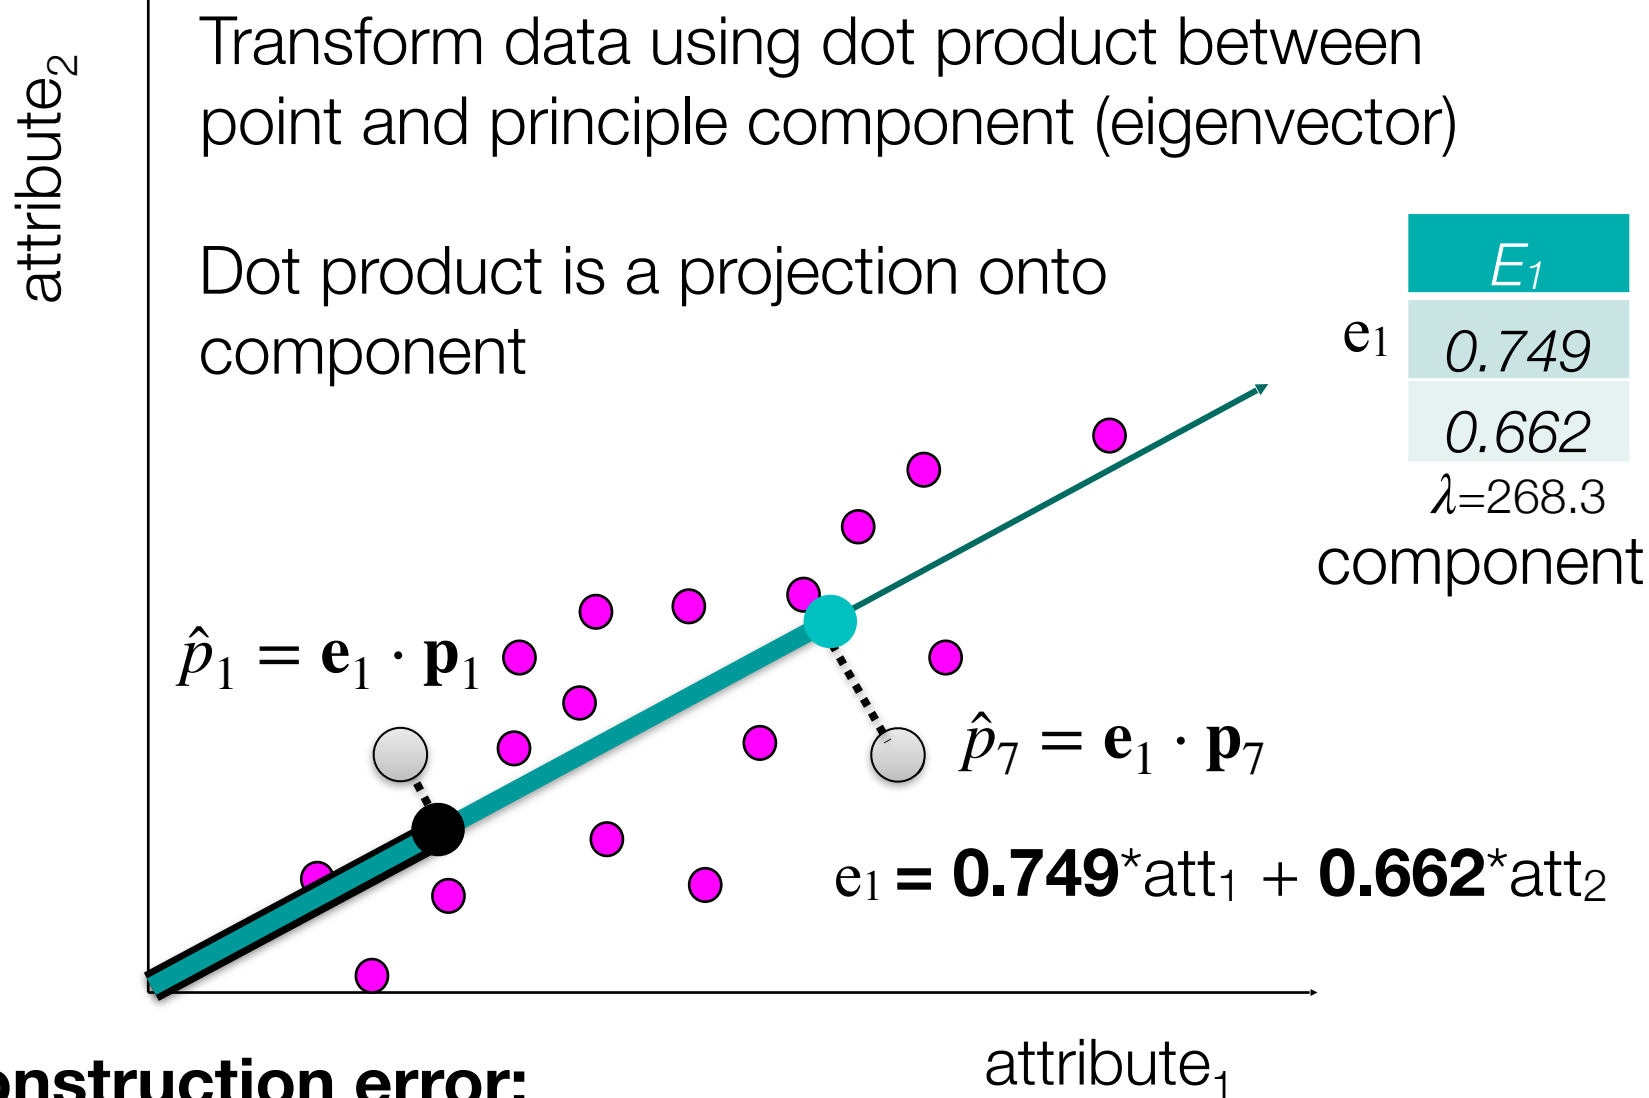
|       |       |
|-------|-------|
| 151.5 | 132.4 |
| 132.4 | 118.3 |

|   | $A_1$ | $A_2$ |
|---|-------|-------|
| 1 | 14    | 12.6  |
| 2 | 26    | 26.6  |
| 3 | 36.3  | 33.3  |
| 4 | 2.5   | 3.6   |
| 5 | 15    | 17.4  |
| 6 | 8     | 11    |

|   | $A'_1$ | $A'_2$ |
|---|--------|--------|
| 1 | -2.96  | -4.82  |
| 2 | 9.03   | 9.18   |
| 3 | 19.33  | 15.88  |
| 4 | -14.46 | -13.82 |
| 5 | -1.96  | -0.02  |
| 6 | -8.96  | -6.42  |

normalize: zero mean  
optional: unit std

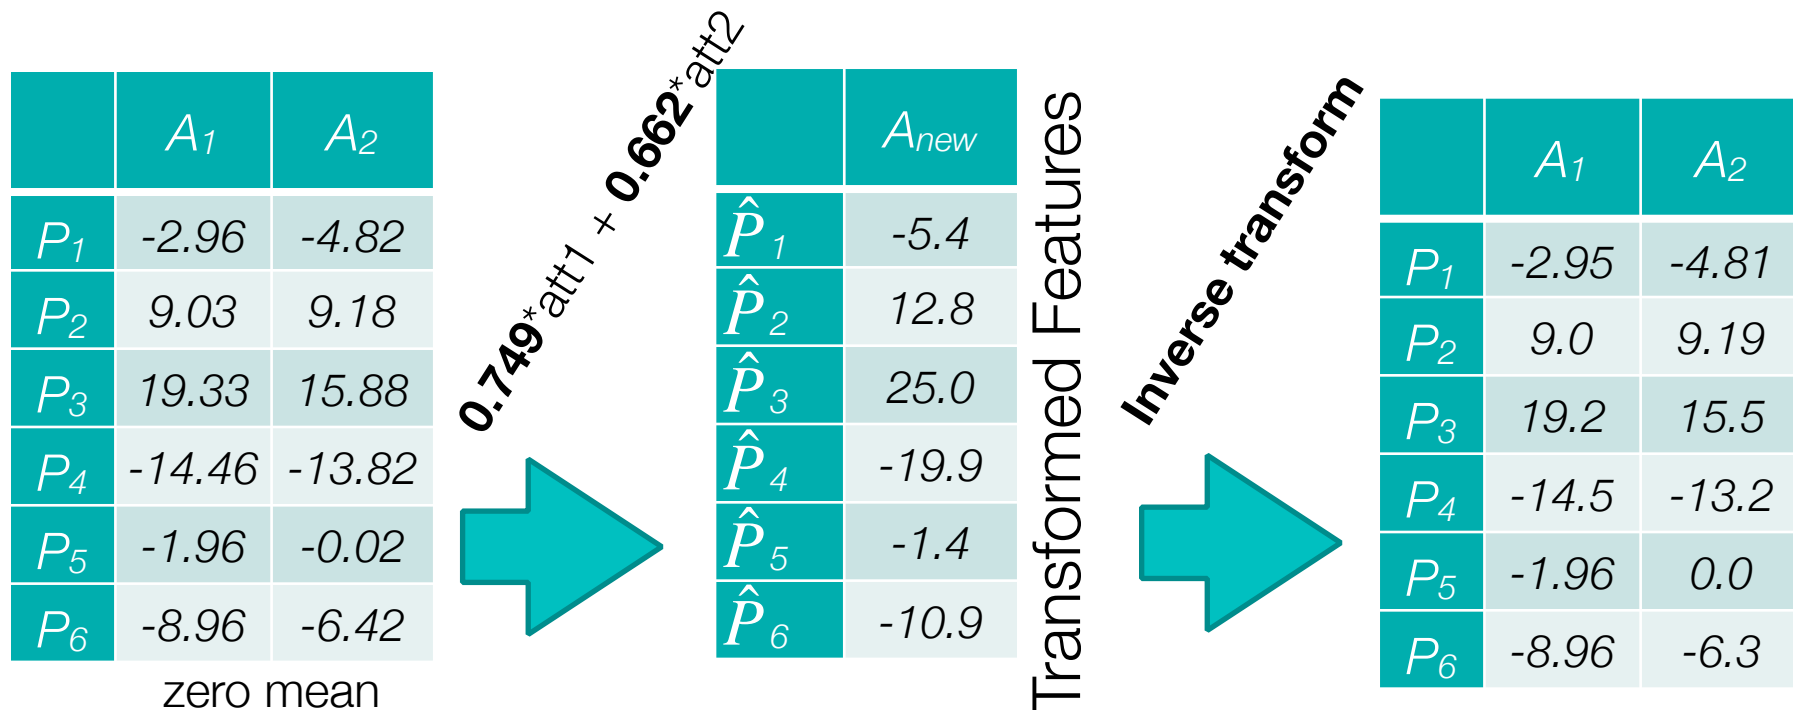
# Dimensionality Reduction: PCA



## Reconstruction error:

difference between projection and original point in 2D space

# Dimensionality Reduction: PCA



This projection is called a **Transform**  
known as the **Karhunen-Loève Transform (KLT)**

**Shown here for two dimensions, but  
could be arbitrarily large!!**

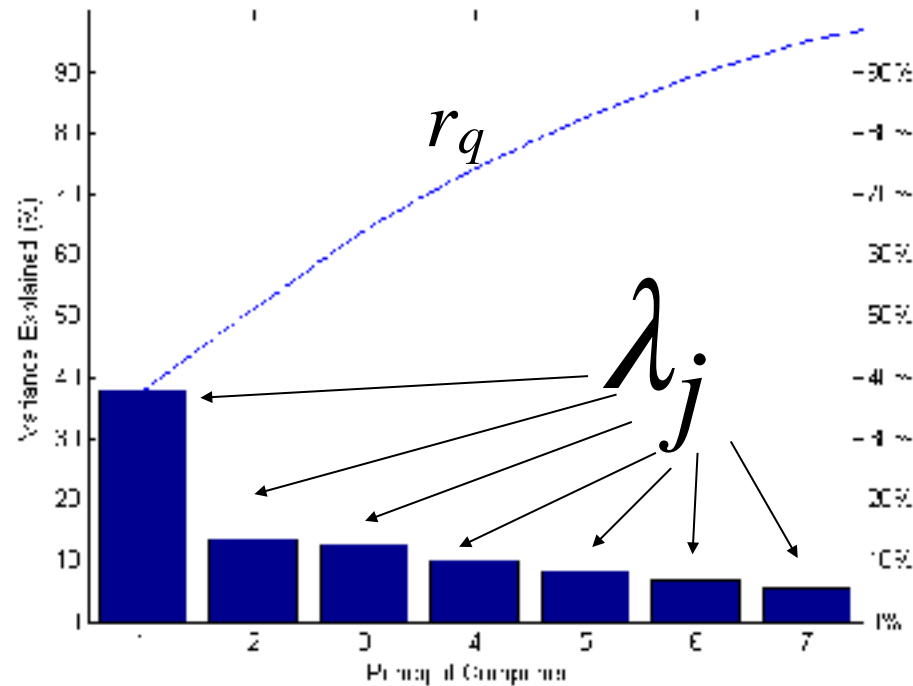
# Explained Variance (scree plot)

- Each principle component **explains** a certain **amount of variation** in the data.
- This explained variation is **encoded** in the **eigenvalues** of each **eigenvector**

sum of  $q$  largest eigenvalues

$$r_q = \frac{\sum_{j=1}^q \lambda_j}{\sum_{\forall i} \lambda_i}$$

sum of all eigenvalues



# Dimensionality Reduction: PCA

Genetic profiles distilled to 2 components

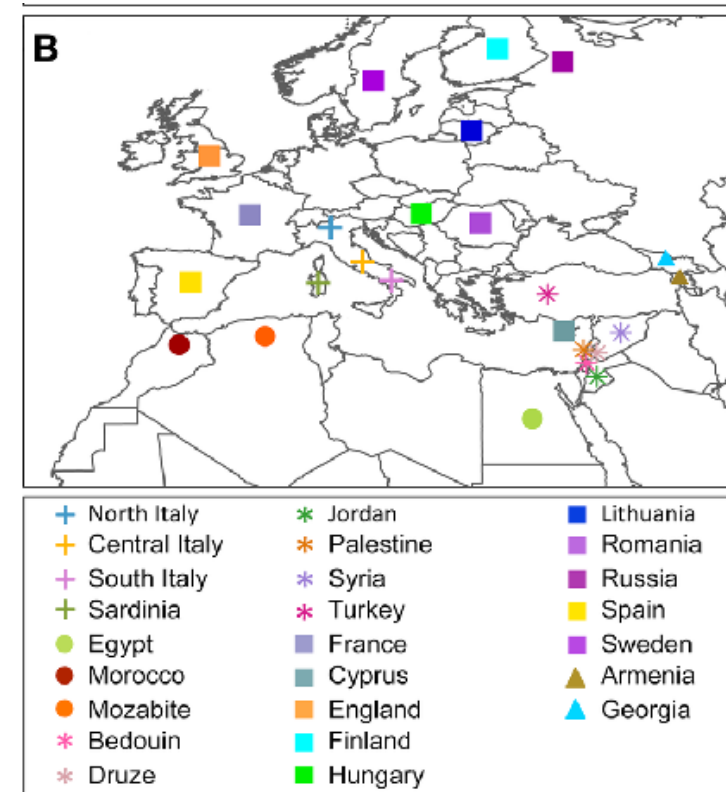
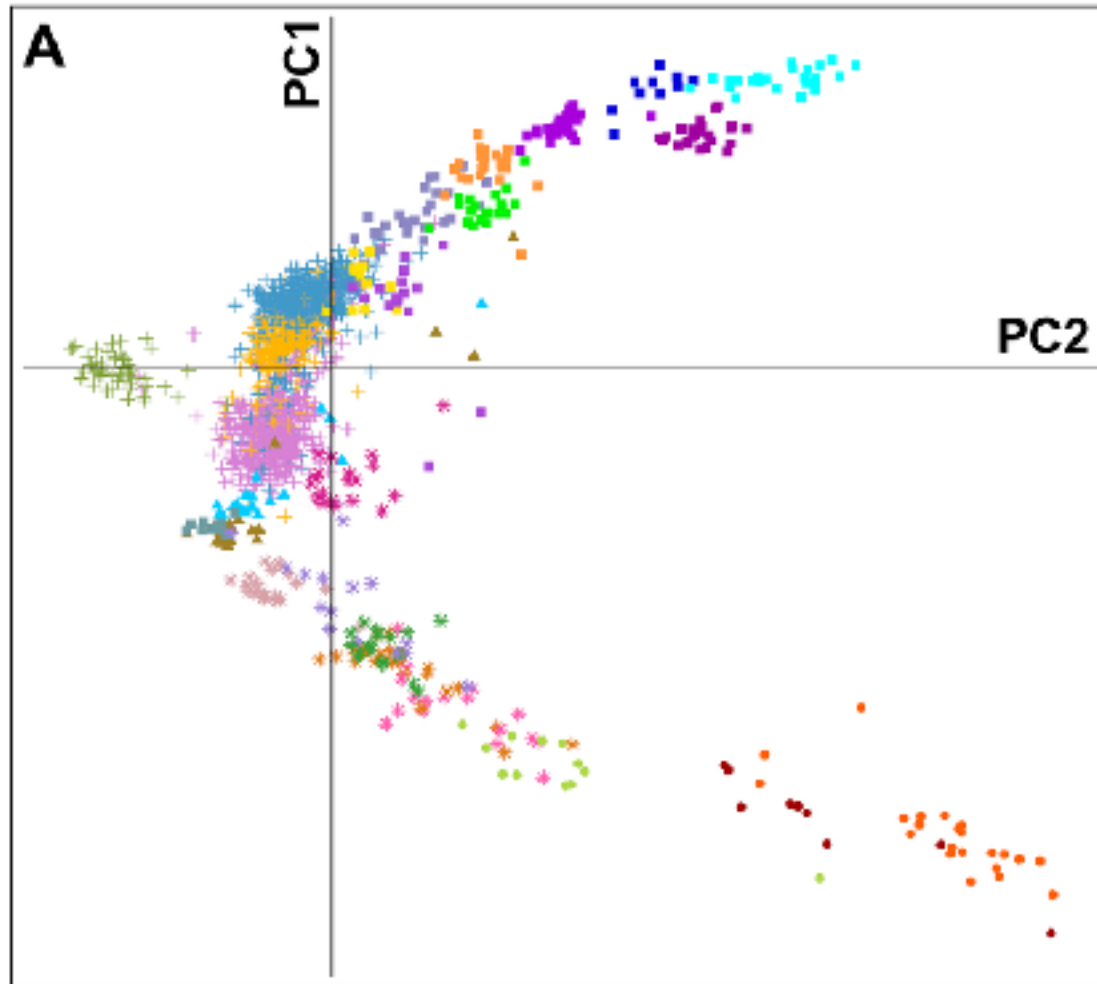


image source: Wikipedia 14

## 04.Dimension Reduction and Images.ipynb

PCA  
biplots



### Other Tutorials:

[http://scikit-learn.org/stable/auto\\_examples/decomposition/plot\\_pca\\_vs\\_lda.html#example-decomposition-plot-pca-vs-lda-py](http://scikit-learn.org/stable/auto_examples/decomposition/plot_pca_vs_lda.html#example-decomposition-plot-pca-vs-lda-py)

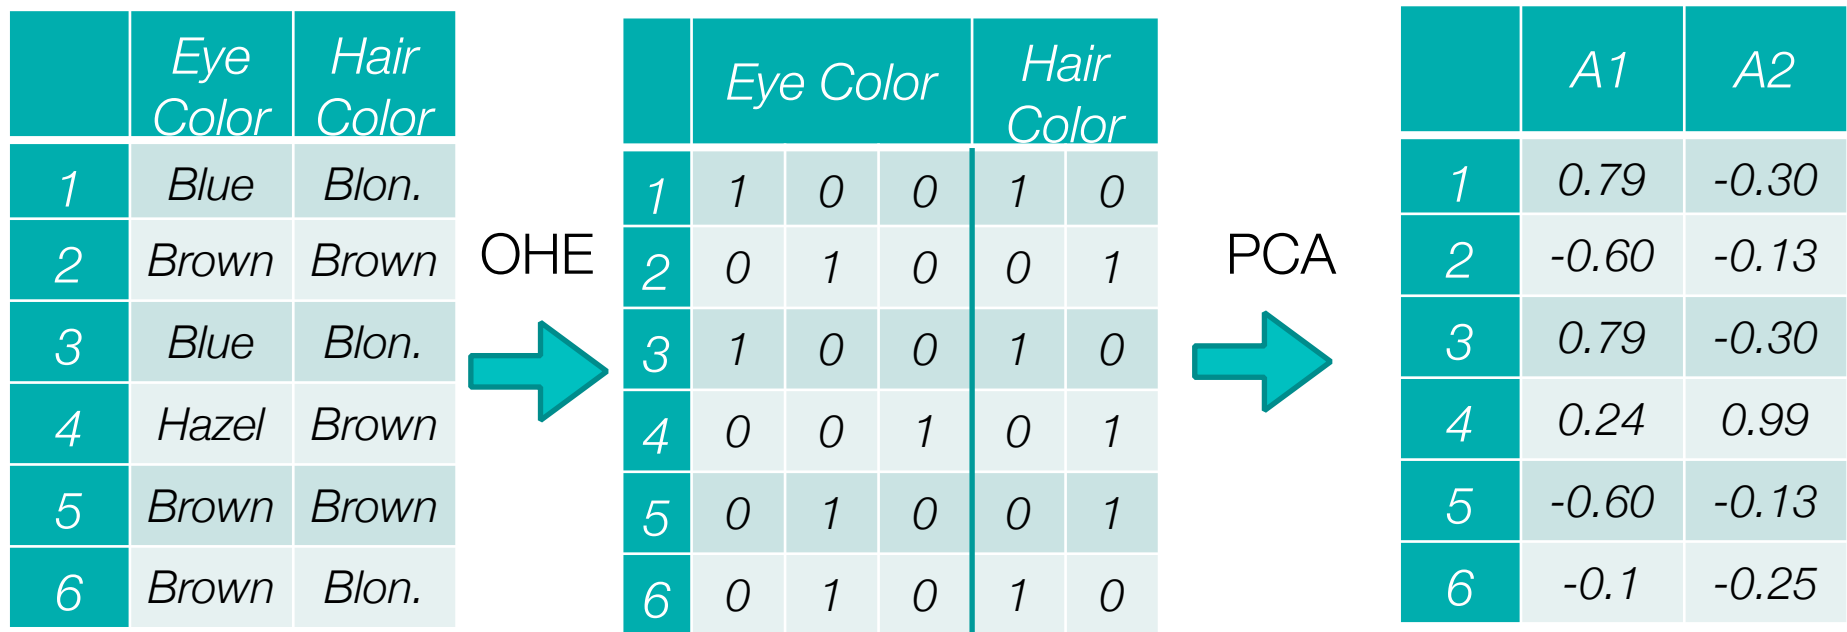
<http://nbviewer.ipython.org/github/ogrisel/notebooks/blob/master/Labeled%20Faces%20in%20the%20Wild%20recognition.ipynb>

# Self Test ML2b.1

Principal Components Analysis works well for categorical data by design.

- A. True
- B. False
- C. It doesn't but people do it anyway

**Better option:** Mutual Correspondence Analysis





# Dimensionality Reduction: Randomized PCA

- **Problem:** PCA on all that data can take a while to compute
  - What if the number of instances is gigantic?
  - What if the number of dimensions is gigantic?
- Can we approximate covariance with a lower rank matrix?
  - By **transforming** our table data,  $\mathbf{A}$ , with another orthogonal matrix,  $\mathbf{Q}$ , we can **approximate** the **covariance matrix**, but with **lower rank**
  - Gives a matrix with typically good enough precision of actual eigenvectors, like using SVD.  $\mathbf{Q}\mathbf{Q}^T\mathbf{A}$  is surrogate

Example Objective

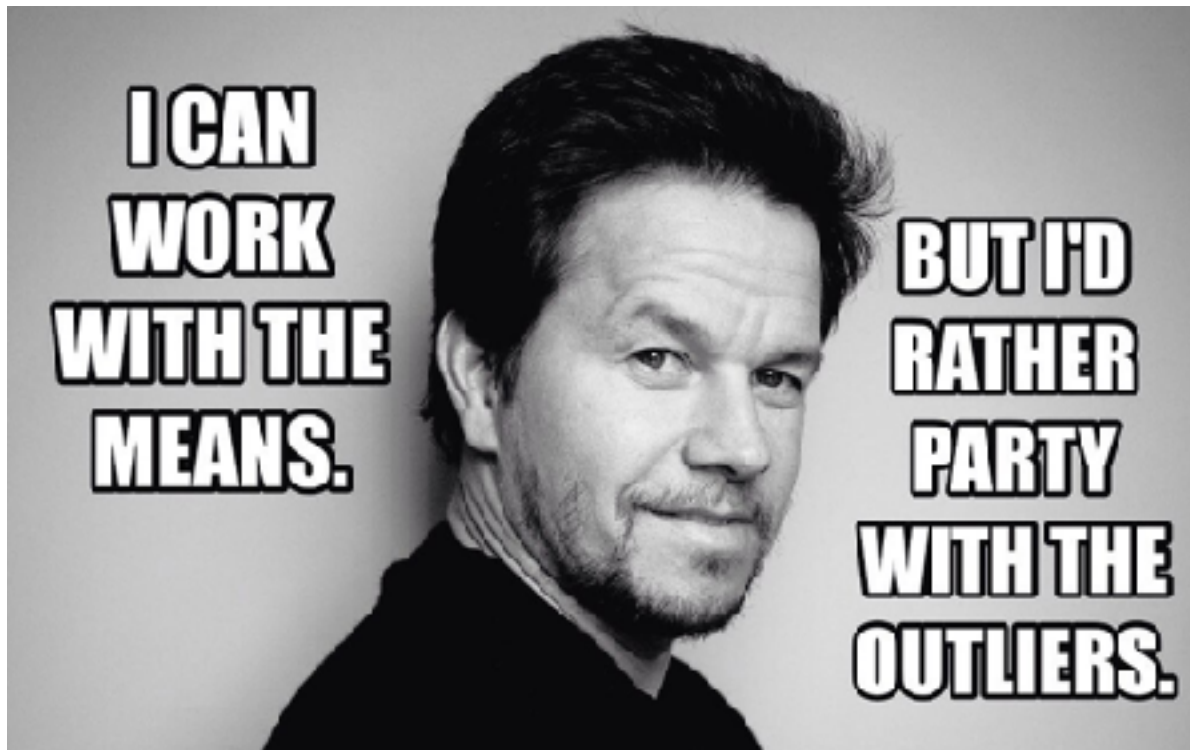
$$\|\mathbf{A} - \underbrace{\mathbf{Q} \cdot \mathbf{Q}^T \mathbf{A}}_{\text{surrogate}}\| < \underbrace{\left(1 + 11\sqrt{k+p} \cdot \min(m,n)\right) \cdot \sigma_{k+1}}_{\text{properties of } \mathbf{A} \text{ and } \mathbf{Q}}$$

Halko, et al., (2009) Finding structure with randomness: Stochastic algorithms for constructing approximate matrix decompositions. <https://arxiv.org/pdf/0909.4061.pdf>

**Just need an intuition about this!!!**

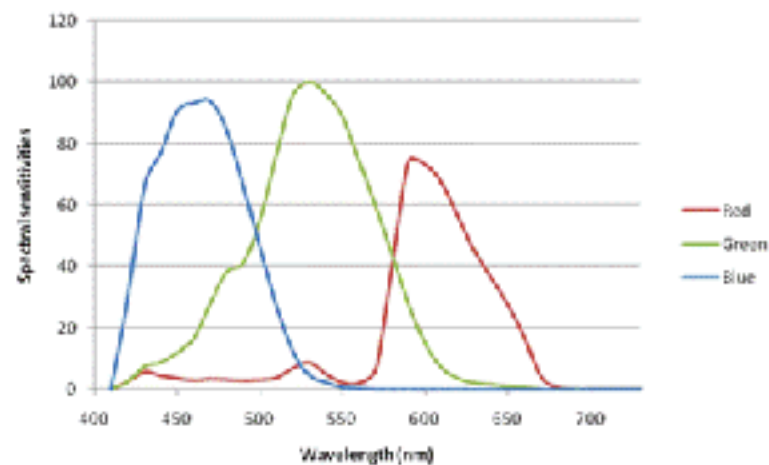
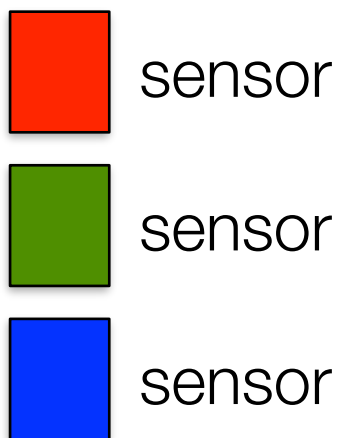
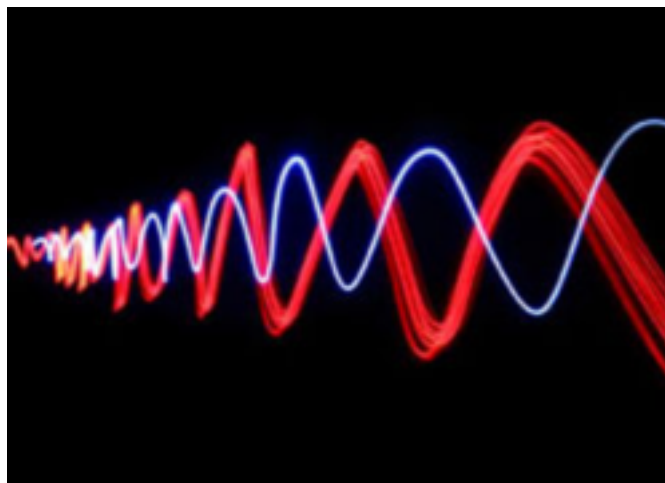
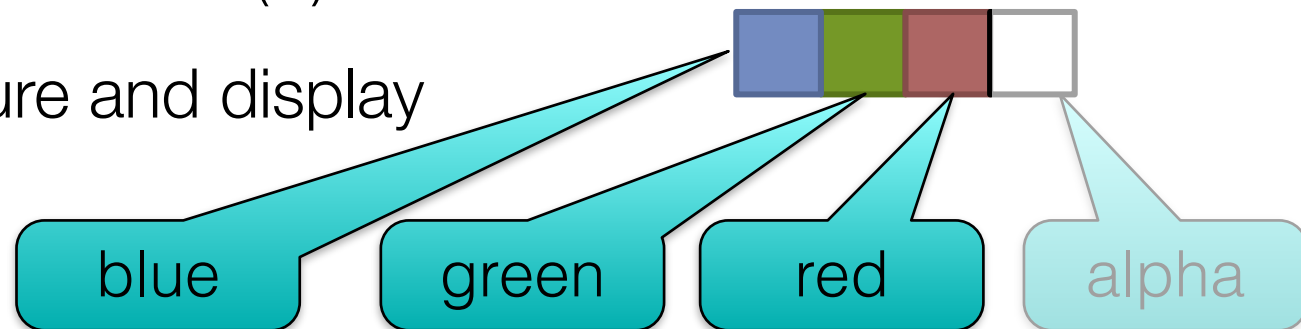
# Image Processing and Representation

Our first @ResearchMark meme



# Images as data

- an image can be represented in many ways
- most common format is a matrix of pixels
  - each “pixel” is BGR(A)
- used for capture and display



# Image Representation

- need a compact representation

- **grayscale**

$$0.3 \cdot R + 0.59 \cdot G + 0.11 \cdot B,$$

“luminance”

gray

|   |   |   |   |   |   |
|---|---|---|---|---|---|
| 1 | 4 | 2 | 5 | 6 | 9 |
| 1 | 4 | 2 | 5 | 5 | 9 |
| 1 | 4 | 2 | 8 | 8 | 7 |
| 3 | 4 | 3 | 9 | 9 | 8 |
| 1 | 0 | 2 | 7 | 7 | 9 |
| 1 | 4 | 3 | 9 | 8 | 6 |
| 2 | 4 | 2 | 8 | 7 | 9 |

Numpy Matrix  
`image[rows, cols]`

|   |   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|---|
|   |   | R |   |   |   |   |   |   |
|   |   | G | 1 | 4 | 2 | 5 | 6 | 9 |
| B |   | 1 | 4 | 2 | 5 | 6 | 9 | 9 |
|   | 1 | 4 | 2 | 5 | 6 | 9 | 9 | 7 |
|   | 1 | 4 | 2 | 5 | 5 | 9 | 7 | 8 |
|   | 1 | 4 | 2 | 8 | 8 | 7 | 8 | 9 |
|   | 3 | 4 | 3 | 9 | 9 | 8 | 9 | 6 |
|   | 1 | 0 | 2 | 7 | 7 | 9 | 6 | 9 |
|   | 1 | 4 | 3 | 9 | 8 | 6 | 9 |   |
|   | 2 | 4 | 2 | 8 | 7 | 9 |   |   |

Numpy Matrix  
`image[rows, cols, channels]`

# Image Representation, Features

**Problem:** need to represent image as table data

- need a compact representation

|   |   |   |   |   |   |
|---|---|---|---|---|---|
| 1 | 4 | 2 | 5 | 6 | 9 |
| 1 | 4 | 2 | 5 | 5 | 9 |
| 1 | 4 | 2 | 8 | 8 | 7 |
| 3 | 4 | 3 | 9 | 9 | 8 |
| 1 | 0 | 2 | 7 | 7 | 9 |
| 1 | 4 | 3 | 9 | 8 | 6 |
| 2 | 4 | 2 | 8 | 7 | 9 |

# Image Representation, Features

**Problem:** need to represent image as table data

- need a compact representation

**Solution:** row concatenation (also, vectorizing)

|       |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
|-------|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| Row 1 | 1 | 4 | 2 | 5 | 6 | 9 | 1 | 4 | 2 | 5 | 5 | 9 | 1 | 4 | 2 | 8 | 8 | 7 | 3 |
| Row 2 | 1 | 4 | 2 | 8 | 8 | 7 | 3 | 4 | 3 | 9 | 9 | 8 | 1 | 4 | 2 | 5 | 5 | 9 | 1 |
| ...   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| Row N | 9 | 4 | 6 | 8 | 8 | 7 | 4 | 1 | 3 | 9 | 2 | 1 | 1 | 5 | 2 | 1 | 5 | 9 | 1 |

# Self test: 3a-1

- When vectorizing images into table data, each “feature column” corresponds to:
  - a. the value (color) of a pixel
  - b. the spatial location of a pixel in the image
  - c. the size of the image
  - d. the spatial location and color channel of a pixel in an image

Row N

|   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 9 | 4 | 6 | 8 | 8 | 7 | 4 | 1 | 3 | 9 | 2 | 1 | 1 | 5 | 2 | 1 | 5 | 9 | 1 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

Images Representation  
in PCA and  
Randomized PCA



04.Dimension Reduction and Images.ipynb



# For Next Lecture

- Next Lecture:
  - Finish Dimension Reduction Demo
  - Crash-course Image Feature Extraction