

# Lecture Notes for **Machine Learning in Python**



Professor Eric Larson  
**Dimensionality Reduction and Images**

# Class Logistics and Agenda

- **Logistics:**
  - Lab grading...
  - Do **quiz one!!**
  - Coldfront Allocation
  - **Next Time:** Flipped Module
    - Turn in one per team (HTML), please include team member names from canvas
- **Agenda**
  - Common Feature Extraction Methods for Images
  - Begin Town Hall, if time

# Class Overview, by topic

Table Data  
Visualization

Numpy, Pandas, Seaborn  
Overviews with some in-depth discussion

Dimension  
Reduction and  
Image Processing

Scikit-learn, Scikit Image,  
Intuition only, Some mathematics

Linear and  
Logistic  
Regression

Numpy, Recreate API for Scikit-learn  
Detailed mathematics for simple optimization  
intuition for advanced optimization

Neural Networks  
and Back Prop.

Numpy  
Detailed mathematics for NN operations

Wide and Deep  
Networks

Convolutional  
Networks

Recurrent  
Networks

Keras, Tensorflow  
Intuition, Detailed implement.

Ethics in  
Language Models

ConceptNet  
Case studies

# Last time...



$E_1$	$E_2$	covariance	
0.749	0.662	151.5	132.4
0.662	-0.749	132.4	118.3
$\lambda=268.3$	$\lambda=1.57$		

	$A_1$	$A_2$
1	14	12.6
2	26	26.6
3	36.3	33.3
4	2.5	3.6
5	15	17.4
6	8	11

	$A_1$	$A_2$
1	-2.96	-4.82
2	9.03	9.18
3	19.33	15.88
4	-14.46	-13.82
5	-1.96	-0.02
6	-8.96	-6.42

normalize: zero mean  
optional: unit std


$$r_q = \frac{\sum_{j=1}^q \lambda_j}{\sum_{j=1}^p \lambda_j}$$



- an image can be represented in many ways
- most common format is a matrix of pixels
  - each "pixel" is BGR(A)
- used for capture and display



- sensor
- sensor
- sensor



# Review: Image Representation, Features

**Problem:** need to represent image as table data

- need a compact representation

1	4	2	5	6	9
1	4	2	5	5	9
1	4	2	8	8	7
3	4	3	9	9	8
1	0	2	7	7	9
1	4	3	9	8	6
2	4	2	8	7	9

# Review: Image Representation, Features

**Problem:** need to represent image as table data

- need a compact representation

**Solution:** row concatenation (also, vectorizing)

Row 1	1	4	2	5	6	9	1	4	2	5	5	9	1	4	2	8	8	7	3
Row 2	1	4	2	8	8	7	3	4	3	9	9	8	1	4	2	5	5	9	1
...																			
Row N	9	4	6	8	8	7	4	1	3	9	2	1	1	5	2	1	5	9	1

## “Refresher” Demo

Images Representation  
in PCA and  
Randomized PCA



04.Dimension Reduction and Images.ipynb


# Features of Images





# Extracting Features: Convolution

- For images:
  - kernel (matrix of values)
  - slide kernel across image, pixel by pixel
  - multiply and accumulate



## **This Example:**

3x3 Kernel (dark)

Ignoring edges of input

Input Image is 5x5

Output is then 3x3

# Convolution

$$\sum \left( \mathbf{I} \left[ i \pm \frac{r}{2}, j \pm \frac{c}{2} \right] \odot \mathbf{k} \right) = \mathbf{O}[i, j] \quad \text{output image at pixel } i, j$$

input image slice centered in  $i, j$   
with range  $r \times c$

kernel of size,  $r \times c$   
usually  $r=c$

0	0	0	0	0	0	0	0	0
0	1	2	3	4	12	9	8	0
0	5	2	3	4	12	9	8	0
0	5	2	1	4	10	9	8	0
0	7	2	1	4	12	7	8	0
0	7	2	1	4	14	9	8	0
0	5	2	3	4	12	7	8	0
0	5	2	1	4	12	9	8	0
0	0	0	0	0	0	0	0	0

input image,  $\mathbf{I}$

0	0	0
2	3	4
2	3	4

1	2	1
2	4	2
1	2	1

kernel  
filter,  $\mathbf{k}$   
 $3 \times 3$   
 $r \times c$

20	21	36	...	...	...	...
...	...	...	...	...	...	...
...	...	...	...	...	...	...
...	...	...	...	...	...	...
...	...	...	...	...	...	...
...	...	...	...	...	...	...
...	...	...	...	...	...	...
...	...	...	...	...	...	...

output image,  $\mathbf{O}$

# Convolution Examples

## Self test:

What does this do?

A. move left pixel to center



B. move right to center

C. blur

0	0	0
1	0	0
0	0	0

## Blur

1	1	1
1	1	1
1	1	1




## Vertical Edges

-1	0	1
-1	0	1
-1	0	1



## Sharpen

0	-1	0
-1	5	-1
0	-1	0



Convolution is linear

with  $\alpha = 5$

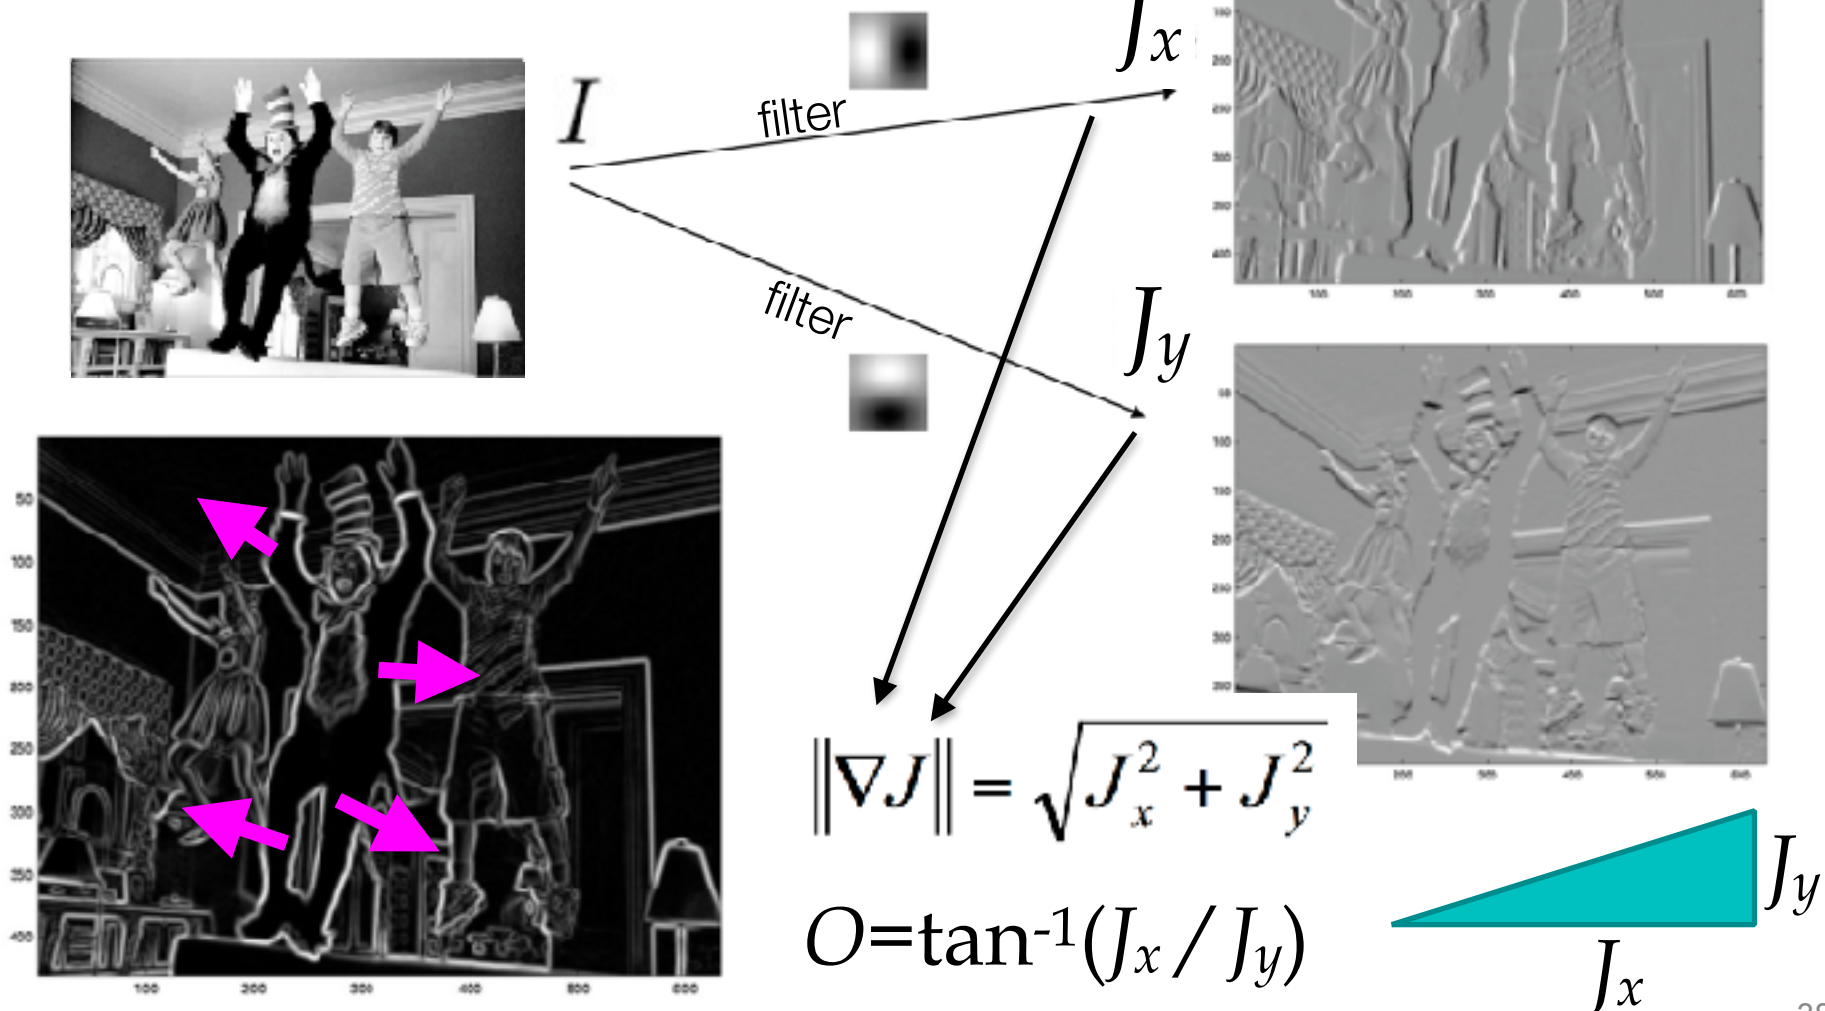
$$\alpha \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} + \left( \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \begin{bmatrix} 0 & 1 & 0 \\ 1 & 1 & 1 \\ 0 & 1 & 0 \end{bmatrix} \right) = \begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

$\alpha$  Image + (Image - Blur) = Sharp



<https://setosa.io/ev/image-kernels/>

# Common operations


- the gradient (2D derivative)



# DAISY: same features, regardless of orientation





invariant to rotations!



1. Select  $u, v$  pixel location in image and radius
2. Take histogram of average gradient magnitudes in circle for each orientation  $\tilde{h}_{\Sigma}(u, v)$
3. Select circles in a ring,  $R_1$
4. For each circle on the ring, take another histogram  $\tilde{h}_{\Sigma}(\mathbf{I}_O(u, v, R_1))$
5. Repeat for more rings:  $R_2, R_3$
6. Save all histograms as “descriptors”  
 $[\tilde{h}_{\Sigma}(\cdot), \tilde{h}_{\Sigma}(\mathbf{I}_1(\cdot, R_1)), \tilde{h}_{\Sigma}(\mathbf{I}_2(\cdot, R_1)) \dots \tilde{h}_{\Sigma}(\mathbf{I}_7(\cdot, R_2)), \tilde{h}_{\Sigma}(\mathbf{I}_8(\cdot, R_2))]$
7. Concatenate as “feature” vector at that pixel location

# Efficient DAISY, Orient x Circle Radius convolutions



take histogram of convolved images at points  $u, v$

one convolution per orientation

one convolve per ring size

Daisy Operator at  $u_0, v_0$  is Concatenated  $||\text{Histograms}||$

$$\mathcal{D}(u_0, v_0) = \text{take } \mathbf{normalized} \text{ histogram of magnitudes}$$

$$[ \tilde{h}_Z(u_0, v_0), \tilde{h}_Z(\mathbf{l}_1(u_0, v_0, R_1)), \tilde{h}_Z(\mathbf{l}_2(u_0, v_0, R_1)) \dots \tilde{h}_Z(\mathbf{l}_7(u_0, v_0, R_2)), \tilde{h}_Z(\mathbf{l}_8(u_0, v_0, R_2)) ]$$

**Tola et al.** "Daisy: An efficient dense descriptor applied to wide- baseline stereo." Pattern Analysis and Machine Intelligence, IEEE



# An intuition for the future: DAISY workflow

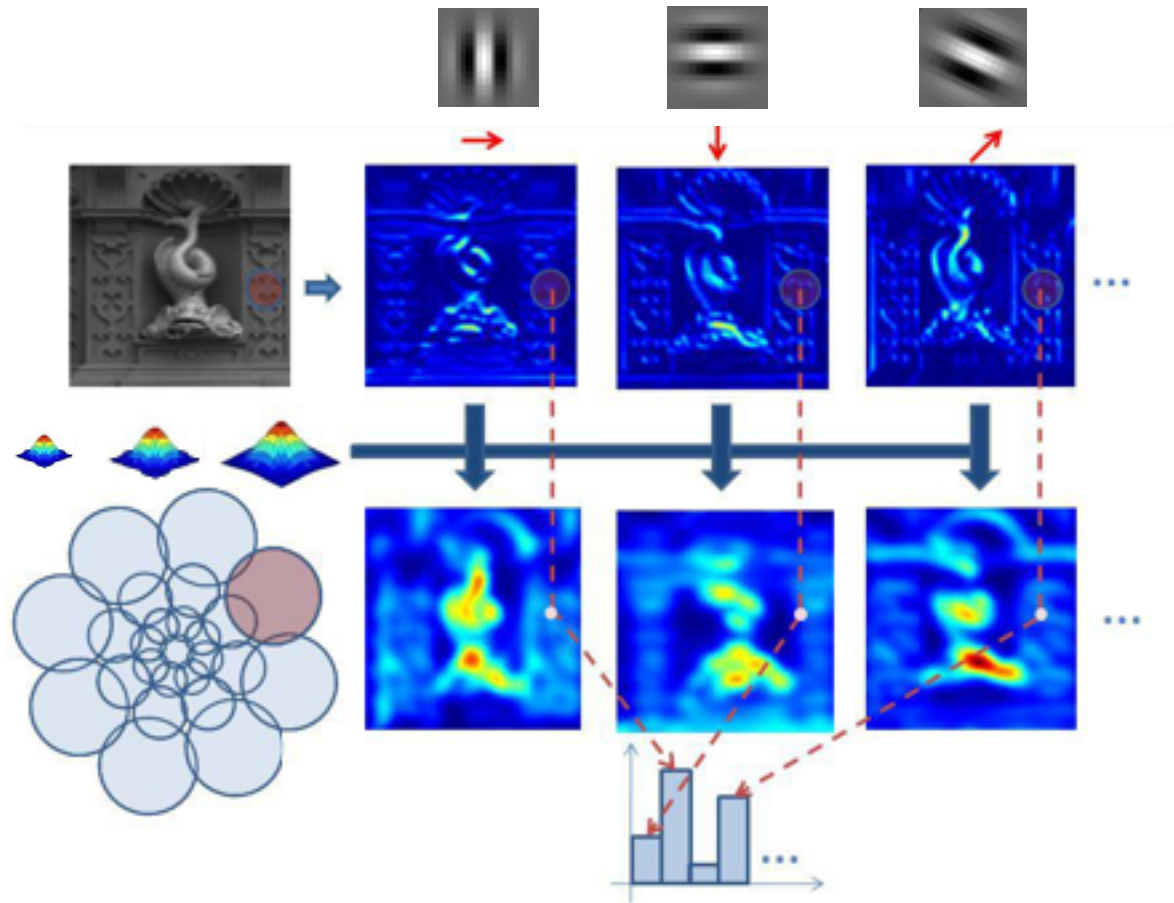
Convolutions on  
Image




Convolutions of  
Previous  
Convolutions



Statistics  
of Convolution  
Magnitudes




# Hyper Parameters in DAISY, need selection



Num rings

Num orientations within each circle

*Num of bins in histogram*



```
daisy(img, step=180, radius=58, rings=2, histograms=6,  
      orientations=8, visualize=True)
```

**Params**

step, radius, num rings, num histograms per ring,  
orientations, *bins per histogram*



# Classification with Daisy

- For each image:
  - Calculate daisy matrix (operator values)
  - Flatten into row
- Now we have a Table of Daisy Features (for each image)
- Separate Table into train and test
- Train your favorite classifier
  - Maybe a nearest neighbor classifier



Gradients  
DAISY



**Other Tutorials:**

[http://scikit-image.org/docs/dev/auto\\_examples/](http://scikit-image.org/docs/dev/auto_examples/)

# Matching versus Bag of Features

- Not a difference of vectors, but a percentage of matching points



- SURF, ORB, SIFT, DAISY

# Feature Matching

## Matching test image to source dataset

1. Choose src image from dataset
2. Take keypoints of src image
3. Take keypoints of test image
4. For each kp in src:
  1. Match with closest kp in test
  2. *How to define match?*
5. Count number of matches between images
6. Determine if src and test are similar based on number of matches
7. Repeat for new src image in dataset
8. Once all images measured, choose best match as the target for the test image



## Scikit-image Implementation

### match\_descriptors

```
skimage.feature.match_descriptors(descriptors1, descriptors2, metric=None, p=2,  
max_distance=inf, cross_check=True, max_ratio=1.0)
```

[\[source\]](#)

Brute-force matching of descriptors.

For each descriptor in the first set this matcher finds the closest descriptor in the second set (and vice-versa in the case of enabled cross-checking).

# Town Hall for Lab 2, Images

- **Quiz is live:** Image Processing!
- **Next Time:** Logistic Regression

