

Lecture Notes for **Machine Learning in Python**



Professor Eric Larson
MLP History

Class Logistics and Agenda

- Logistics:
 - Grading Update (Lab 3 due this weekend)
 - **Next time: Flipped Module on back propagation (work on your own time)**
- Multi Week Agenda:
 - Today: Neural Networks History, up to 1980 and Multi-layer Architectures
 - **Flipped:** Programming Multi-layer training
 - More Neural Networks,
 - Town Hall, Lab 4 (after flipped)
 - **Flipped:** Cross Validation

Class Overview, by topic

Table Data
Visualization

Numpy, Pandas, Seaborn
Overviews with some in-depth discussion

Dimension
Reduction and
Image Processing

Scikit-learn, Scikit Image,
Intuition only, Some mathematics

Linear and
Logistic
Regression

Numpy, Recreate API for Scikit-learn
Detailed mathematics for simple optimization
intuition for advanced optimization

Neural Networks
and Back Prop.

Numpy
Detailed mathematics for NN operations

Wide and Deep
Networks

Convolutional
Networks

Recurrent
Networks

Keras, Tensorflow
Intuition, Detailed implement.

Ethics in
Language Models

ConceptNet
Case studies

Lab 3, Town Hall (if needed)

Hey, what's your
name?



Right, what a stupid question.
I apologize, silly me.
I recognize the logo now.



...Anyway, I'm

Program received signal SIGSEGV,
Segmentation fault.



```
#0 0x00000000005009e0 in Lang::Name  
    (this=0x601172 <name>) at name.cpp:9
```



```
#1 0x0000000000500980 in  
    __static_initialization_and_destruction_0  
    (__initialize_p=1, __priority=65535)  
    at name.cpp:19
```



```
#2 0x00000000005009aa in  
    _GLOBAL__sub_name_print () at name.cpp:19
```



Python!




A History of Neural Networks




Neurons


- From biology to modeling:



McCulloch and Pitts, 1943



dendrite



input

$$\mathbf{X} \cdot \mathbf{W} = a$$

for each neuron

logic gates of the mind




Warren McCulloch



Walter Pitts

Neurons

- McCulloch and Pitts, 1943
- Donald Hebb, 1949
 - Hebb's Law: close neurons fire together
 - neurons "learn"
 - easier synaptic
 - basis of neural



I was infatuated with the idea of **brainwashing** and controlling minds of others! I also invented a number of **torture procedures** like sensory deprivation and **isolation tanks**—and carried out a number of secret studies on real people!!



Donald O. Hebb



Warren McCulloch




Walter Pitts

Rosenblatt's perceptron, 1957




Frank Rosenblatt



Axon Functions

hard limit




$$\begin{aligned} a &= -1 & z < 0 \\ a &= 1 & z \geq 0 \end{aligned}$$

linear




$$a = z$$

sigmoid




$$a = \frac{1}{1 + \exp(-z)}$$

dendrite




input

The Mark 1



PERCEPTRON




Perceptron Learning Rule:
~Stochastic Gradient Descent


Layers Notation for Table Data

One Neuron

dendrite



input



$\mathbf{x}^{(i)}$ One **row** from Table data as a
input **column** to model

$\mathbf{a}=\mathbf{x}$

$$\mathbf{x}^{(i)} = \begin{bmatrix} x_1 \\ \vdots \\ x_j \\ \vdots \\ x_N \end{bmatrix}^{(i)} = [\mathbf{a}^{(1)}]^{(i)}$$

$$\mathbf{z} = \mathbf{W} \cdot \mathbf{x}^{(i)} + \mathbf{b}$$

one neuron \rightarrow


$$\mathbf{W} = \begin{bmatrix} w_{1,1} & \dots & w_{1,N} \\ w_{2,1} & \dots & w_{2,N} \\ \vdots & & \vdots \\ w_{S,1} & \dots & w_{S,N} \end{bmatrix} \quad \mathbf{b} = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_S \end{bmatrix}$$

$$[\mathbf{z}^{(1)}]^{(i)} = \mathbf{W}^{(1)} \cdot [\mathbf{a}^{(1)}]^{(i)} + \mathbf{b}^{(1)}$$

$$\mathbf{a}^{next} = \phi(\mathbf{z}^{current})$$

$$\mathbf{a}^{(next)} = \begin{bmatrix} \phi(z_1^{curr}) \\ \vdots \\ \phi(z_N^{curr}) \end{bmatrix} \rightarrow \mathbf{a}^{(L)} = \begin{bmatrix} \phi(z_1^{L-1}) \\ \vdots \\ \phi(z_N^{L-1}) \end{bmatrix}$$

Generic Multiple Layers Notation



$$\mathbf{a}^{(L+1)} = \phi(\mathbf{z}^{(L)})$$


$$\mathbf{a}^{(final)} \text{ size=unique classes, } C$$

$$\mathbf{z}^{(L)} = \mathbf{W}^{(L)} \cdot \mathbf{a}^{(L)} + \mathbf{b}^{(L)}$$

$$\mathbf{z}^{(L)} = \mathbf{W}^{(L)} \cdot \phi(\mathbf{z}^{(L-1)}) + \mathbf{b}^{(L)}$$


$$\mathbf{z}^{(L)} = \mathbf{W}^{(L)} \cdot \phi(\mathbf{W}^{(L-1)} \cdot \phi(\mathbf{z}^{(L-2)}) + \mathbf{b}^{(L-1)}) + \mathbf{b}^{(L)}$$

Multiple layers notation



- **Self test:** How many parameters need to be trained in the above network?
 - A. $[(N+1) \times S^1] + [(S^1 + 1) \times S^2] + [(S^2 + 1) \times S^3]$
 - B. $|W^{(1)}| + |W^{(2)}| + |W^{(3)}| + |b^{(1)}| + |b^{(2)}| + |b^{(3)}|$
 - C. can't determine from diagram
 - D. it depends on the sizes of intermediate variables, $z^{(i)}$


Compact feedforward notation



$$\mathbf{X}^T = \left[\begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_N \end{bmatrix}^{(1)}, \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_N \end{bmatrix}^{(2)}, \dots, \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_N \end{bmatrix}^{(M)} \right] = \left[\begin{bmatrix} a_0^{(1)} \\ a_1^{(1)} \\ \vdots \\ a_N^{(1)} \end{bmatrix}^{(1)}, \begin{bmatrix} a_0^{(1)} \\ a_1^{(1)} \\ \vdots \\ a_N^{(1)} \end{bmatrix}^{(2)}, \dots, \begin{bmatrix} a_0^{(1)} \\ a_1^{(1)} \\ \vdots \\ a_N^{(1)} \end{bmatrix}^{(M)} \right] = \mathbf{A}^{(1)}$$

Table Data Table Data, in Neural Net Notation

Compact feedforward notation



$$\mathbf{Z}^{(L)} = \mathbf{W}^{(L)} \cdot \mathbf{A}^{(L)} + \mathbf{b}^{(L)}$$

$$\mathbf{Z}^{(L)} = \mathbf{W}^{(L)} \cdot \phi(\mathbf{Z}^{(L-1)}) + \mathbf{b}^{(L)}$$


$$[\mathbf{z}^{(L)}]^{(i)} = \mathbf{W}^{(L)} \cdot [\mathbf{a}^{(L)}]^{(i)} + \mathbf{b}^{(L)}$$

$$\begin{bmatrix} z_2^{(L)} \\ \vdots \\ z_{S^L}^{(L)} \end{bmatrix} = \mathbf{W}^{(L)} \cdot \begin{bmatrix} a_1^{(L)} \\ \vdots \\ a_{S^{L-1}}^{(L)} \end{bmatrix} + \mathbf{b}^{(L)}$$


$$\begin{bmatrix} \begin{bmatrix} z_1^{(L)} \\ z_2^{(L)} \\ \vdots \\ z_{S^L}^{(L)} \end{bmatrix}^{(1)}, \begin{bmatrix} z_1^{(L)} \\ z_2^{(L)} \\ \vdots \\ z_{S^L}^{(L)} \end{bmatrix}^{(2)}, \dots, \begin{bmatrix} z_1^{(L)} \\ z_2^{(L)} \\ \vdots \\ z_{S^L}^{(L)} \end{bmatrix}^{(M)} \end{bmatrix} = \mathbf{W}^{(L)} \cdot \begin{bmatrix} \begin{bmatrix} a_0^{(L)} \\ a_1^{(L)} \\ \vdots \\ a_{S^{L-1}}^{(L)} \end{bmatrix}^{(1)}, \begin{bmatrix} a_0^{(L)} \\ a_1^{(L)} \\ \vdots \\ a_{S^{L-1}}^{(L)} \end{bmatrix}^{(2)}, \dots, \begin{bmatrix} a_0^{(L)} \\ a_1^{(L)} \\ \vdots \\ a_{S^{L-1}}^{(L)} \end{bmatrix}^{(M)} \end{bmatrix} + \mathbf{b}^{(L)}$$

\mathbf{b} is broadcast added

Historical Training of Neural Network Architectures




When a binary classification model outputs 0.5



One Layer Linear Architectures

- Adaline network, Widrow and Hoff, 1960



Marcian "Ted" Hoff



Bernard Widrow

Simplify Objective Function:


$$J(\mathbf{W}) = \|\mathbf{Y} - \hat{\mathbf{Y}}\|^2 \longrightarrow J(\mathbf{w}) = \|\mathbf{Y} - \mathbf{A} \cdot \mathbf{w}\|^2$$

Need gradient $\nabla J(\mathbf{w})$ for update equation $\mathbf{w} \leftarrow \mathbf{w} + \eta \nabla J(\mathbf{w})$

We have been using the **Widrow-Hoff Learning Rule**

One Layer Linear Architectures

- Adaline network, Widrow and Hoff, 1960



Marcian "Ted" Hoff



Bernard Widrow


Need gradient $\nabla J(\mathbf{w})$ for update equation $\mathbf{w} \leftarrow \mathbf{w} + \eta \nabla J(\mathbf{w})$

For case $S=1$, \mathbf{W} has only one row, \mathbf{w} this is just **linear regression**...

$$J(\mathbf{w}) = \sum_{i=1}^M (y^{(i)} - \mathbf{x}^{(i)} \cdot \mathbf{w})^2$$
$$\mathbf{w} = (\mathbf{X}^T \cdot \mathbf{X})^{-1} \cdot \mathbf{X}^T \cdot \mathbf{y}$$



From Regression to Classification



No longer regression,
 \mathbf{Y} is a category

ground truth \mathbf{Y} is
one-hot encoded!

Ground Truth
Labels


$$\mathbf{Y} = \begin{bmatrix} \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_C \end{bmatrix}^{(1)} & \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_C \end{bmatrix}^{(2)} & \dots & \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_C \end{bmatrix}^{(M)} \end{bmatrix} \rightarrow \begin{bmatrix} \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}^{(1)} & \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix}^{(2)} & \dots & \begin{bmatrix} 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix}^{(M)} \end{bmatrix}$$

Need objective Function, minimize MSE $J(\mathbf{W}) = \left\| \mathbf{Y} - \hat{\mathbf{Y}} \right\|^2$

$$J(\mathbf{W}) = \left\| \underbrace{\begin{bmatrix} \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_C \end{bmatrix}^{(1)} & \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_C \end{bmatrix}^{(2)} & \dots & \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_C \end{bmatrix}^{(M)} \end{bmatrix}}_{\mathbf{Y}} - \underbrace{\begin{bmatrix} \begin{bmatrix} a_1^{(L)} \\ a_2^{(L)} \\ \vdots \\ a_C^{(L)} \end{bmatrix}^{(1)} & \begin{bmatrix} a_1^{(L)} \\ a_2^{(L)} \\ \vdots \\ a_C^{(L)} \end{bmatrix}^{(2)} & \dots & \begin{bmatrix} a_1^{(L)} \\ a_2^{(L)} \\ \vdots \\ a_C^{(L)} \end{bmatrix}^{(M)} \end{bmatrix}}_{\hat{\mathbf{Y}}} \right\|^2$$

One Layer Classification

- Rosenblatt's perceptron, 1957




Self Test - If this is a binary classification problem, how large is S , the length of $\hat{\mathbf{y}}$ and number of rows in \mathbf{W} ?

- A. Can't determine
- B. 2
- C. 1
- D. N

One Layer Classification

- Modern Perceptron network



$$g(z) = \phi(z) = \frac{1}{1 + \exp(-z)}$$

$$g(\mathbf{w} \cdot \mathbf{x}) = \phi(\mathbf{w} \cdot \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w} \cdot \mathbf{x})}$$

Need gradient $\nabla J(\mathbf{w})$ for update equation $\mathbf{w} \leftarrow \mathbf{w} + \eta \nabla J(\mathbf{w})$

For case $S=1$, this is just **logistic regression...**
and **we have already solved this!**


$$\mathbf{w} \leftarrow \mathbf{w} + \eta \cdot \text{mean} \left(\underbrace{(\mathbf{y} - g(\mathbf{X} \cdot \mathbf{w})) \odot \mathbf{X}}_{\mathbf{y}_{diff}} \right)_{cols}$$



What happens when $S > 1$?



One Layer Architectures of Many Classes



$$J(\mathbf{w}_{row=1}) = \sum_i \left(y_1^{(i)} - \phi(\mathbf{w}_{row=1} \cdot \mathbf{x}^{(i)}) \right)^2$$

... for each class/row ...

$$J(\mathbf{w}_{row=C}) = \sum_i \left(y_C^{(i)} - \phi(\mathbf{w}_{row=C} \cdot \mathbf{x}^{(i)}) \right)^2$$

$$J(\mathbf{W}) = \left\| \mathbf{Y} - \hat{\mathbf{Y}} \right\|^2$$

$$J(\mathbf{W}) = \left\| \mathbf{Y} - \phi(\mathbf{W} \cdot \mathbf{X}^T) \right\|^2$$

$$\mathbf{Y} = \begin{bmatrix} \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_C \end{bmatrix}^{(1)} & \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_C \end{bmatrix}^{(2)} & \dots & \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_C \end{bmatrix}^{(M)} \end{bmatrix} \rightarrow \begin{bmatrix} \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}^{(1)} & \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix}^{(2)} & \dots & \begin{bmatrix} 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix}^{(M)} \end{bmatrix}$$

Each target class and row of \mathbf{W} can be independently optimized




$$\hat{\mathbf{Y}} = \begin{bmatrix} \begin{bmatrix} \phi(\mathbf{w}_{row=1} \cdot \mathbf{x}^{(1)}) \\ \phi(\mathbf{w}_{row=2} \cdot \mathbf{x}^{(1)\cdot}) \\ \vdots \\ \phi(\mathbf{w}_{row=C} \cdot \mathbf{x}^{(1)}) \end{bmatrix}^{(1)} & \begin{bmatrix} \phi(\mathbf{w}_{row=1} \cdot \mathbf{x}^{(2)}) \\ \phi(\mathbf{w}_{row=2} \cdot \mathbf{x}^{(2)\cdot}) \\ \vdots \\ \phi(\mathbf{w}_{row=C} \cdot \mathbf{x}^{(2)}) \end{bmatrix}^{(2)} & \dots & \begin{bmatrix} \phi(\mathbf{w}_{row=1} \cdot \mathbf{x}^{(M)}) \\ \phi(\mathbf{w}_{row=2} \cdot \mathbf{x}^{(M)\cdot}) \\ \vdots \\ \phi(\mathbf{w}_{row=C} \cdot \mathbf{x}^{(M)}) \end{bmatrix}^{(M)} \end{bmatrix}$$

which is one versus-all!

Singel Layer Early Architectures: Summary

- Adaline network, Widrow and Hoff, 1960
 - linear regression, iterative updates
- Perceptron
 - NN *with sigmoid*: logistic regression
- Multi-class perceptron: one versus all
- **But what about when we have more than one layer?**





Math, math...



and look!

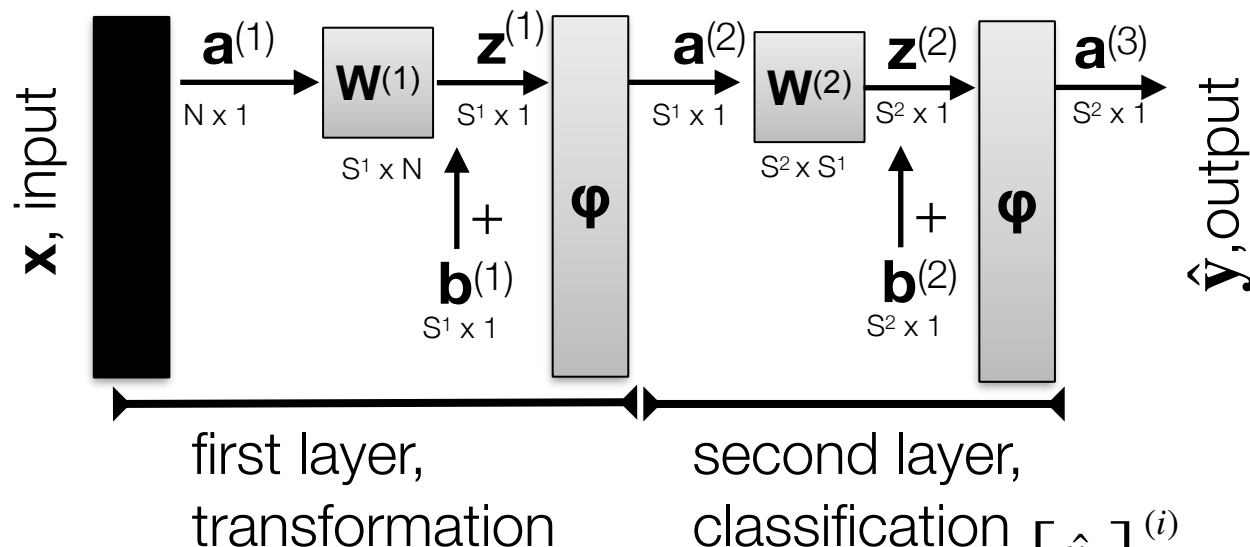


More math...

Beyond Single Layer Networks

Moving to multiple layers...

- The multi-layer perceptron (MLP):
 - two layers shown, but could be arbitrarily many layers



each element of $\hat{\mathbf{y}}$ is no longer independent.

$\mathbf{W}^{(1)}$ used for all classes so we cannot optimize using one versus all 😞

$$\begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \vdots \\ \hat{y}_C \end{bmatrix}^{(i)} = \begin{bmatrix} \phi \left(\phi([\mathbf{z}^{(1)}]^{(i)}) \cdot \mathbf{w}_{\text{row}=1}^{(2)} + b_1^{(2)} \right) \\ \phi \left(\phi([\mathbf{z}^{(1)}]^{(i)}) \cdot \mathbf{w}_{\text{row}=2}^{(2)} + b_2^{(2)} \right) \\ \vdots \\ \phi \left(\phi([\mathbf{z}^{(1)}]^{(i)}) \cdot \mathbf{w}_{\text{row}=C}^{(2)} + b_C^{(2)} \right) \end{bmatrix}$$

$$[\mathbf{z}^{(1)}]^{(i)} = \mathbf{W}^{(1)} \cdot [\mathbf{a}^{(1)}]^{(i)} + \mathbf{b}^{(1)}$$

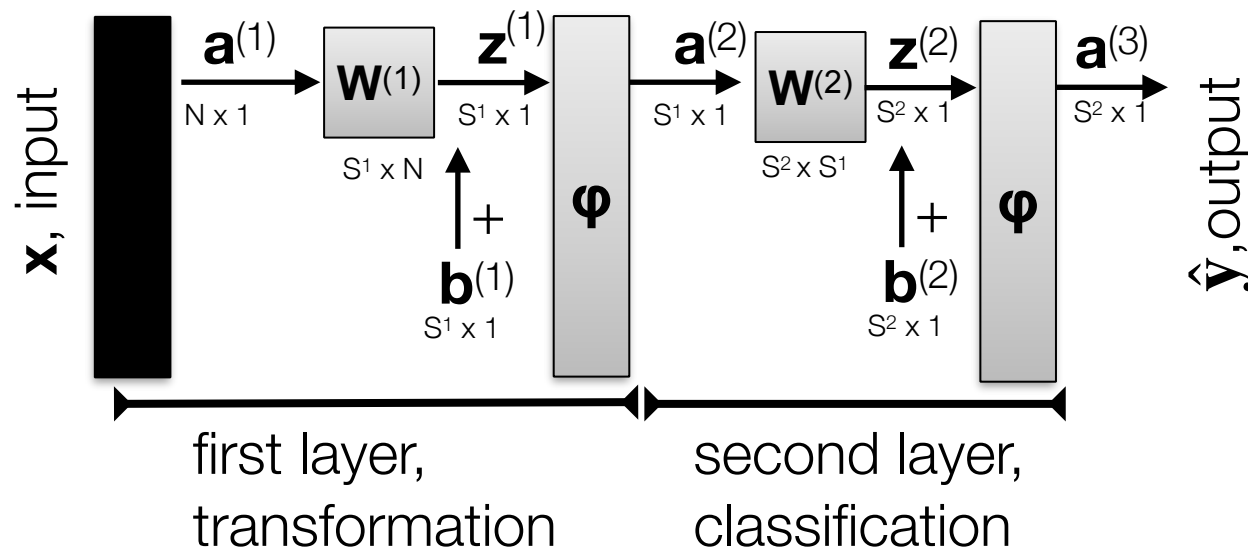


Back propagation

- Optimize all weights of network at once, using chain rule many times...
- Steps:
 - Forward propagate to get all $\mathbf{Z}^{(l)}$, $\mathbf{A}^{(l)}$
 - Get final layer gradient
 - Back propagate sensitivities (chain rule)
 - Update each $\mathbf{W}^{(l)}$



**Back-propagation
is solved in flipped
assignment!!**



$$J(\mathbf{W}) = \| \mathbf{Y} - \hat{\mathbf{Y}} \|^2$$
$$w_{i,j}^{(l)} \leftarrow w_{i,j}^{(l)} - \eta \frac{\partial J(\mathbf{W})}{\partial w_{i,j}^{(l)}}$$

Back propagation

Backprop

$$\bar{z}^{(l)} = W^{(l)} \bar{A}^{(l-1)}$$

$$\bar{A}^{(l)} = \phi'(\bar{z}^{(l)})$$

$$W^{(l)} \leftarrow W^{(l)} + \eta \frac{\partial J(w)}{\partial W^{(l)}}$$

OR

$$w_{ij}^{(l)} \leftarrow w_{ij}^{(l)} + \eta \frac{\partial J(w)}{\partial w_{ij}^{(l)}}$$

THIS WILL BE OUR TEMPORARY VARIABLE

WE AGREE

$$\frac{\partial J}{\partial w_{ij}^{(l)}} = \frac{\partial J}{\partial z_i^{(l)}} \frac{\partial z_i^{(l)}}{\partial w_{ij}^{(l)}}$$

$$= \frac{\partial J}{\partial z_i^{(l)}} \frac{\partial}{\partial w_{ij}^{(l)}} \left(\sum_{k=1}^{S^{(l-1)}} \bar{z}_k^{(l-1)} w_{kj}^{(l)} \right)$$

DUMMY VARIABLE

$z_i^{(l)} =$



You are ready to begin back propagation!