

Lecture Notes for **Machine Learning in Python**



Professor Eric Larson
Neural Network, Wrap Up

Class Logistics and Agenda

- Logistics
 - hiring undergraduates!
 - grading update
- Agenda:
 - Demo and Review
 - Universality
 - Parameter Searching
 - Statistical Testing
 - Lab 4 Town Hall

Class Overview, by topic

Table Data
Visualization

Numpy, Pandas, Seaborn
Overviews with some in-depth discussion

Dimension
Reduction and
Image Processing

Scikit-learn, Scikit Image,
Intuition only, Some mathematics

Linear and
Logistic
Regression

Numpy, Recreate API for Scikit-learn
Detailed mathematics for simple optimization
intuition for advanced optimization

Neural Networks
and Back Prop.

Numpy
Detailed mathematics for NN operations

Wide and Deep
Networks

Convolutional
Networks

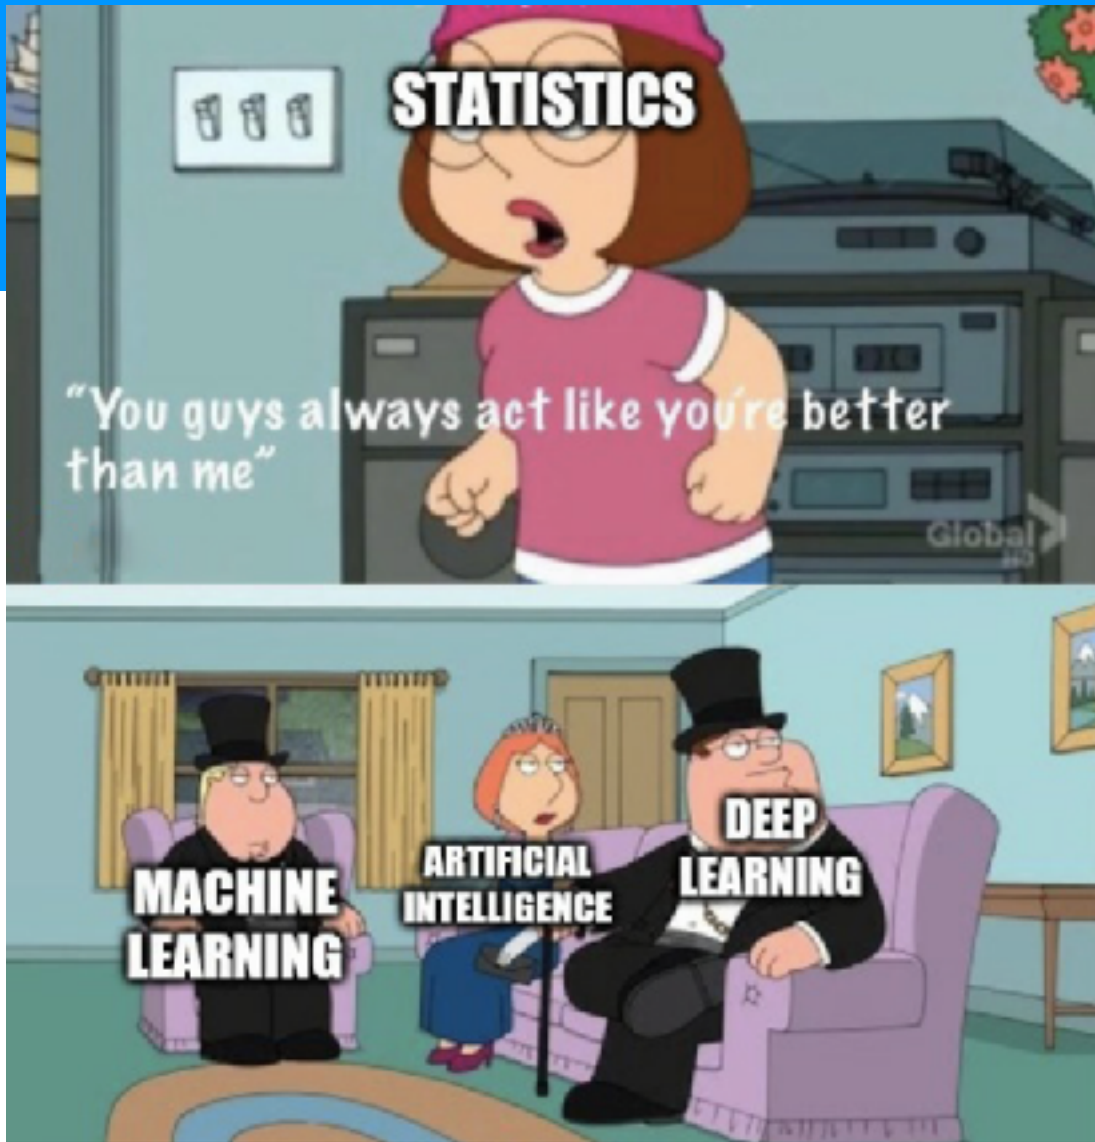
Sequential
Networks

Keras, Tensorflow
Intuition, Detailed implement.

Ethics in
Language Models

ConceptNet
Case studies

Review



Review

$$\mathbf{W}_{k+1} = \mathbf{W}_k - \rho_k$$

- Cross entropy $\mathbf{V}^{(2)} = \mathbf{A}^{(3)} - \mathbf{Y}$
new final layer update

- Momentum $\rho_k = \alpha \nabla J(\mathbf{W}_k) + \beta \nabla J(\mathbf{W}_{k-1})$

- Nesterov's Momentum $\rho_k = \beta \nabla J(\underbrace{\mathbf{W}_k + \alpha \nabla J(\mathbf{W}_{k-1})}_{\text{step twice}}) + \alpha \nabla J(\mathbf{W}_{k-1})$

- Mini-batching

← all data →

	batch 1	batch 2	batch 3	batch 4	batch 5	batch 6	batch 7	batch 8	batch 9
Epoch 1									
Epoch 2									
Epoch 3									
Epoch 4									
...									

shuffle ordering each epoch and update W 's after each batch


- Learning rate adjustment (eta)

Stair step


$$\eta_e = \eta_0 \cdot d^{\lfloor \frac{e}{e_d} \rfloor}$$

e_d epochs between reductions

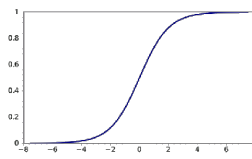
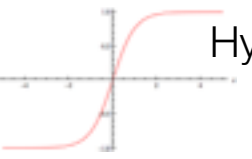
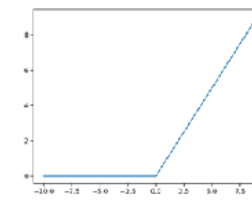
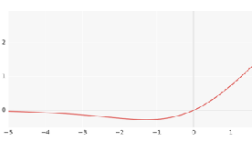
$0 < d < 1$



$$\eta_e = \eta_{min} + \frac{1}{2}(\eta_{max} - \eta_{min}) \left(1 + \cos \left(\frac{e}{e_{max}} \pi \right) \right)$$



Review: Activations Summary

	Definition	Derivative	Weight Init (Uniform Bounds)
 <p>Sigmoid</p>	$\phi(z) = \frac{1}{1 + e^{-z}}$	$\nabla \phi(z) = a(1 - a)$	$w_{ij}^{(L)} \sim \pm 4 \sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$ $b^{(l)} = [-2]$
 <p>Hyperbolic Tangent</p>	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	$\nabla \phi(z) = \frac{4}{(e^z + e^{-z})^2}$	$w_{ij}^{(L)} \sim \pm \sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$ $b^{(l)} = [0]$
 <p>ReLU</p>	$\phi(z) = \begin{cases} z, & \text{if } z > 0 \\ 0, & \text{else} \end{cases}$	$\nabla \phi(z) = \begin{cases} 1, & \text{if } z > 0 \\ 0, & \text{else} \end{cases}$	$w_{ij}^{(L)} \sim \pm \sqrt{2} \sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$
 <p>SiLU</p>	$\phi(z) = \frac{z}{1 + e^{-z}}$	$\nabla \phi(z) = \phi(z) + \sigma(z) \cdot (1 - \phi(z))$	$b^{(l)} = [0]$

Review of Adaptive Strategies $\mathbf{W}_{k+1} = \mathbf{W}_k - \eta \cdot \rho_k$

Adjust each element of gradient by the steepness (for each layer):

- AdaGrad

all operations are per element

$$\rho_k = \frac{1}{\sqrt{\mathbf{G}_k + \epsilon}} \odot \nabla J(\mathbf{W}_k)$$

where

$$\mathbf{G}_k = \gamma \cdot \mathbf{G}_{k-1} + \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k)$$

- RMSProp

all operations are per element

$$\rho_k = \frac{1}{\sqrt{\mathbf{V}_k + \epsilon}} \odot \nabla J(\mathbf{W}_k)$$

$$\mathbf{G}_k = \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k)$$

$$\mathbf{V}_k = \gamma \cdot \mathbf{V}_{k-1} + (1 - \gamma) \cdot \mathbf{G}_k$$

- AdaDelta

all operations are per element

$$\rho_k = \frac{\mathbf{M}_k}{\sqrt{\mathbf{V}_k + \epsilon}}$$

$$\mathbf{M}_{k+1} = \gamma \cdot \mathbf{M}_k + (1 - \gamma) \cdot \nabla J(\mathbf{W}_k)$$

- AdaM

update momentum

$$\mathbf{M}_{k+1} \leftarrow \beta_1 \cdot \mathbf{M}_k + (1 - \beta_1) \cdot \nabla J(\mathbf{W}_k)$$

$$\hat{\mathbf{M}}_k \leftarrow \frac{\mathbf{M}_k}{(1 - [\beta_1]^k)}$$

normalizer momentum

$$\mathbf{V}_{k+1} \leftarrow \beta_2 \cdot \mathbf{V}_k + (1 - \beta_2) \cdot \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k)$$

$$\hat{\mathbf{V}}_k \leftarrow \frac{\mathbf{V}_k}{(1 - [\beta_2]^k)}$$

full update

$$\mathbf{W}_{k+1} \leftarrow \mathbf{W}_k - \eta \cdot \frac{\hat{\mathbf{M}}_k}{\sqrt{\hat{\mathbf{V}}_k + \epsilon}}$$

exploitation, boosting

08a. Practical_NeuralNetsWithBias.ipynb

Momentum

Learning Rate Adaptation

Cross Entropy

Smarter Weight Initialization

Adaptive training with AdaGrad

ReLU Nonlinearities

Sklearn Comparison




Revisiting Universality (if time)


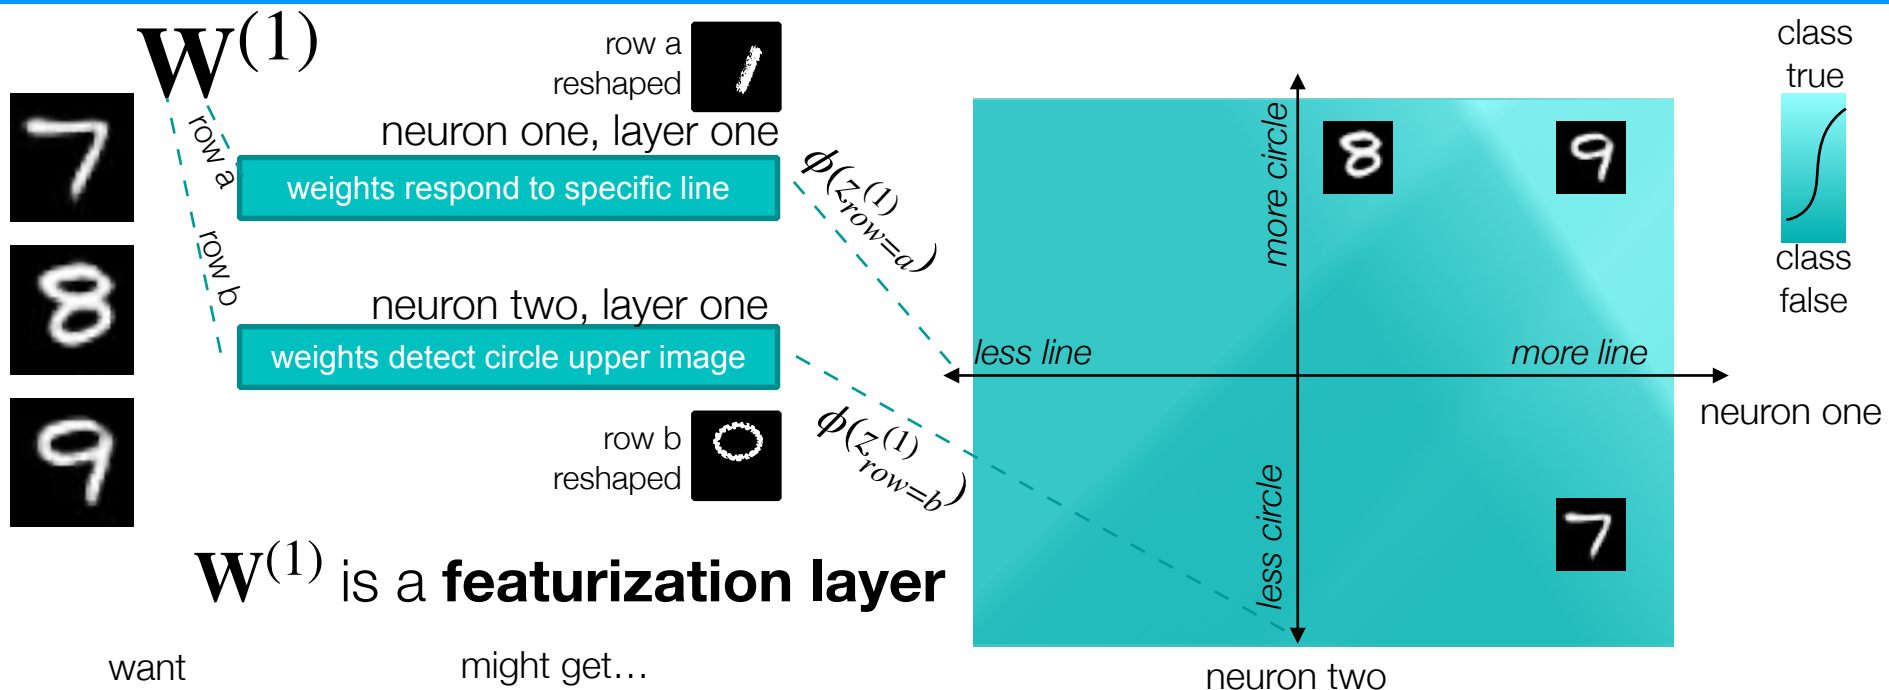
- Neural networks can separate any data through multiple layers. The true realization of Rosenblatt:

"Given an elementary α -perceptron, a stimulus world W , and any classification $C(W)$ for which a solution exists; let all stimuli in W occur in any sequence, provided that each stimulus must reoccur in finite time; then beginning from an arbitrary initial state, an error correction procedure will always yield a solution to $C(W)$ in finite time..."


- **Universality:** No matter what function we want to compute, we know that there is a neural network which can do the job.



Universality




$W^{(2)}$ is **classification layer**, like one versus all logistic regression



- One nonlinear hidden layer with an output layer can **perfectly train any problem with enough data**, but might be memorizing...
- ... could be better to have **even more layers** for more generalizing features

Grid Searching



Review: Grid Searching

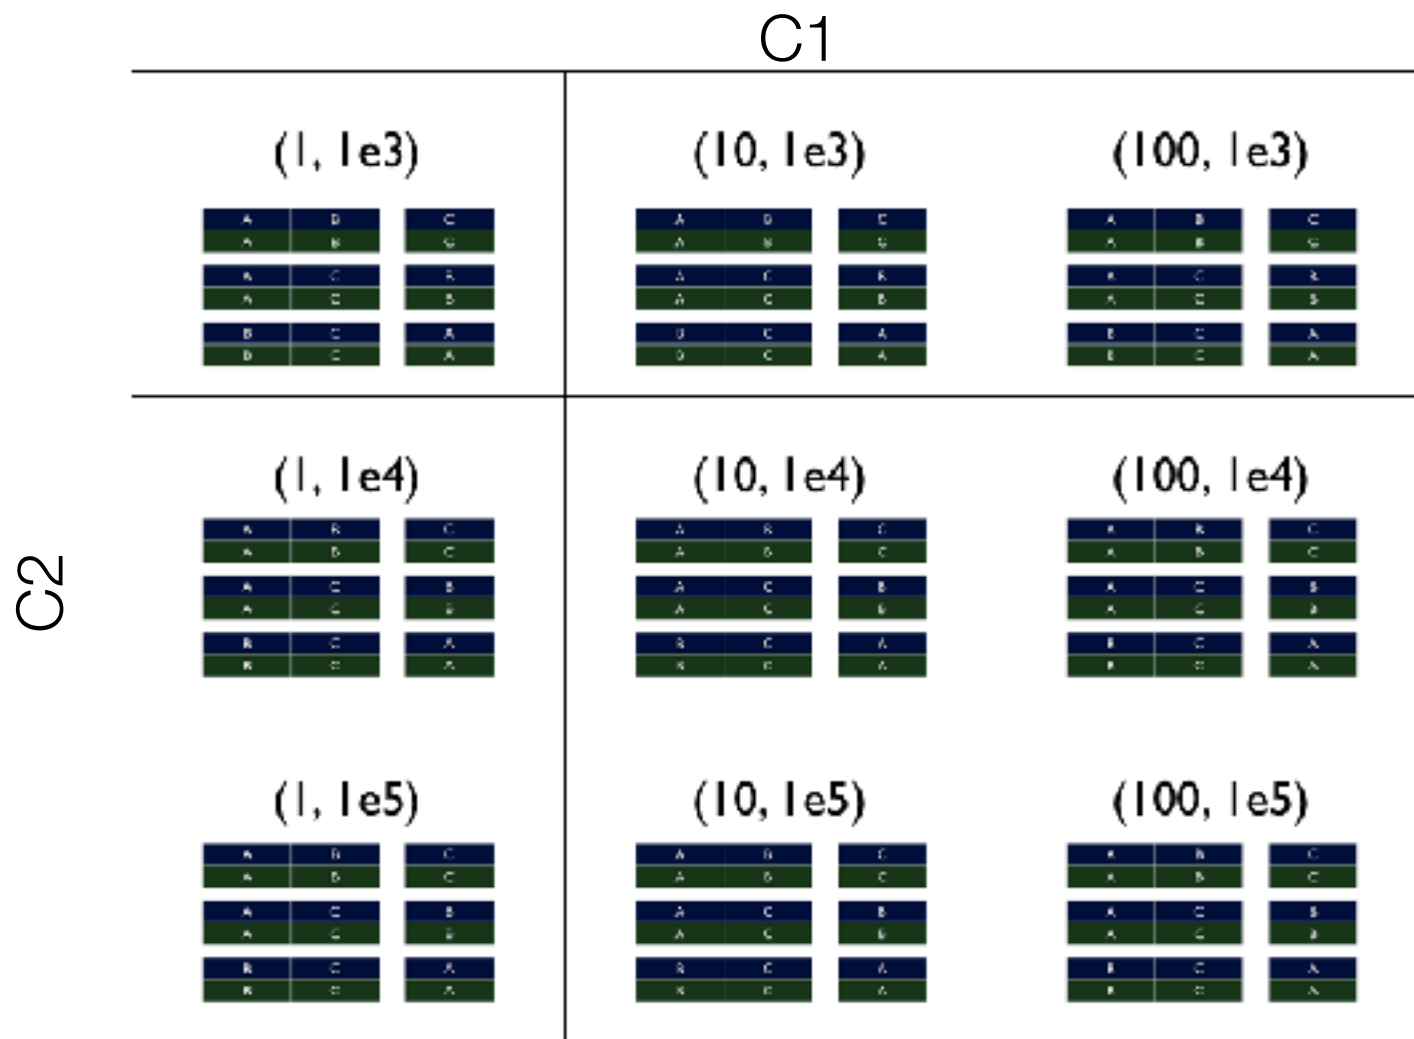
Trying to find the best parameters

NN: C1=[1, 10, 100] C2=[1e3,1e4,1e5]

		C1	
C2	(1, 1e3)	(10, 1e3)	(100, 1e3)
	(1, 1e4)	(10, 1e4)	(100, 1e4)
	(1, 1e5)	(10, 1e5)	(100, 1e5)

Review: Grid Searching

For each value, want to run cross validation...



Review: Grid Searching

Could perform iteratively

C1

(1, 1e3)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

(10, 1e3)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

(100, 1e3)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

(1, 1e4)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

(10, 1e4)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

(100, 1e4)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

(1, 1e5)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

(10, 1e5)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

(100, 1e5)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

C2

Review: Grid Searching

or at random...

C1

(1, 1e3)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

(10, 1e3)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

(100, 1e3)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

(1, 1e4)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

(10, 1e4)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

(100, 1e4)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

(1, 1e5)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A


(10, 1e5)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

(100, 1e5)

A	B	C
A	B	C
A	C	B
A	C	B
B	C	A
B	C	A

If random
can also draw
at random from
a distribution!



C1 drawn from
 $\mathcal{N}(\mu = 50, \sigma = 20)$

Review: Grid Searches in Scikit-learn

```
>>> from sklearn import svm, datasets
>>> from sklearn.model_selection import GridSearchCV
>>> iris = datasets.load_iris()
>>> parameters = {'kernel':('linear', 'rbf'), 'C':[1, 10]}
>>> svc = svm.SVC()
>>> clf = GridSearchCV(svc, parameters)
>>> clf.fit(iris.data, iris.target)
GridSearchCV(estimator=SVC(),
              param_grid={'C': [1, 10], 'kernel': ('linear', 'rbf')})
```



[Key Features](#) [Code Examples](#) [Installation](#) [Blog](#) [Videos](#) [Paper](#) [Community](#)

Optuna is framework agnostic. You can use it with any machine learning or deep learning framework.

[Quick Start](#) [PyTorch](#) [PyTorch](#) [Chainer](#) [TensorFlow](#) [Keras](#) [MXNet](#) [Scikit-Learn](#) [XGBoost](#) [LightGBM](#)

values, sampled

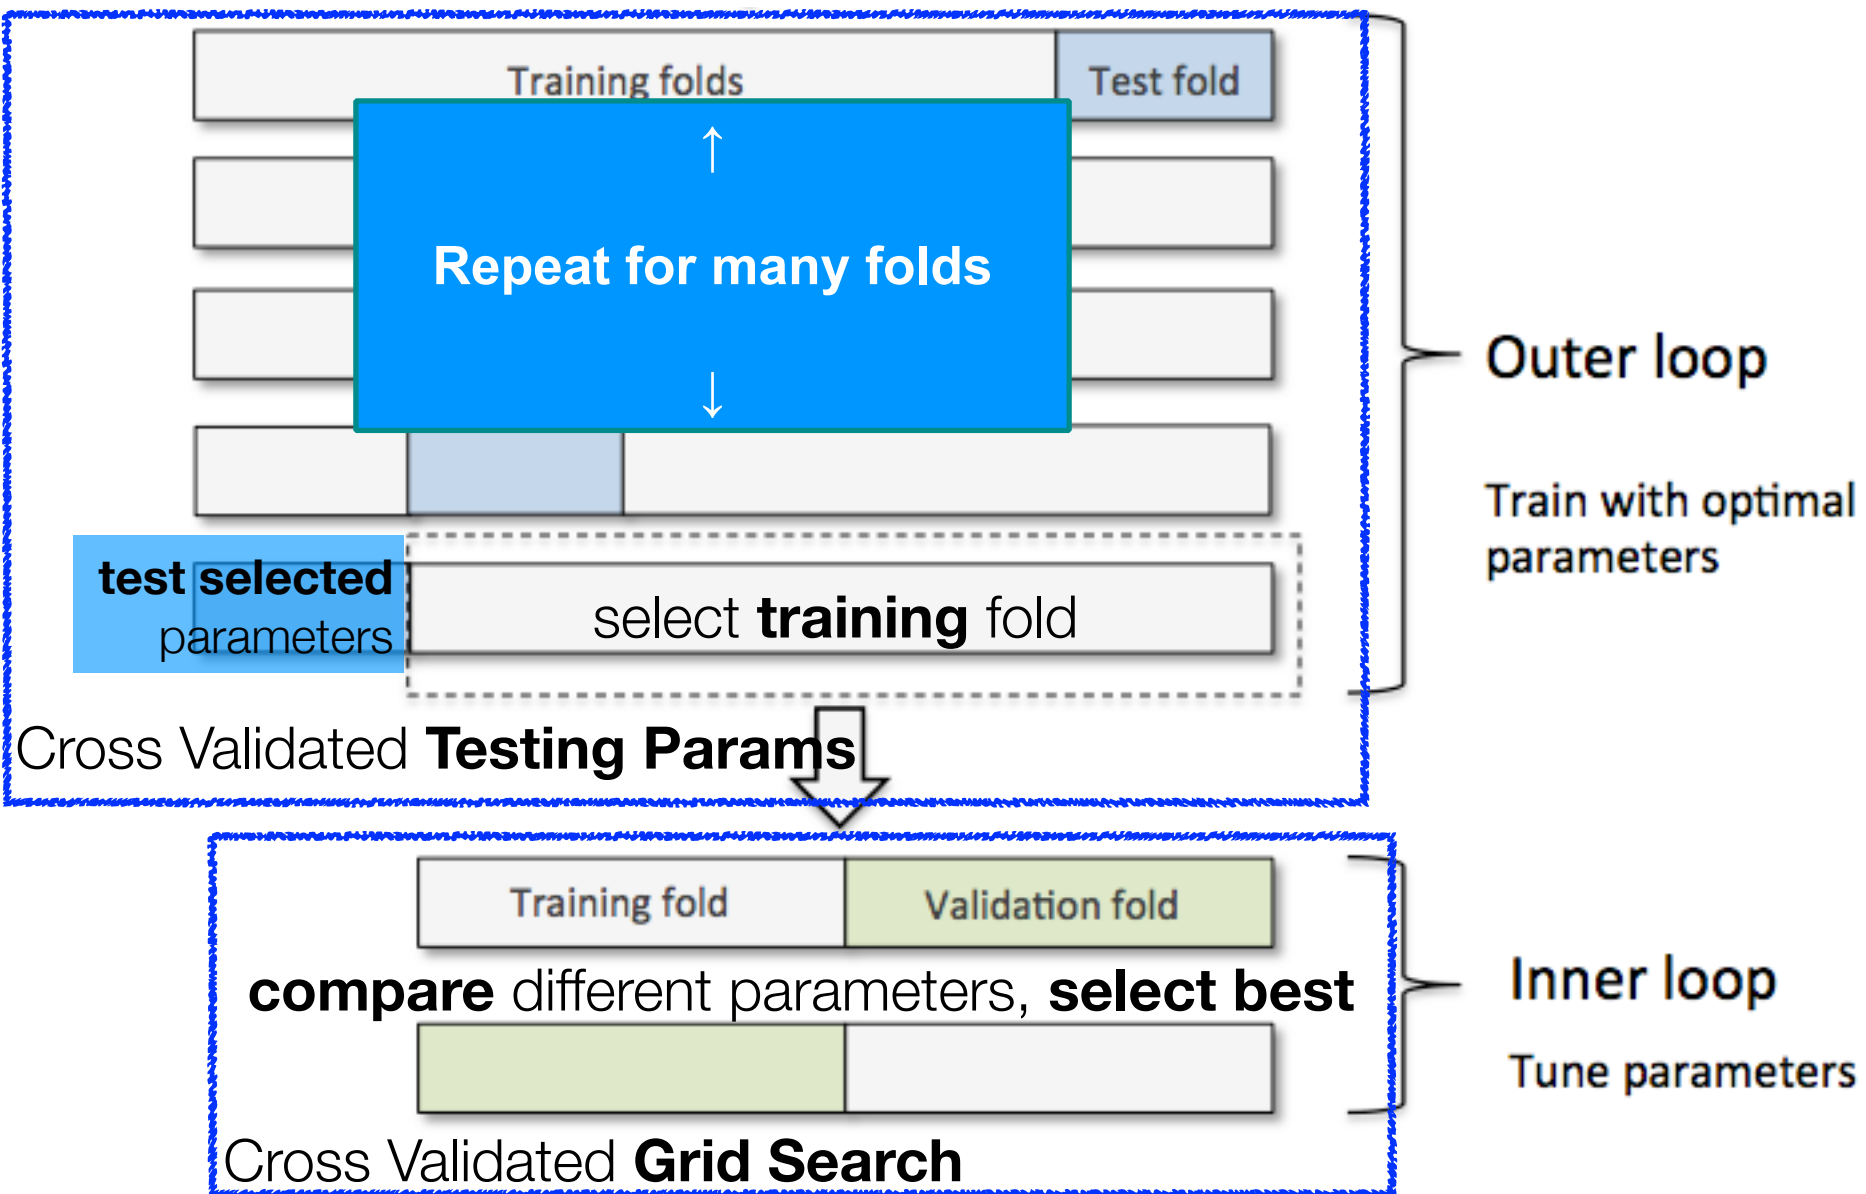
```
>>> from sklearn.linear_model import LogisticRegression
>>> from sklearn.model_selection import RandomizedSearchCV
>>> from scipy.stats import uniform
>>> iris = load_iris()
>>> logistic = LogisticRegression(solver='saga', tol=1e-2, max_iter=200,
                                random_state=0)
>>> distributions = dict(C=uniform(loc=0, scale=4),
                        ... penalty=['l2', 'l1'])
>>> clf = RandomizedSearchCV(logistic, distributions, random_state=0)
>>> search = clf.fit(iris.data, iris.target)
>>> search.best_params_
{'C': 2..., 'penalty': 'l1'}
```


Review: Self Test

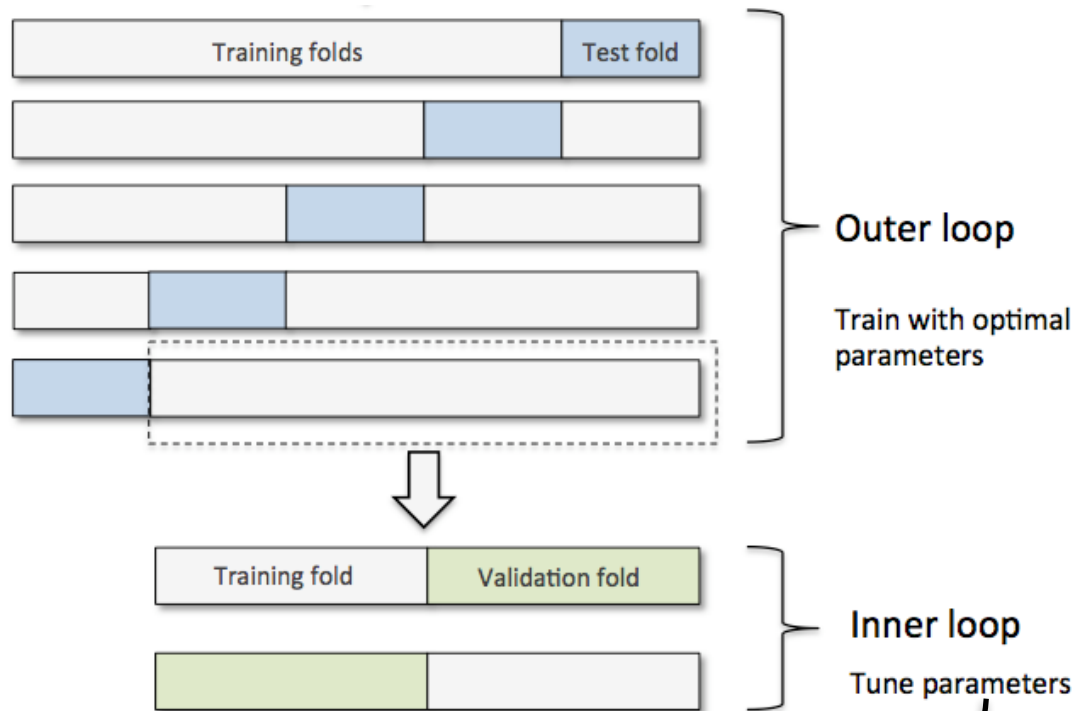
- Using the grid search parameters and testing on the same set...
- Is this **data snooping**?
 - **A. True**, this is snooping because it uses test set to define parameters
 - **B. True**, this is snooping because we can no longer reliably define the expected performance on new data

How can we define expected performance when using cross validation in a grid search?

Review: Nested Cross Validation



Review: Nested Cross Validation: Hyper-parameters




```
gs = GridSearchCV(estimator=pipe_svc,  
                  param_grid=param_grid,  
                  scoring='accuracy',  
                  cv=2)  
  
scores = cross_val_score(gs, X_train, y_train, scoring='accuracy', cv=5)  
print('CV accuracy: %.3f +/- %.3f' % (np.mean(scores), np.std(scores)))
```

Self Test

- **What is the end goal of nested cross-validation?**
 - A. To determine hyper parameters
 - B. To estimate generalization performance
 - C. To estimate generalization performance when performing hyper parameter tuning
 - D. To estimate the variation in tuned hyper parameters

Statistical Tests



Flipped Module, Model Comparison

- How do we compare two (or more) trained models to one another?
- Are models different enough to prefer one model over another?

Comparing Performance of 2 Models

- Given two models, M_1 and M_2 , which is better?
 - M_1 is tested on D_1 (size= n_1), found error rate = e_1
 - M_2 is tested on D_2 (size= n_2), found error rate = e_2
 - Assume D_1 and D_2 are independent
 - If n_1 and n_2 are sufficiently large, then

$$e_1 \sim N(\mu_1, \sigma_1^2)$$

$$e_2 \sim N(\mu_2, \sigma_2^2)$$

- Approximate:
$$\hat{\sigma}_i^2 = \frac{e_i(1 - e_i)}{n_i}$$

variance estimate comes from **binomial distrib**
which is approximated well by **normal distrib**

Comparing Performance of 2 Models

- To test if performance difference is statistically significant: $d = e_1 - e_2$ ← estimate of the mean difference
 - $d \sim N(d_0, \alpha)$ where d_0 is the true difference
 - Since D_1 and D_2 are independent, their variance adds up:

$$\sigma_d^2 = \sigma_{e_1}^2 + \sigma_{e_2}^2 = \frac{e_1(1 - e_1)}{n_1} + \frac{e_2(1 - e_2)}{n_2}$$

Folded statistical comparisons

- Each learning algorithm may produce k models:
 - L_1 may produce $M_{11}, M_{12}, \dots, M_{1k}$
 - L_2 may produce $M_{21}, M_{22}, \dots, M_{2k}$
- If models are generated on the same test sets D_1, D_2, \dots, D_k (e.g., via cross-validation)

- For each set: compute $d_j = e_{1j} - e_{2j}$, the j th difference
- d_j has mean d and variance α

$$\sigma_d^2 = \frac{1}{k-1} \sum_{j=1}^k (d_j - \bar{d})^2$$

$$d_i = \bar{d} \pm \frac{1}{\sqrt{k}} t_{1-\alpha, k-1} \sigma_d$$

now we can bound to
← get a better idea about
how the criterion varies

McNemar Testing for Comparing Performance

Few assumptions, **Null hypothesis**: predictions are not different!

McNemar and Edwards, 1948

	Model 2 correct	Model 2 wrong
Model 1 correct	A	B
Model 1 wrong	C	D

$$\chi^2 \approx \frac{(|B - C| - 1)^2}{B + C}$$

If predictions are drawn from the same distributions, then this equation follows χ **squared statistic with one DOF**

Steps:

1. Compare each model's predictions on **the same test data** (2x2 matrix)
2. Calculate χ^2 statistic
3. Look up *critical value* associated with χ^2 statistic for given confidence
4. Are you confident enough to **reject the null hypothesis** that the performance is the same ($p < 0.05$)?

One caveat: Statistical power depends upon $B+C$, which might be small, even with lots of test data.

McNemar Example

Model 1	Model 2	Label	Matrix
T-shirt	T-shirt	T-shirt	A
Sneaker	T-shirt	Sneaker	B
T-shirt	Pullover	Pullover	C
Sneaker	Sneaker	Sneaker	A
T-shirt	Sneaker	Sneaker	C
Pullover	Pullover	T-shirt	D
Pullover	T-shirt	Pullover	B
Sneaker	Sneaker	Sneaker	A
Sneaker	Sneaker	Sneaker	A

	Model 2 correct	Model 2 wrong
Model 1 correct	4 ^A	2 ^B
Model 1 wrong	2 ^C	1 ^D

McNemar and Edwards, 1948

$$\chi^2 \approx \frac{(|B - C| - 1)^2}{B + C}$$

$$\chi^2 = \frac{(|2 - 2| - 1)^2}{2 + 2} = 0.25$$

Confidence	0.90	0.95	0.99
1 DOF, Critical Value	2.706	3.841	6.635

<https://www.itl.nist.gov/div898/handbook/eda/section3/eda3674.htm>

Since $0.25 < 3.841$, we cannot reject the null hypothesis. This means **we should not say the models' performance are different** based on the evidence.

Self Test:

- You have trained **three different models** on prediction of child poverty ratings. Each model is trained on the same data and **tested on the same single split** of the data.
- **Can you use a McNemar test for selecting a model?**
 - **A. Yes.** McNemar testing can be used for any testing of any model without any assumptions.
 - **B. Yes.** McNemar testing can be applied pairwise for the three models.
 - **C. No.** McNemar testing cannot be used for more than two models.
 - **D. No.** McNemar testing can only tell if the models are different, we still cannot tell which one is best

Town Hall



Next Time:
Deep Learning