


# Lecture Notes for **Machine Learning in Python**




Professor Eric Larson  
**Tensorflow Deep Dive**

# Lecture Agenda

- Logistics:
  - CS 5/7325 in Spring
  - Grading update
- Agenda:
  - More Introduction to TensorFlow
    - Tensors, Tf.Data
    - Deep APIs
  - Wide and Deep Networks
  - Town Hall

# Class Overview, by topic



# Last Time

- Up to this point: back propagation saved AI winter for NN (Hinton and others!)
- 80's, 90's, 2000's: convolutional networks for image processing start to get deeper
  - but back propagation no longer does great job at training them
- SVMs and Random Forests gain traction...
  - The second AI winter begins, research in NN plummets
- 2004: Hinton secures funding from CIFAR in 2004 Hinton rebrands: Deep Learning
- 2006: Auto-encoding and Restricted Boltzmann Machines
- 2007: Deep networks are more efficient when pre-trained
- 2009: GPUs decrease training time by 70 fold...
- 2010: Hinton's students go to internships with Microsoft, Google, and IBM, making their speech recognition systems faster, more accurate and deployed in only 3 months...
- 2012: Hinton Lab, Google, IBM, and Microsoft jointly publish paper, popularity sky-rockets for deep learning methods
- 2011-2013: Ng and Google run unsupervised feature creation on YouTube videos (becomes computer vision benchmark)
- 2012+: Pre-training is not actually needed, just solutions for vanishing gradients (like ReLU, SiLU, initializations, more data, GPUs)



1949, Hebb's Law  
Close neuron fire together



1960, Widrow & Hoff  
Adaline Network



1969, Minsky & Papert  
Linear Models are Doomed



1986, Rumelhart & Hinton  
Back-propagation



2003, Vapnik  
Kernel SVMs



2012, Hinton, Fei-Fei Li  
CNNs win ImageNet



2015, Google  
Tensorflow Open Source

1940

1960

1980

2000

2020

Period of Discovery

First AI Winter

Golden Age of NN

2nd AI Winter

Age of Deep Learning

1943, McCulloch & Pitts  
Logic Gates of The Mind

1957, Rosenblatt  
Perceptron

1969, Minsky & Papert  
Linear Models are Doomed

2001, Breiman  
Random Forests


2011, Bengio  
Init and ReLU

2015, Google  
Tensorflow Open Source



TensorFlow

# Last Time



```

Input = tf.placeholder() # size is 28x28
Input = tf.reshape(Input, [784,1])
classes = tf.placeholder()

W_sm = tf.Variable(...)
b_sm = tf.Variable(...)
W_hl = tf.Variable(...)
b_hl = tf.Variable(...)

trainable_variables =
    [W_sm,b_sm,W_hl,b_hl]

def model_forward(Input):
    A_hl = tf.relu( tf.matmul(Input,W_hl) + b_hl )
    A_sm = tf.matmul(A_hl,W_sm) + b_sm
    return A_sm

y_pr = tf.softmax(A_sm)
loss = tf.sparse_softmax_cross_entropy_with_logits

opt = tf.train.SGDOptimizer(learning_rate=0.01)

for features, labels in train_data:
    with tf.GradientTape() as tape:
        yhat = model_forward(features)
        loss_val = loss(labels, yhat)
    grads = tape.gradient(loss_val, trainable_variables)
    opt.apply_gradients(zip(grads, trainable_variables))

```

# Self Test

- The computation graph in tensorflow:
  - A. Can run one operation at a time
  - B. Can be “compiled” for efficiency
  - C. Has one input path (e.g., features) and one output path (e.g., predictions)
  - D. All of the Above

# Using Keras and Tensorflow



# Keras Programming Interfaces

```
from tensorflow import keras
```

- **Keras Sequential API**

- great for simple, feed forward models

- **Keras Functional API**

- build models through series of nested functions
- each “function” represents an operation in the NN

- **Keras Classes (Inheritance)**

- good for more advanced functionality

Primary ML software tool used by top-5 teams on Kaggle in each competition (n=120)



```
my_model = Sequential([
    Dense(100, activation='sigmoid'),
    Dense(10)
])
```

```
x_in = Input(shape=(1,))
```

```
x = Dense(100, activation='sigmoid')(x_in)
x = Dense(10)(x)
```


```
my_model = Model(inputs = x_in,
                  outputs = x )
```

```
class CustomModel(keras.Model):
    def __init__(self, num_classes=10):
        super().__init__()
        self.dense1 = Dense(100, activation='sigmoid')
        self.output_layer = Dense(num_classes)
```

```
def call(self, inputs):
    x = self.dense1(inputs)
    return self.output_layer(x)
```

```
my_model = CustomModel(10)
```

# Using Keras and Tensorflow



- **Computation Graph:**
    - Compiled models use optimized CG
    - GradientTape uses Eager execution by Default

Revisiting the MLP with Keras




10a. Keras Wide and Deep as TFData.ipynb

Make me slow down if I go too fast!!


# Using Keras and Tensorflow

Option B




## FeatureSpace Methods:

```
integer_categorical  
integer_hashed  
  
string_categorical  
string_hashed  
  
float  
float_discretized  
float_normalized  
float_rescaled
```



# Keras FeatureSpaces




Setting up Feature Spaces




10a. Keras Wide and Deep as TFData.ipynb

# Categorical Feature Embeddings

- One hot encoded data can be made dense through a matrix multiplication,  $\mathbf{a} = \mathbf{W}^{(emb)} \cdot \mathbf{x}_{OHE}$



# Using Embeddings in Keras



Adding Embedding Branches



10a. Keras Wide and Deep as TFData.ipynb