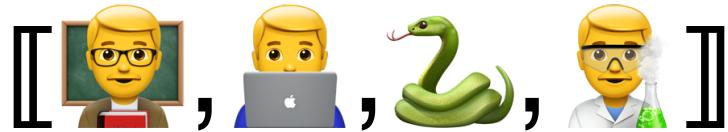


# Lecture Notes for Machine Learning in Python

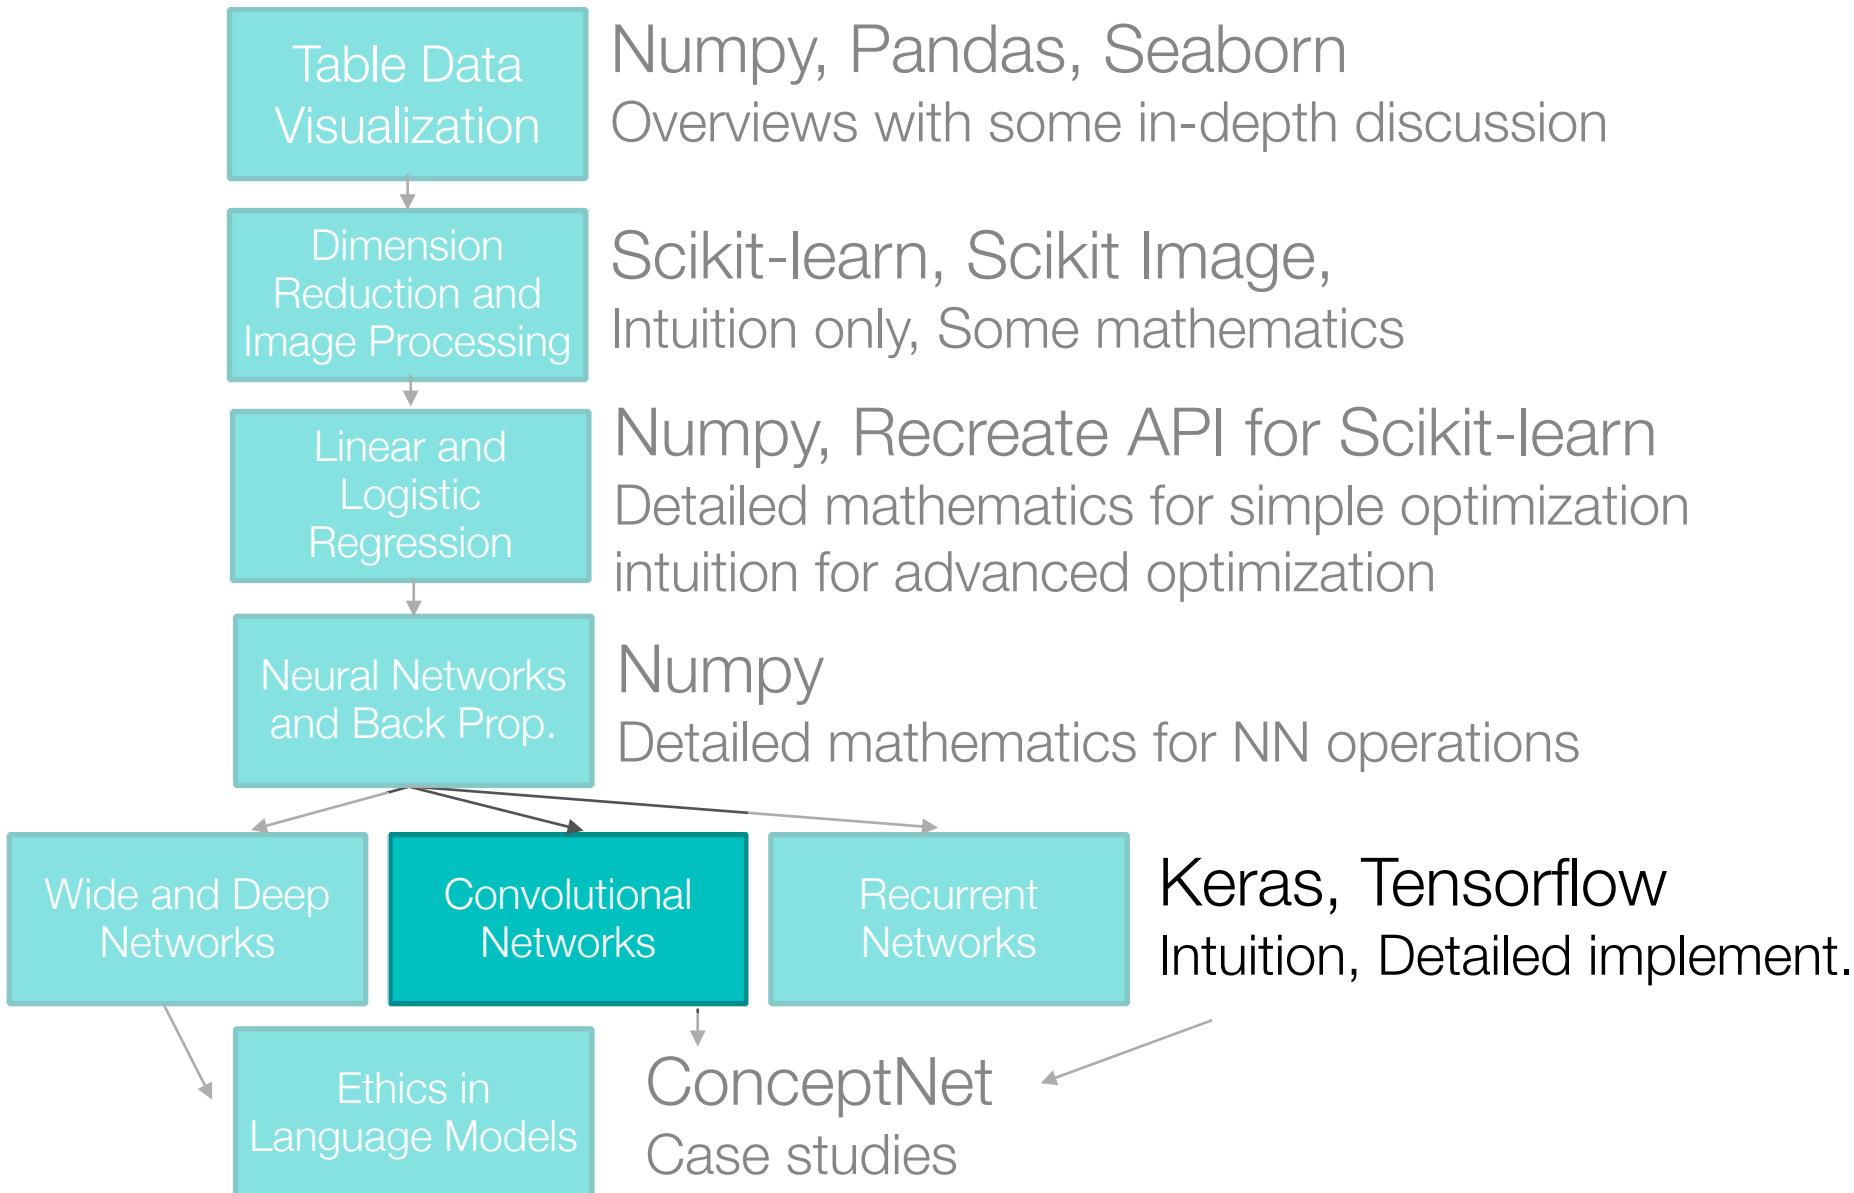


Professor Eric Larson  
**Practical Introductory CNNs**

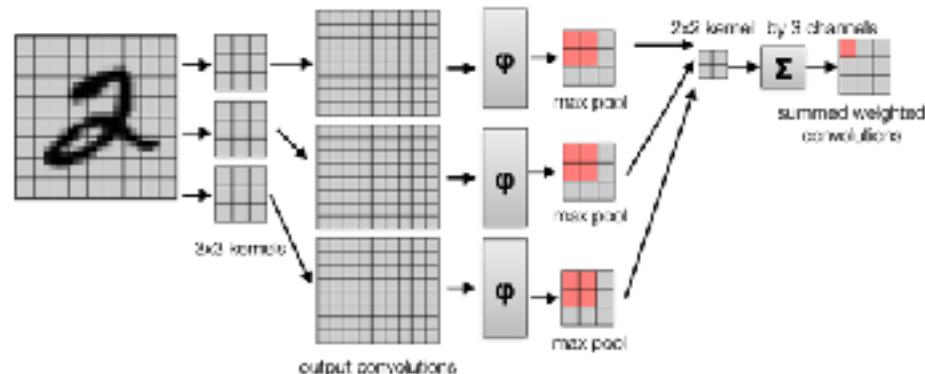
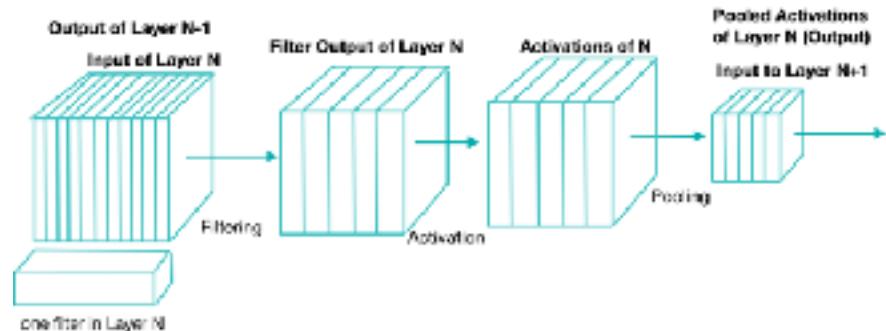
# Class logistics and Agenda

- Wide/Deep Lab (late turn in)
- Agenda:
  - History of CNNs
    - with Modern CNN Architectures
    - and some built-in demos
- Next Time:
  - Finish this lecture
  - and move on to Transformers

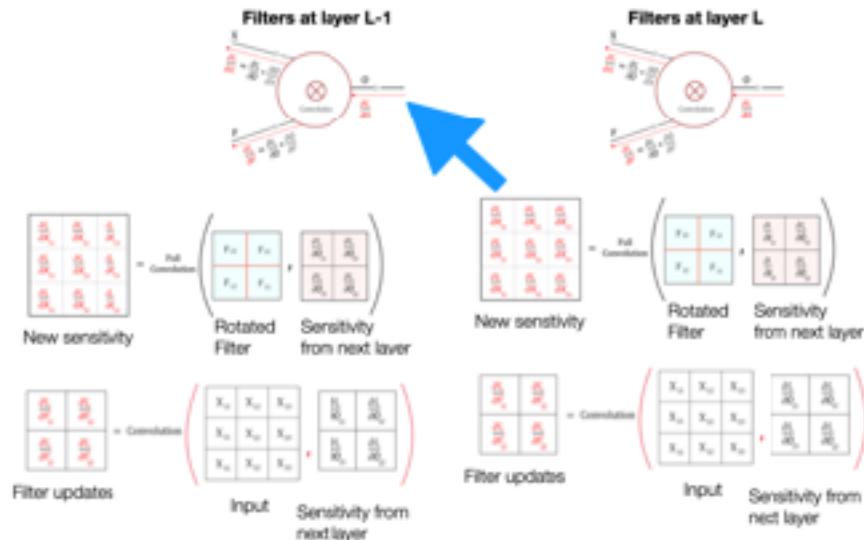
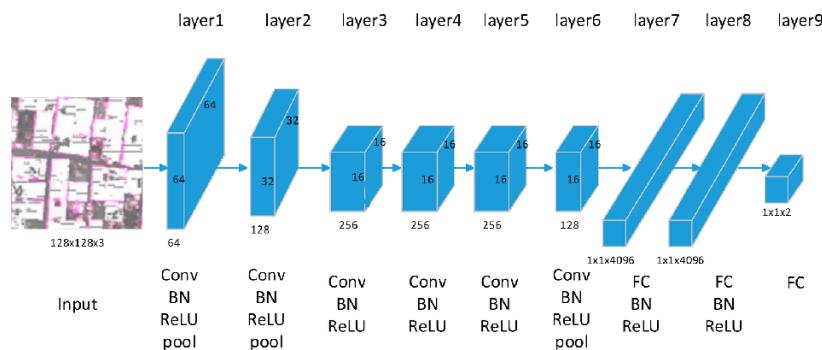
# Class Overview, by topic



# Last Time:



**Structure of Each Tensor:** Channels x Rows x Columns



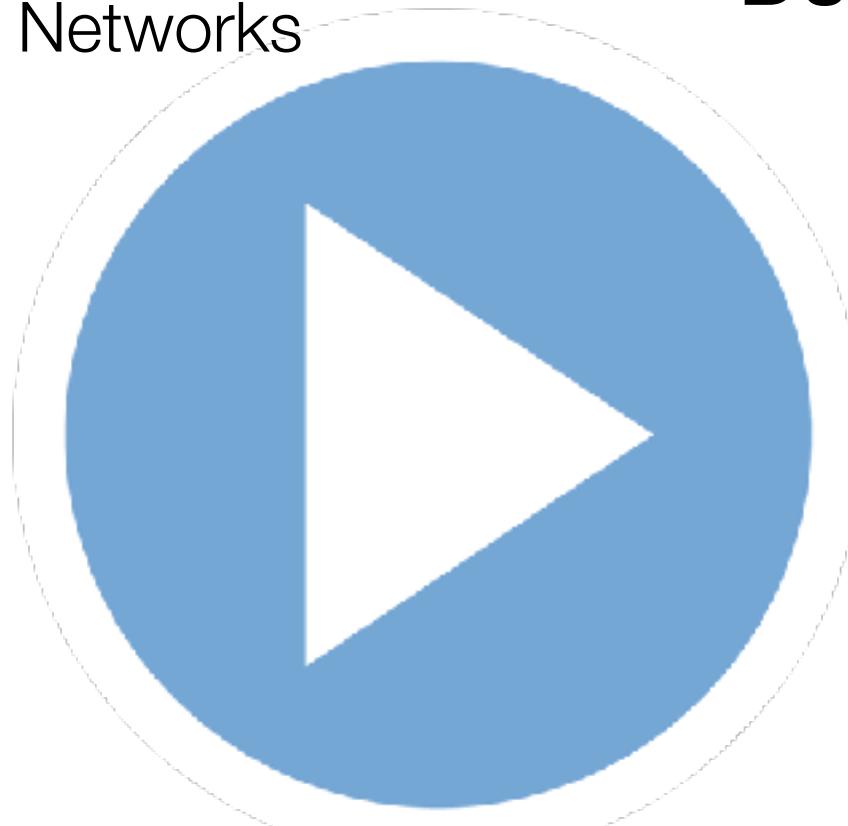
# TensorFlow and Basic CNNs

**Last  
Time!**

Convolutional Neural Networks  
in TensorFlow  
with Keras

with Sequential API!

If needed:  
**Finish  
Demo**



**11. Convolutional Neural Networks.ipynb**

# History of Convolutional Neural Networks



Machine Learning 101

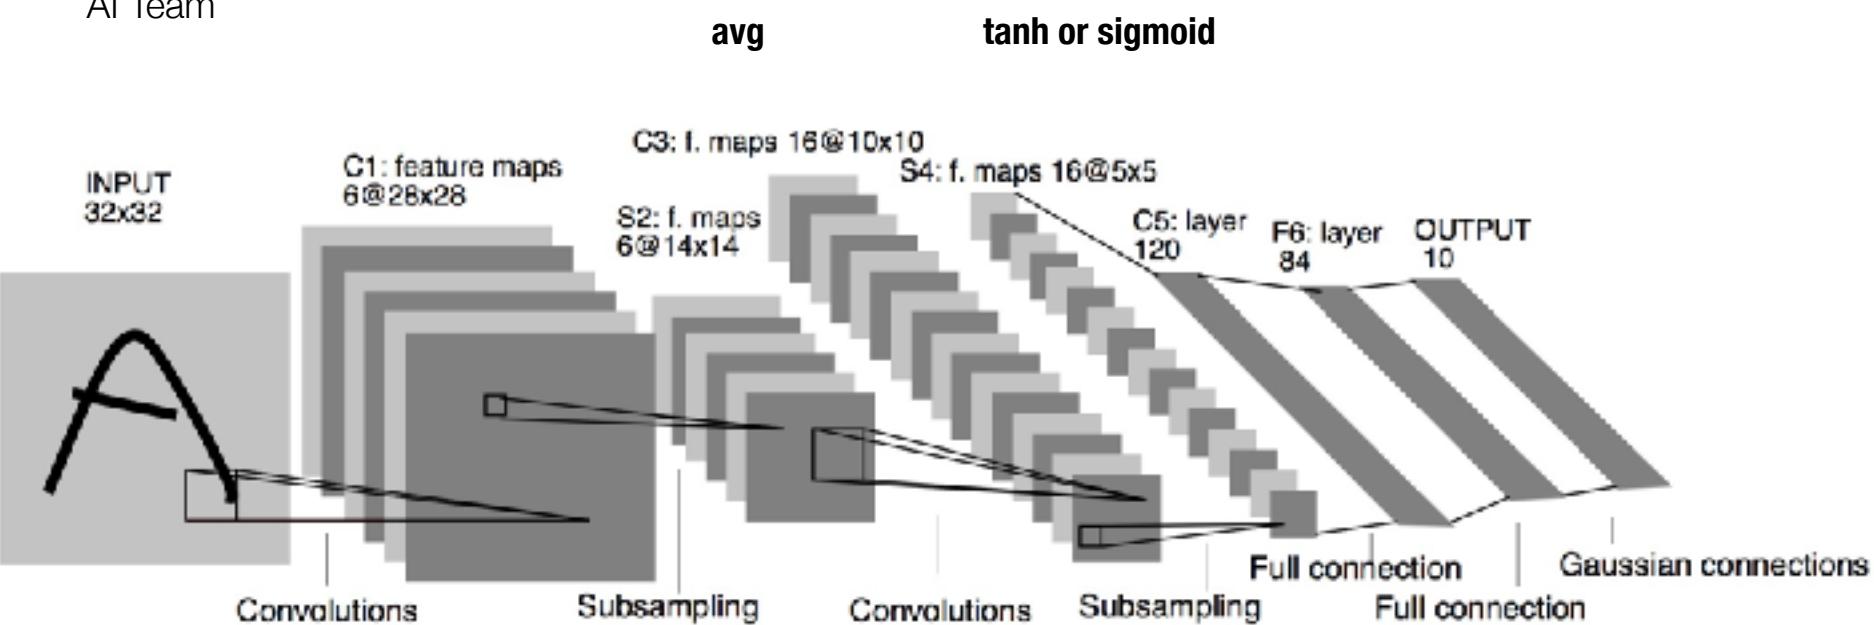
# Types of CNN, 1988-1998

- **LeNet-1** (1988)
  - ~2600 params, not many layers
- **LeNet-5** (1998)
  - 7 layers, gets excellent MNIST performance
- Major contribution, general structure:
  - conv=>pool=>non-linearity=> ...=>MLP



Yann  
LeCun

Heads Facebook  
AI Team



# CNN History

- List of major breakthroughs from 1998 through 2010 for convolutional networks:



- 2010



# Types of CNN, 2010



Dan  
Ciresan

AI Researcher  
IDSA, Switzerland

- **Ciresan Net**
- Publishes code for running CNN via GPU
  - Subsequently wins 5 international competitions
    - from stop signs => cancer detection
- Major contribution: NVIDIA parallelized training algorithms

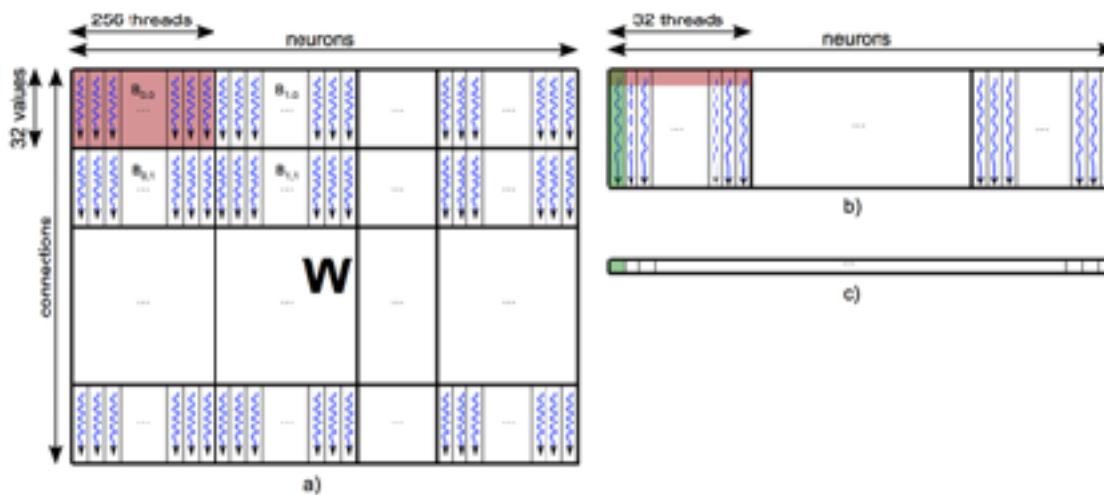
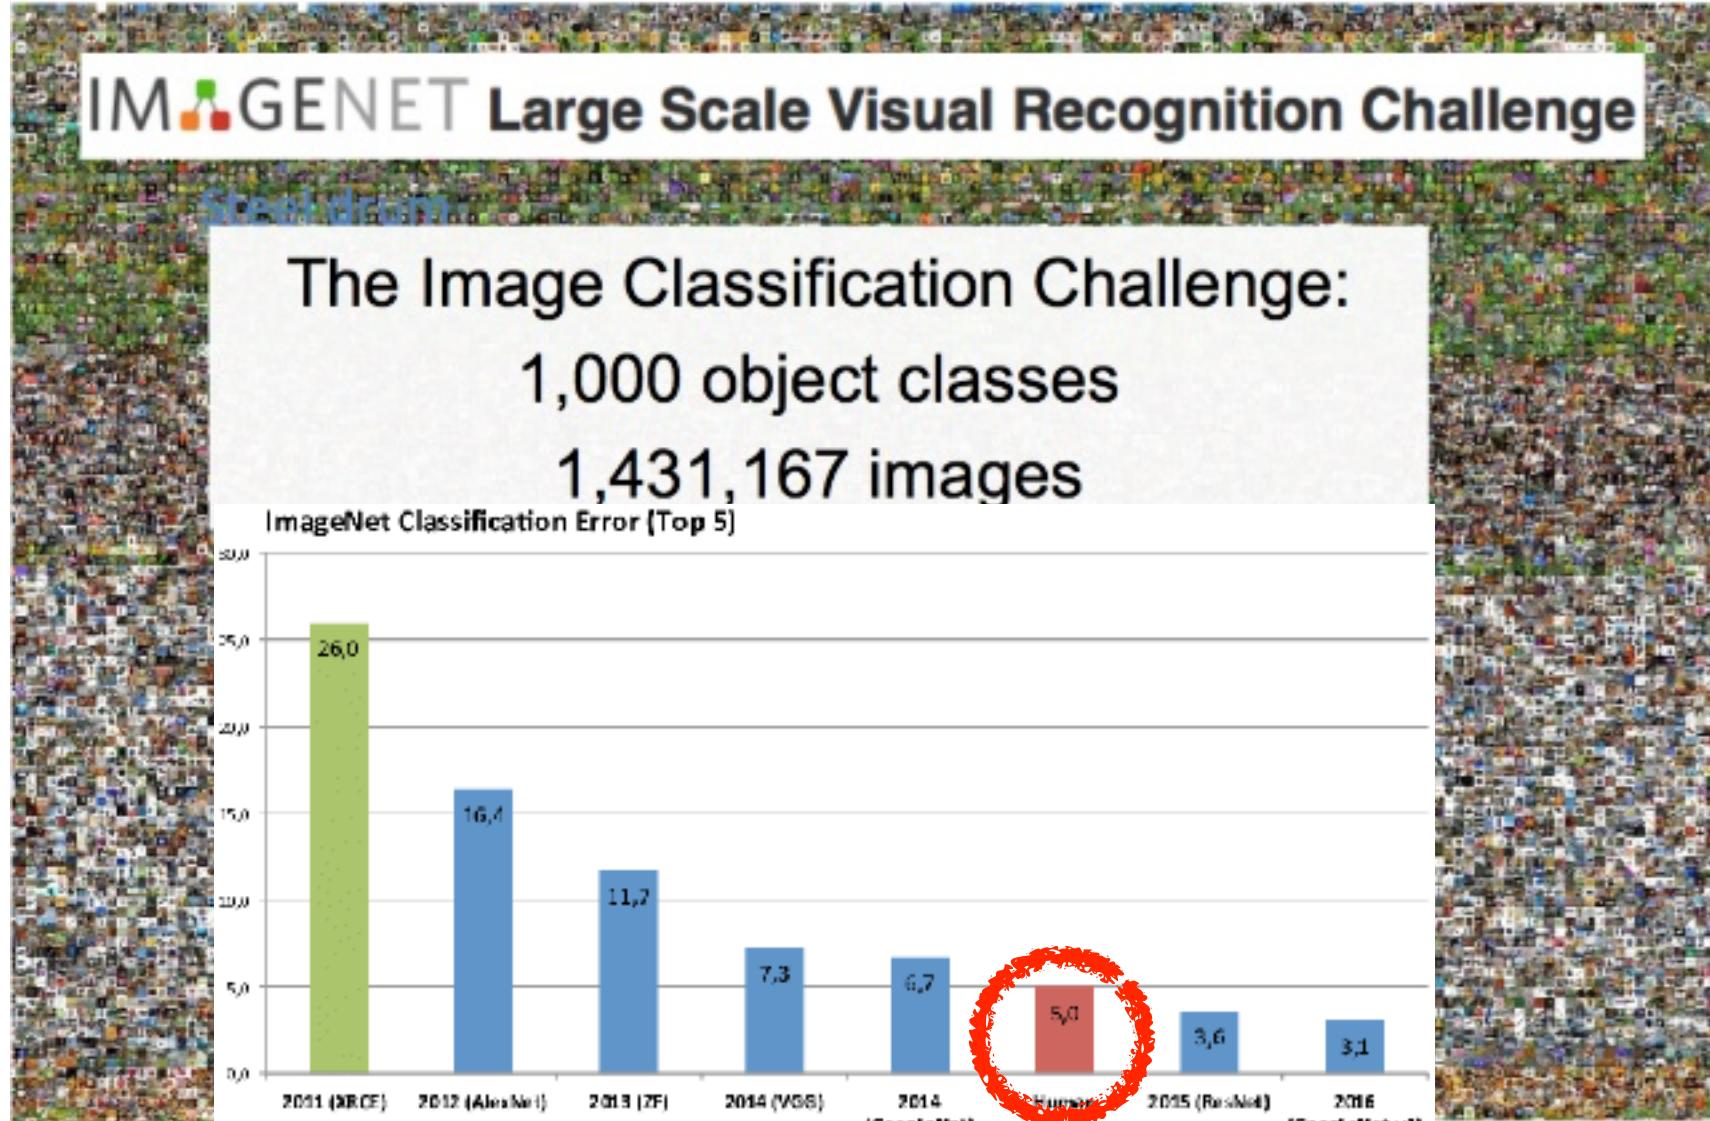


Figure 2: Forward propagation: a) mapping of kernel 1 grid onto the padded weight matrix; b) mapping the kernel 2 grid onto the partial dot products matrix; c) output of forward propagation.

<https://arxiv.org/pdf/1003.0358.pdf>

# ImageNet Competition (2010-2016)



[https://www.researchgate.net/figure/Winner-results-of-the-ImageNet-large-scale-visual-recognition-challenge-LSVRC-of-the\\_fig7\\_324476862](https://www.researchgate.net/figure/Winner-results-of-the-ImageNet-large-scale-visual-recognition-challenge-LSVRC-of-the_fig7_324476862)

<https://www.slideshare.net/nmhkahn/case-study-of-convolutional-neural-network-61556303>

# Types of CNN, 2012

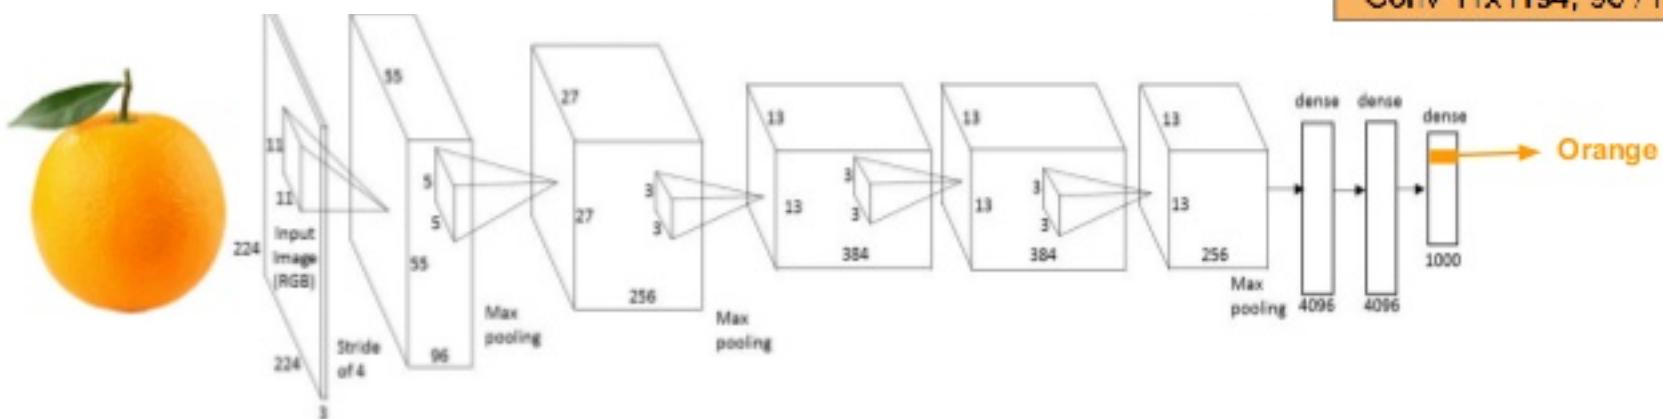
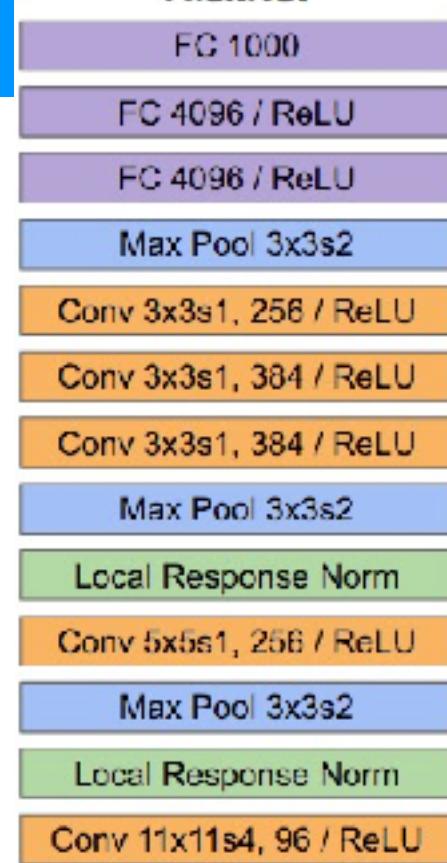
AlexNet



Alex  
Krizhevsky

Google

- **AlexNet**, Hinton is mentor
  - wins ImageNet competition
- Major contributions:
  - dropout for regularization
  - systematic use of ReLU
  - data expansion
  - ***overlapping max pool***



# AlexNet, in Code

```

cnn2 = Sequential(name='AlexNet_Aug')

cnn2.add( Input([64,64,3]) )

# add in augmentations directly
cnn2.add( RandomFlip("horizontal") ) # flip horizontally
cnn2.add( RandomRotation(0.05) ) # rotate by 5%
cnn2.add( RandomTranslation(height_factor=0.1, width_factor=0.1) )
cnn2.add( RandomBrightness(factor=0.1, value_range=(0.0, 1.0)) ) #
cnn2.add( RandomContrast(0.1) ) # add or decrease contrast

cnn2.add(Conv2D(filters=96,
                kernel_size=(11,11),
                strides=(2,2), # add some initial downsampling
                padding='same',
                kernel_initializer='he_uniform',
                kernel_regularizer=l2(1e-6),
                activation='relu')) # more compact syntax

    ...

cnn2.add(Conv2D(filters=256,
                kernel_size=(5,5),
                padding='same',
                kernel_initializer='he_uniform',
                kernel_regularizer=l2(1e-5),
                activation='relu')) # more compact syntax
cnn2.add(MaxPooling2D(pool_size=(2, 2)))

cnn2.add(Conv2D(filters=256,
                kernel_size=(3,3),
                padding='same',
                kernel_initializer='he_uniform',
                kernel_regularizer=l2(1e-5),
                activation='relu')) # more compact syntax
cnn2.add(MaxPooling2D(pool_size=(3, 3)))

```

AlexNet
FC 1000
FC 4096 / ReLU
FC 4096 / ReLU
Max Pool 3x3s2
Conv 3x3s1, 256 / ReLU
Conv 3x3s1, 384 / ReLU
Conv 3x3s1, 384 / ReLU
Max Pool 3x3s2
Local Response Norm
Conv 5x5s1, 256 / ReLU
Max Pool 3x3s2
Local Response Norm
Conv 11x11s4, 96 / ReLU
# add one layer on flattened output
cnn2.add(Flatten())
cnn2.add(Dense(1024, activation='relu', kernel_initializer='he_uniform', kernel_regularizer=l2(1e-4)))
cnn2.add(Dense(512, activation='relu', kernel_initializer='he_uniform', kernel_regularizer=l2(1e-4)))
cnn2.add(Dropout(0.5)) # add some dropout for regularization
cnn2.add(Dense(NUM_CLASSES, activation='softmax', kernel_initializer='glorot_uniform'))
opt = keras.optimizers.Adam(beta_1=.9, beta_2=.999, epsilon=1e-08)
# Let's train the model
cnn2.compile(loss='sparse_categorical_crossentropy', optimizer=opt, metrics=['accuracy'])

## 12a. More Advanced CNN Techniques as TFData.ipynb

# Warning



# Types of CNN, 2013



Karen Simonyan  
and  
Andrew Zisserman



- Oxford **VGG Net** (Visual Geometry Group)
- Major contributions:
  - small cascaded kernels
  - way more layers (19 versus ~7)
  - “emulates” biology “better” 😊
  - trained on NVIDIA GPUs for 2-3 weeks

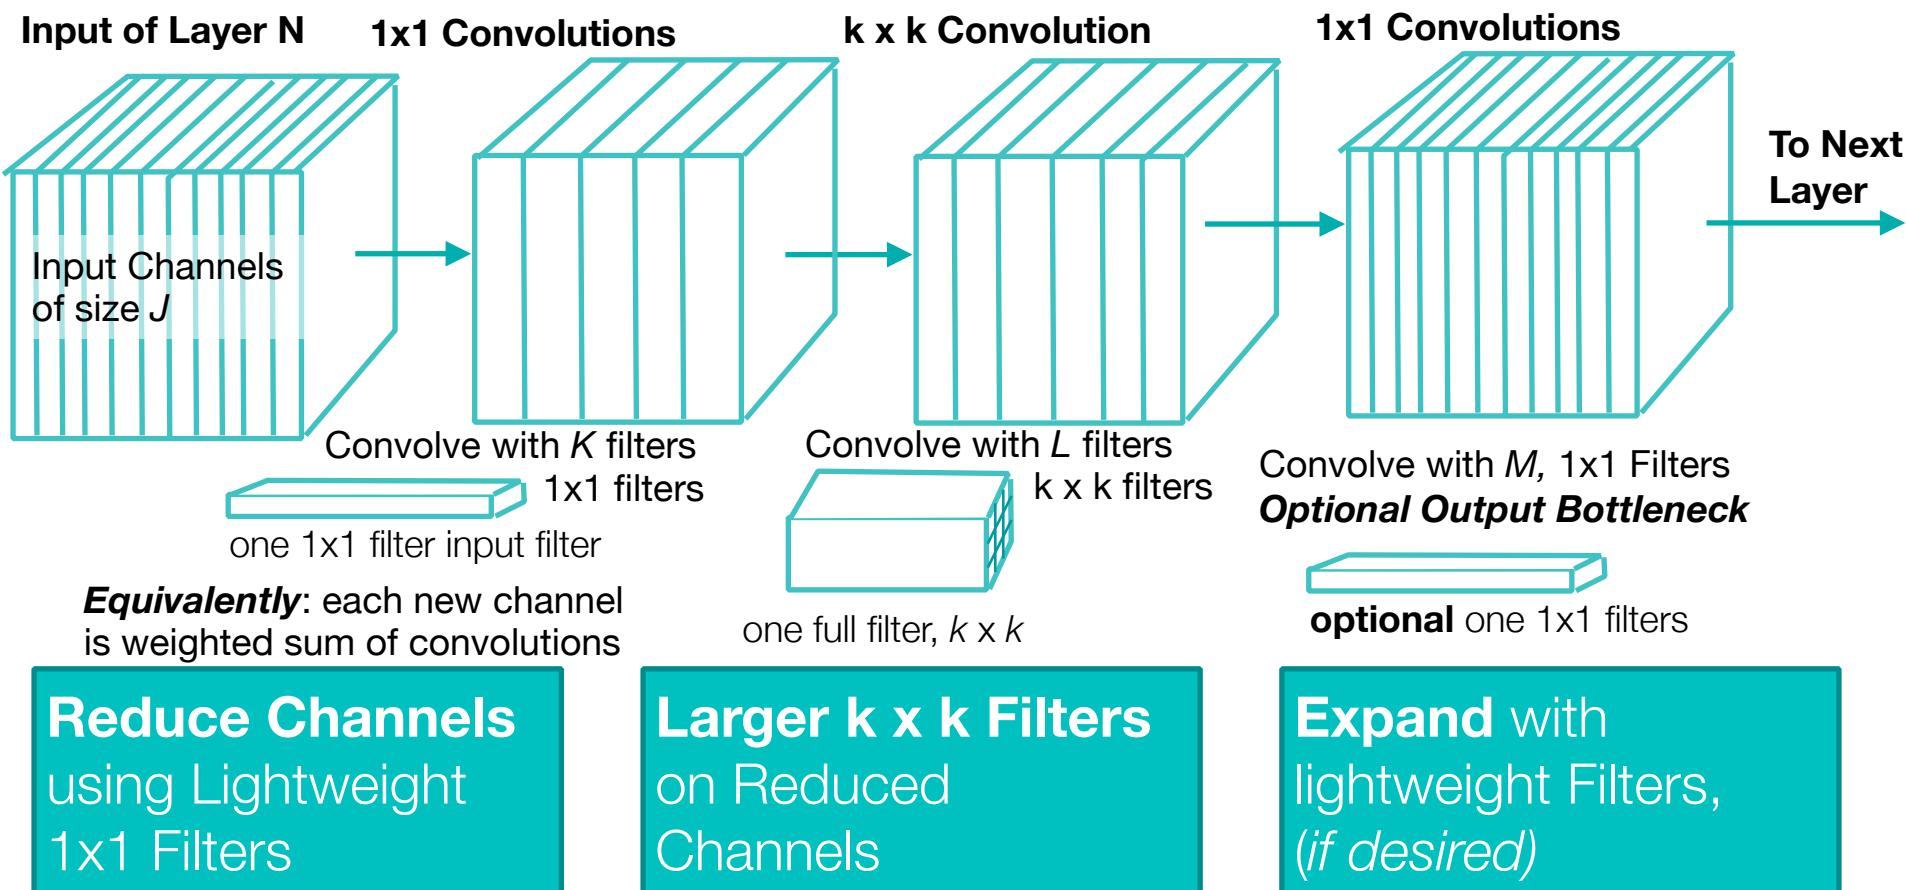
<http://www.robots.ox.ac.uk/~vgg/practicals/cnn/>

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64	conv3-64	conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128	conv3-128	conv3-128
maxpool					
conv3-256	conv3-256	conv3-256	conv3-256	conv3-256	conv3-256
conv3-256	conv3-256	conv3-256	conv3-256	conv3-256	conv3-256
maxpool					
conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Table 2: Number of parameters (in millions).

Network	A,A-LRN	B	C	D	E
Number of parameters	133	133	134	138	144

# 2014 Network in Network, expanded view



**Example:** Train 32 filters with kernel size of 3x3, input channels of 256 (outputs size of 32)

- No reduction:  $(32 \times 3 \times 3 \times 256) = 73,728$  (no bottleneck)
- Reduction to 128:  $(128 \times 1 \times 1 \times 256) + (32 \times 3 \times 3 \times 128) = 69,632$  (reduce 128)
- Reduction to 64:  $(64 \times 1 \times 1 \times 256) + (32 \times 3 \times 3 \times 64) = 34,816$  (reduce 64)

# Bottleneck, in Code

```
# we can also use sequential like a function builder,
parallel_3x3a = keras.Sequential([
    Conv2D(filters=32, kernel_size=(1,1), padding='same',
           kernel_initializer='he_uniform', kernel_regularizer=l2(1e-5)),
    Conv2D(filters=128,kernel_size=(3,3),padding='same',
           kernel_initializer='he_uniform', kernel_regularizer=l2(1e-5))
])
...
# no max pool before next conv layer!!
x = Conv2D(filters=96,
            kernel_size=(11,11),
            strides=(2,2), # add some initial downsampling
            padding='same',
            kernel_initializer='he_uniform',
            kernel_regularizer=l2(1e-6),
            activation='relu')(x)

# now save this tensor to use in multiple branches
input_conv = MaxPooling2D(pool_size=(2, 2), name='branch_point1')(x)
...
x = parallel_3x3a(input_conv)
...
```

## 12a. More Advanced CNN Techniques as TFData.ipynb 55

# Types of CNN, 2014 (parallel pathways)

Research at Google

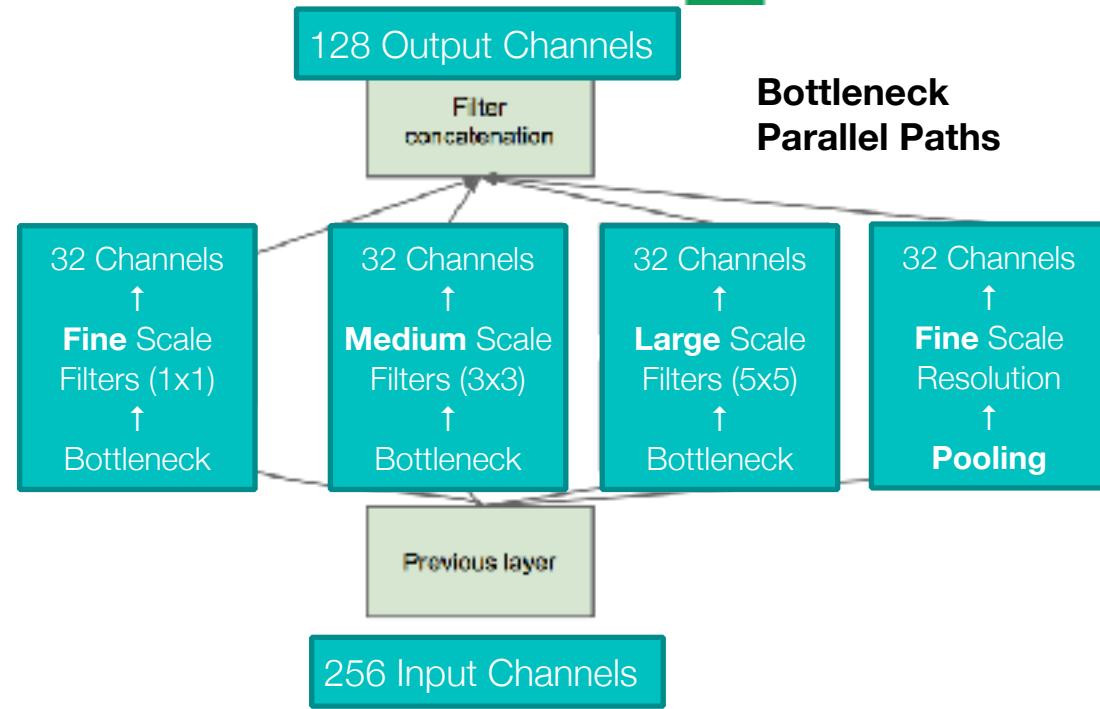
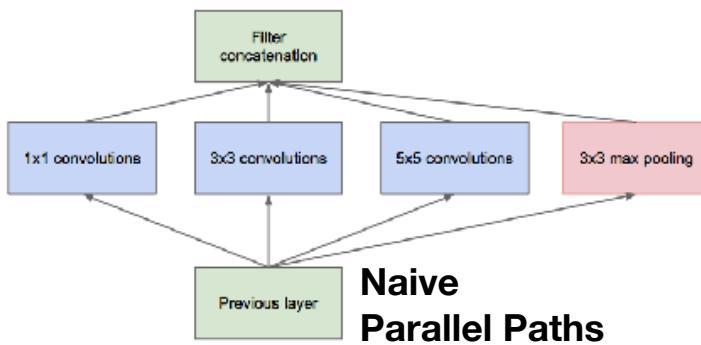
Home Publications People Teams Outreach

Christian Szegedy



Research Area(s)  
Machine Intelligence  
Machine Perception

- GoogLeNet
  - or Inception V1
- Major contribution:
  - multi-scale filtering
  - parallel NiN



<https://arxiv.org/pdf/1409.4842.pdf>

# Types of CNN, 2015 February and December

Research at Google

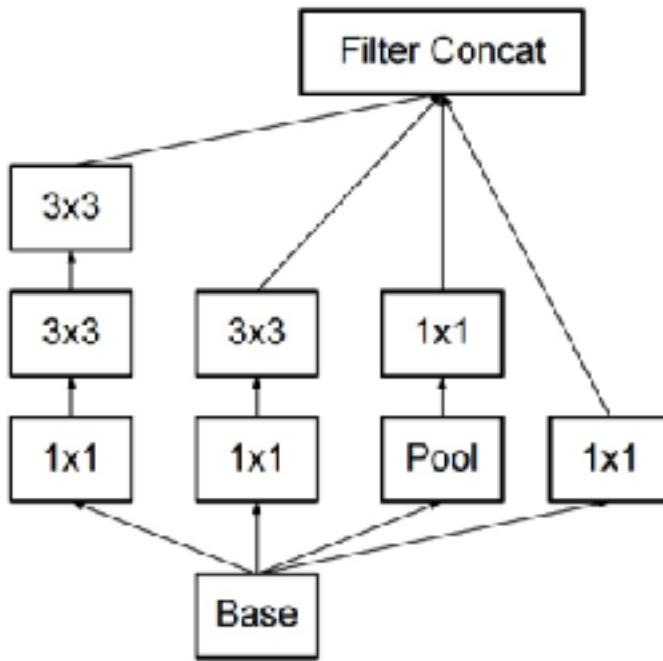
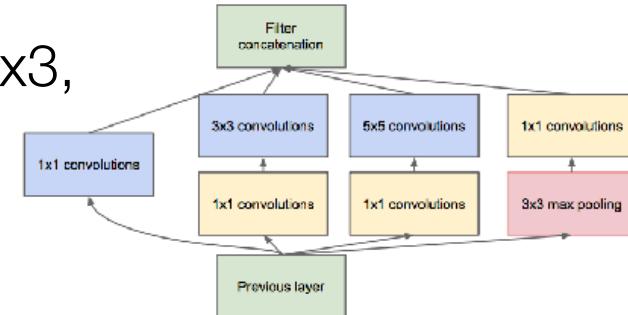
Home Publications People Teams Outreach

Christian Szegedy

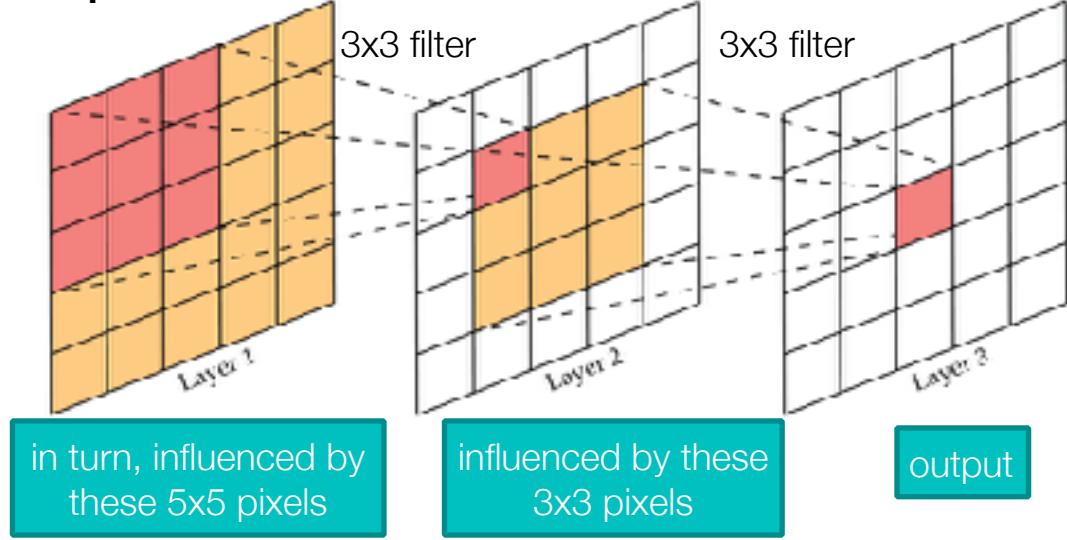


Research Area(s)  
Machine Intelligence  
Machine Perception

- **Inception V2,**
  - Inception V1 with different normalization
- **Inception V3:**
  - replace 5x5 with multiple 3x3,  
maintains receptive field



Receptive Field



<https://arxiv.org/pdf/1512.00567.pdf>

<https://medium.com/@rekalantar/receptive-fields-in-deep-convolutional-networks-43871d2ef2e9>

# Parallel Paths, in Code

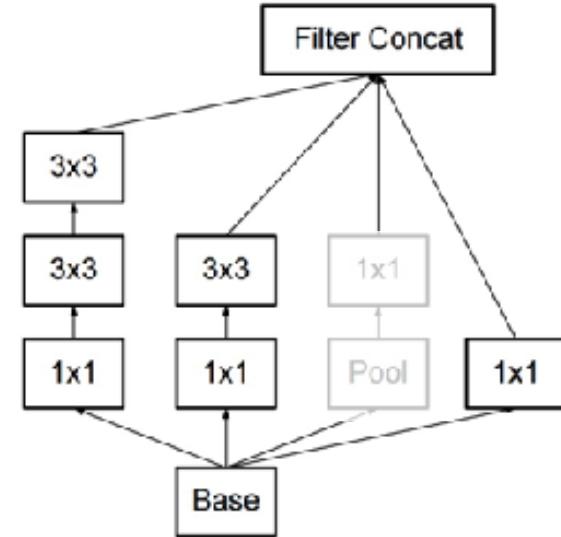
```
# we can also use sequential like a function builder,
parallel_3x3a = keras.Sequential([
    Conv2D(filters=32, kernel_size=(1,1), padding='same',
           kernel_initializer='he_uniform', kernel_regularizer=l2(1e-5)),
    Conv2D(filters=128,kernel_size=(3,3),padding='same',
           kernel_initializer='he_uniform', kernel_regularizer=l2(1e-5))
])

parallel_3x3multa = keras.Sequential([
    Conv2D(filters=32, kernel_size=(1,1), padding='same',
           kernel_initializer='he_uniform', kernel_regularizer=l2(1e-5)),
    Conv2D(filters=64,kernel_size=(3,3),padding='same',
           kernel_initializer='he_uniform', kernel_regularizer=l2(1e-5)),
    Conv2D(filters=64,kernel_size=(3,3),padding='same',
           kernel_initializer='he_uniform', kernel_regularizer=l2(1e-5))
])

conv_1x1a = Conv2D(filters=64,kernel_size=(1,1), padding='same',
                   kernel_initializer='he_uniform', kernel_regularizer=l2(1e-6))

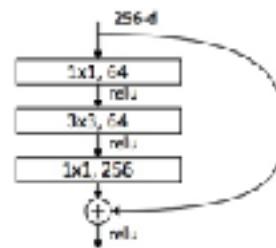
...
# now save this tensor to use in multiple branches
input_conv = MaxPooling2D(pool_size=(2, 2), name='branch_point1')(x)

# now place into parallel input branches
branches = []
x = conv_1x1a(input_conv)
branches.append(x)
x = parallel_3x3a(input_conv)
branches.append(x)
x = parallel_3x3multa(input_conv)
branches.append(x)
# that's it, we just need to average the results
x = Concatenate(axis=-1,name='ens_concat1')(branches)
```

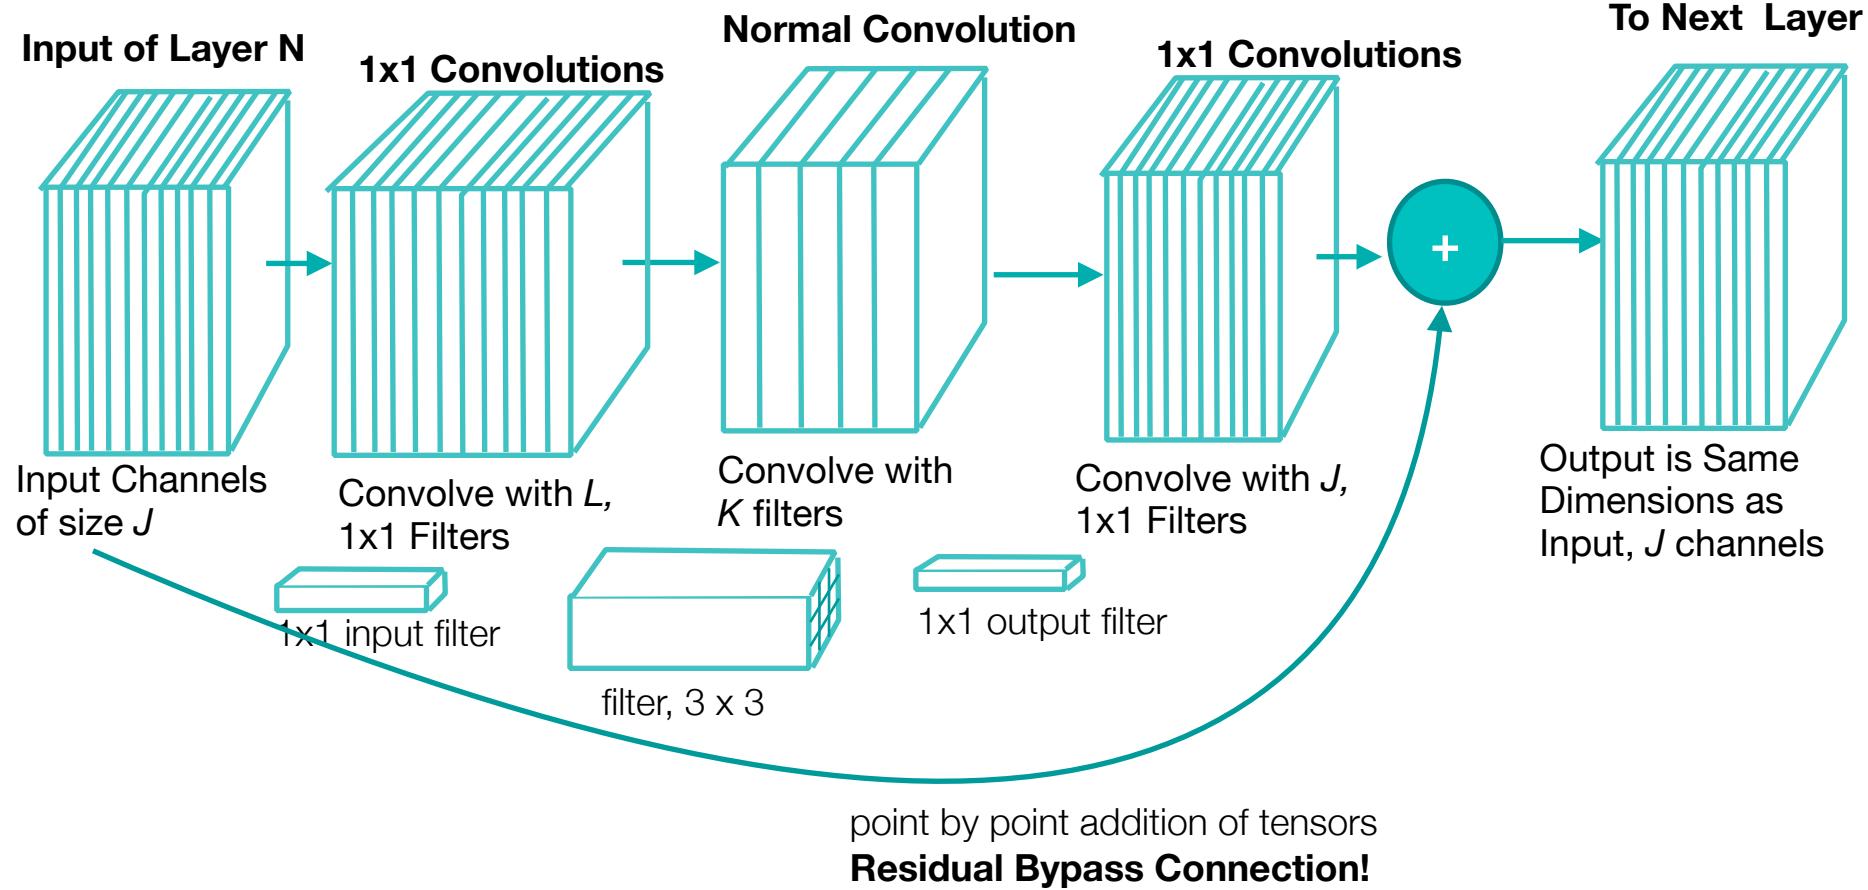


# Types of CNN, 2015 December

- **ResNet**, Major Contributions:
  - bypass pathways
  - “bio-plausible” with feedback 😊



**Back Propagation:** Two paths, including one without ANY operations that cause the gradient to vanish...



# Residual Paths, in Code

```
concat_filt_size = 128

residual_3x3a = keras.Sequential([
    Conv2D(filters=32, kernel_size=(1,1), padding='same',
           kernel_initializer='he_uniform', kernel_regularizer=l2(1e-5)),
    Conv2D(filters=32,kernel_size=(3,3),padding='same',
           kernel_initializer='he_uniform', kernel_regularizer=l2(1e-5)),
    Conv2D(filters=concat_filt_size,kernel_size=(1,1),padding='same',
           kernel_initializer='he_uniform', kernel_regularizer=l2(1e-5))
])

residual_3x3b = keras.Sequential([
    ...
    x = MaxPooling2D(pool_size=(2, 2), name='branch_point1')(x)

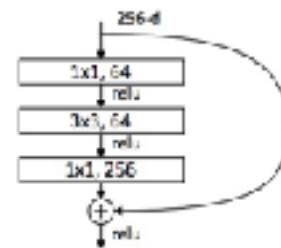
    x_split = Conv2D(filters=concat_filt_size, kernel_size=(1,1), padding='same',
                      kernel_initializer='he_uniform', kernel_regularizer=l2(1e-5))(x)

    x = residual_3x3a(x_split)
    # now add back in the split layer, x_split (residual added in)
    x = Add()([x, x_split])
    x = Activation("relu")(x)

    x_split = MaxPooling2D(pool_size=(2, 2))(x)

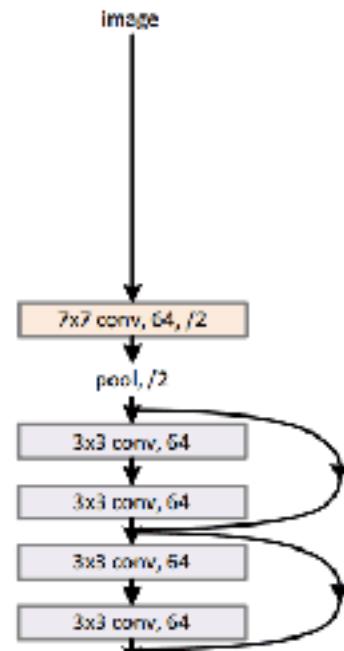
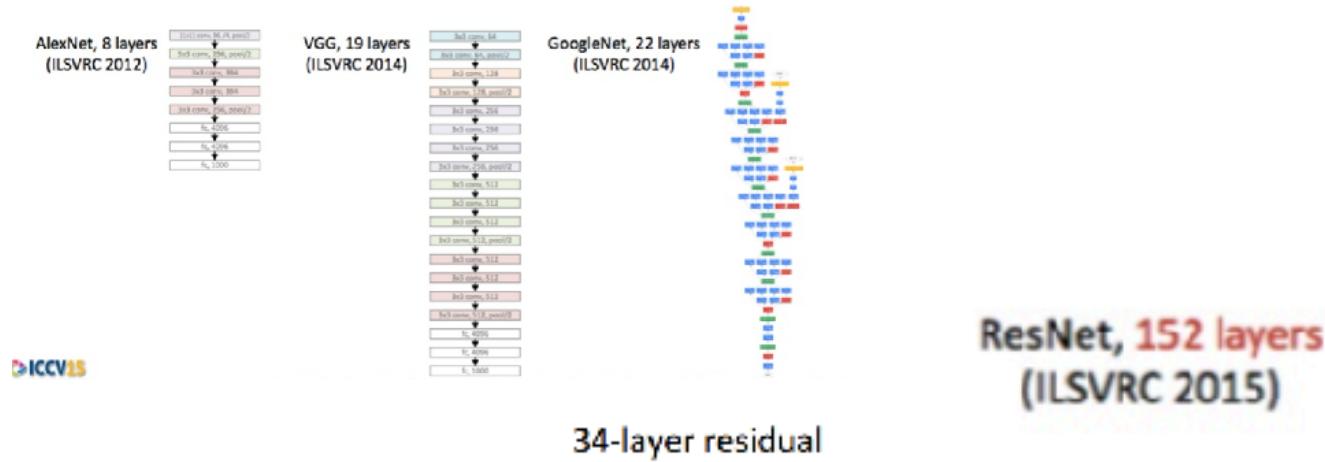
    x = residual_3x3b(x_split)
    # now add back in the split layer, x_split (residual added in)
    x = Add()([x, x_split])
    x = Activation("relu")(x)

    x = MaxPooling2D(pool_size=(2, 2))(x)
```



# How big are these networks?

# How big are these networks?



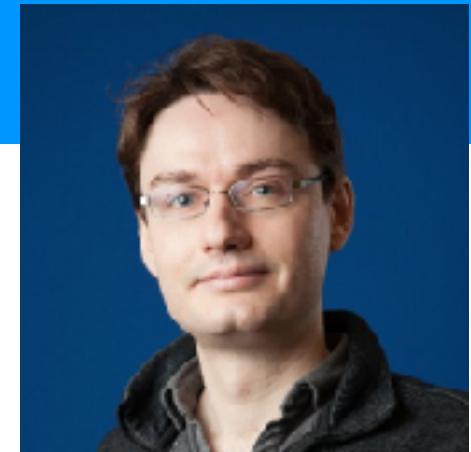
# Transition Period in Convolutional Networks

- 2012 - 2017:
  - Add more layers! 😜
  - How can we train it even deeper? 😎
  - Can we run out of memory? Let's try! 🙄
- 2017-2021:
  - How can we get similar performance with reduced parameters? 🤔
  - How should the number of parameters scale for competing resource? Is there an optimum scaling for a given set of resources? ↘
- 2021 - present: Transformers...

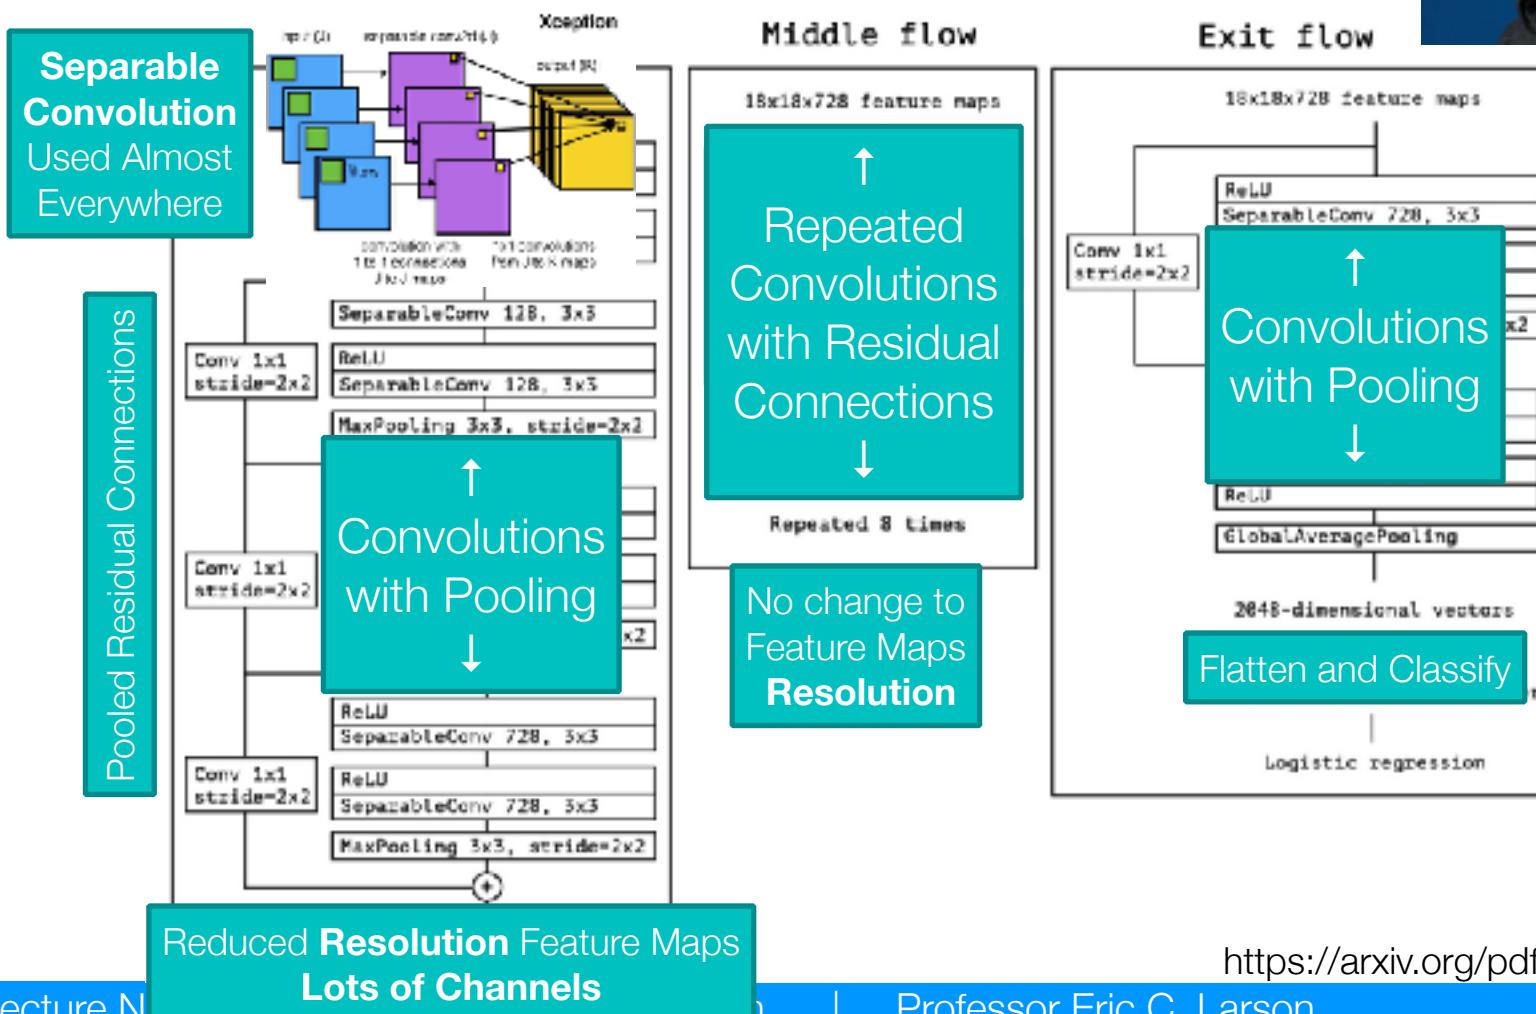
# Types of CNN, 2017

## Xception

- Major Contributions:
- combining branching / residual blocks
- separable convolutions (fewer trainable params)



Francois Chollet  
Google

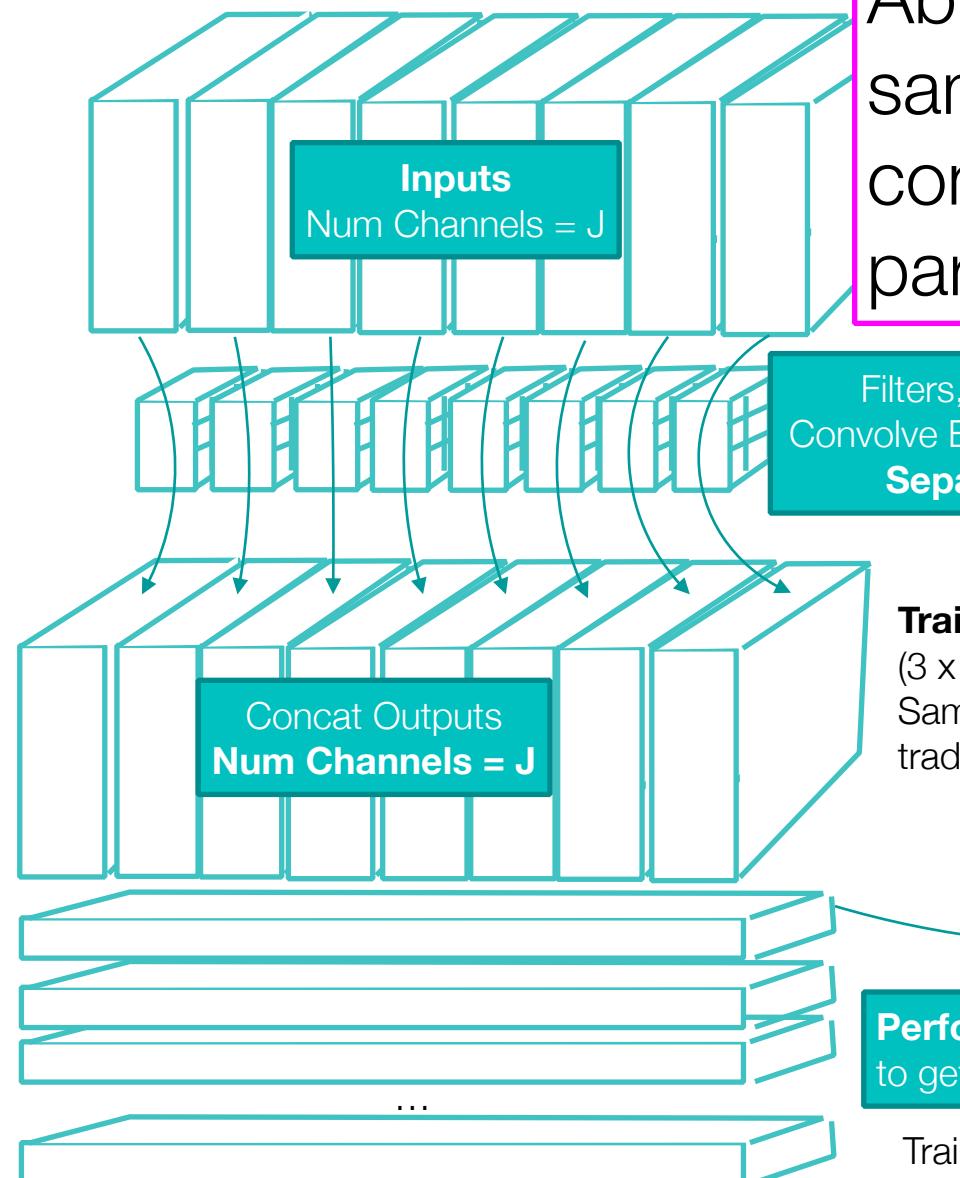


<https://arxiv.org/pdf/1610.02357.pdf> 64

# Separable Convolution

## Separable Convolution:

Able to represent many of the same features as traditional convolution, but with far fewer parameters



Filters, Layer N  
Convolve Each Channel  
**Separately**

**Trainable params:**  
 $(3 \times 3 \times J)$   
Same as *one filter* in  
traditional convolution!

**Perform K (1x1) Convolutions**  
to get K Outputs

Trainable params:  $K \times J$

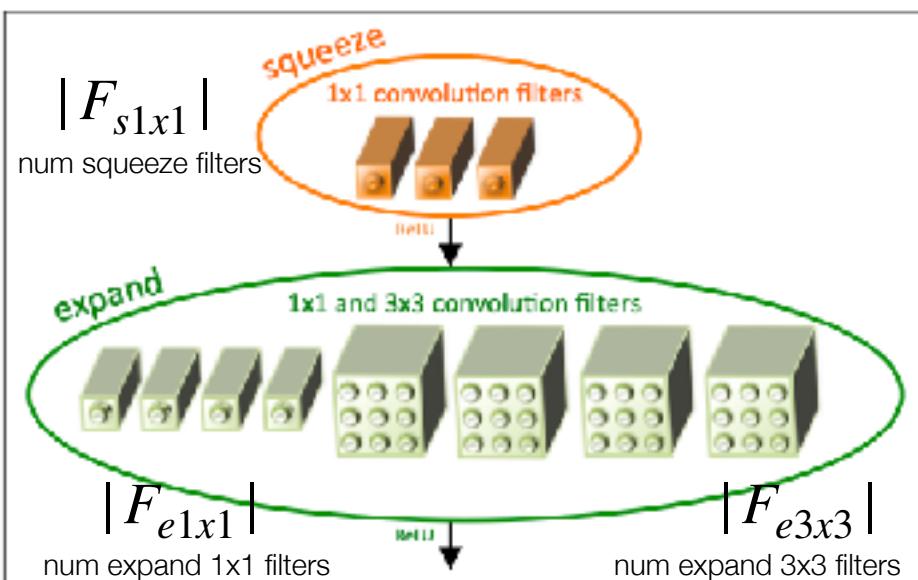
convolution with  
1 to 1 connections  
 $J$  to  $J$  maps

1x1 convolutions  
from  $J$  to  $K$  maps

**K Outputs**

# SqueezeNet (2018)

- How to squeeze and expand in each layer, with residual pathways?
  - Decide bottleneck size from a ratio, *what ratio is best?*
  - Use mostly 1x1 filters, *how many versus 3x3?*

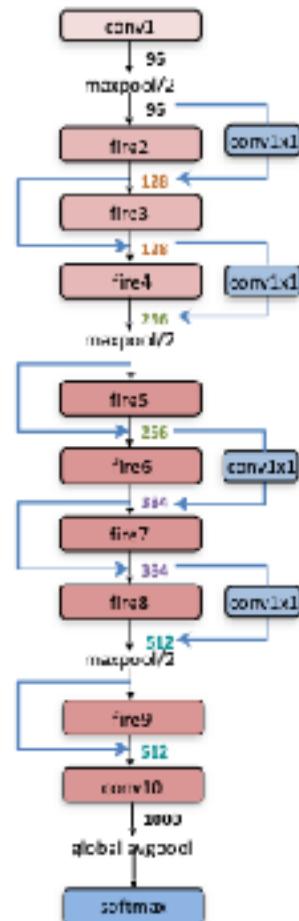


$$SR = \frac{|F_{s1x1}|}{|F_{e1x1}| + |F_{e3x3}|}$$

Controls how much to bottleneck

$$PCT_{3x3} = \frac{|F_{e3x3}|}{|F_{e1x1}| + |F_{e3x3}|}$$

Controls num filter params  
(how many 1x1 versus 3x3)



SQUEEZE NET: ALEXNET-LEVEL ACCURACY WITH 50X FEWER PARAMETERS AND <0.5MB MODEL SIZE

Forrest N. Iandola<sup>1</sup>, Song Han<sup>2</sup>, Matthew W. Moskewicz<sup>1</sup>, Khalid Ashraf<sup>1</sup>,  
William J. Dally<sup>2</sup>, Kurt Keutzer<sup>1</sup>  
<sup>1</sup>DeepScale & UC Berkeley   <sup>2</sup>Stanford University  
{forresti, moskewicz, kashraf, keutzer}@eecs.berkeley.edu  
{songhan, dally}@stanford.edu

**In paper, Good Performance when:**

- SR is 12.5% up to 100%
- PCT<sub>3x3</sub> is 25% up to 100%

# Efficient Net (2019)

Start with something

**Observation 1** – Scaling up any width, depth, or resolution improves accuracy gain diminishes for bigger models

**Observation 2** – In order to pursue efficiency, it is critical to balance all width, depth, and resolution during scaling

Depth Scaling

**Resolution Scaling:** For input different resolutions

$$\text{depth: } d = \alpha^\phi$$

$\phi$  user specified scaling coefficient

$$\text{width: } w = \beta^\phi$$

$$\text{res.: } r = \gamma^\phi$$

$$\text{s.t. } \alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$$

$$\alpha, \beta, \gamma \geq 1$$

$$\begin{aligned}\alpha &= 1.2 \\ \beta &= 1.1 \\ \gamma &= 1.15\end{aligned}$$

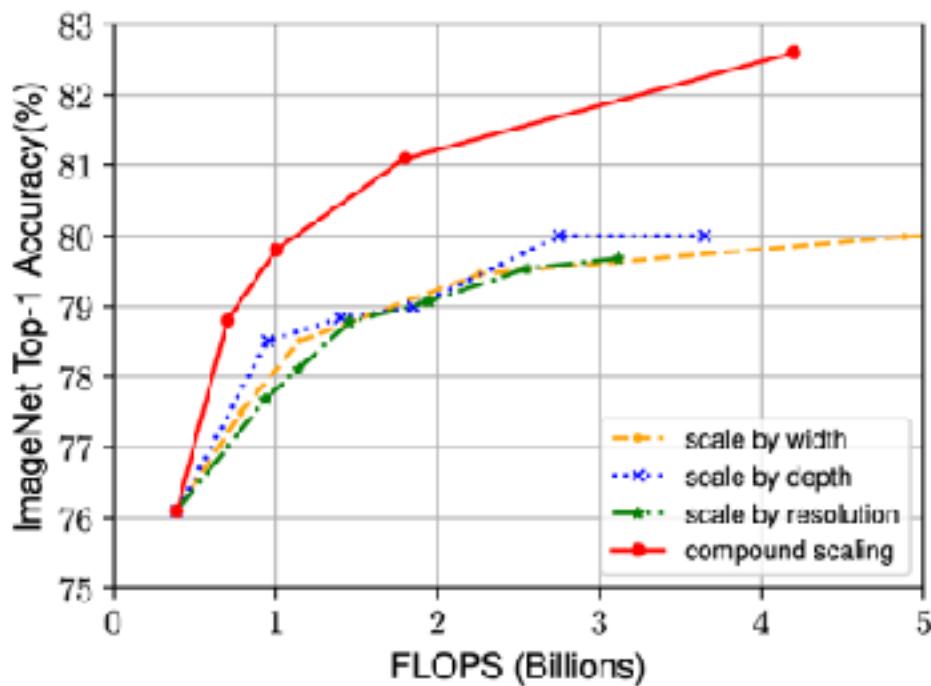


Figure 8. Scaling Up EfficientNet-B0 with Different Methods.

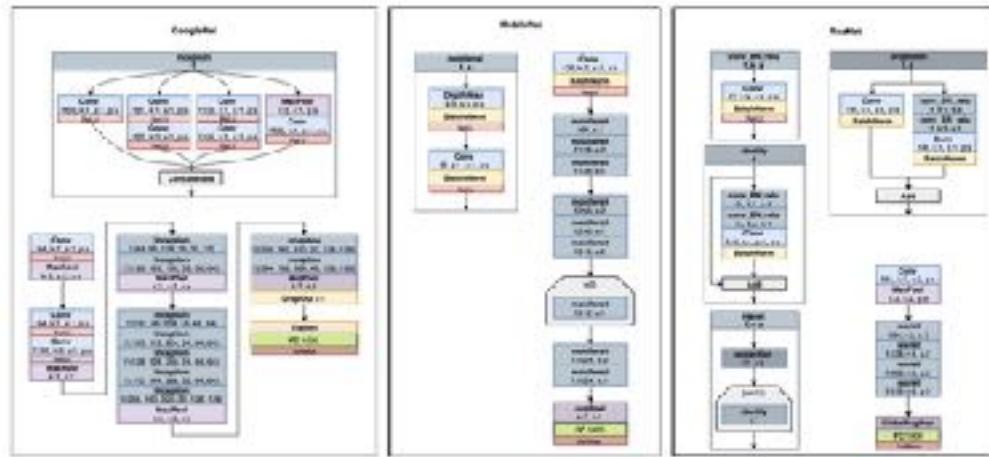
where  $\alpha, \beta, \gamma$  are constants that can be determined by a small grid search. Intuitively,  $\phi$  is a user-specified coefficient that controls how many more resources are available for model scaling, while  $\alpha, \beta, \gamma$  specify how to assign these extra resources to network width, depth, and resolution respectively.

optimal values found in paper!

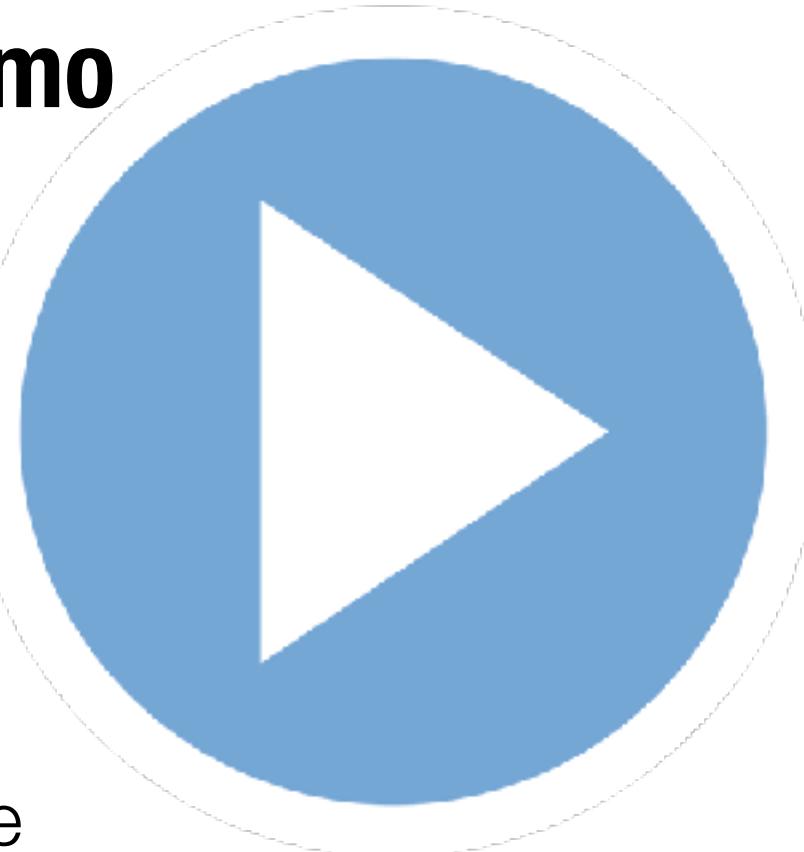
<https://arxiv.org/pdf/1905.11946v5.pdf>

# More Modern CNN Architectures

Even more Convolutional  
Neural Networks  
...in TensorFlow  
...with Keras



# Mostly Self Guided Demo



Transfer Learning Covered Here

## 12. More Advanced CNN Techniques as `TFData.ipynb`

## Next Time:

- Intro to Sequential Neural Network Architectures
  - Word Embeddings, 1D CNNs, Transformers